Rel-Sortability of Time-Unrolled Causal Graphs

Quang-Vinh Dang[1*]

[1*]British University Vietnam, Hung Yen, Vietnam.

Corresponding author(s). E-mail(s): vinh.dq4@buv.edu.vn;

Abstract

Hiding the time labels of an unrolled process does not hide the time order. In a time-unrolled DAG built from a summary graph with self-loops, the relatives of a variable leave a graph-determined footprint of time. We follow this footprint through four questions. First, we compute the rel-sortability of every edge exactly, for every summary graph and horizon, with its limit and a burn-in (the largest finite distance in the summary graph) after which the value no longer depends on the start; the graph statements are machine-checked in Lean 4. The burn-in is minimal for directed chains (proved for relatives, computed for values) and can be too long for other graphs. Second, a fixed start inflates the footprint: strongly connected summary graphs are not protected at short horizons, and the fixed start raised the order accuracy of a relatives-sorting baseline by 0.13. Third, we state when relatives can be estimated from one autocorrelated record. Fourth, controlled tests show that the gain follows the correlations a test can detect, not the population relatives pattern, and that the share of the exact gap that is realised grows with sample size (0.18 at $\boldsymbol{n = 200}$, 0.76 at $\boldsymbol{n = 25{,}000}$). Tuned DAGMA, NOTEARS and GOLEM-NV overtook FLOP on in-trees under the fixed start and lost after a burn-in; variance is involved but does not explain it alone. Real-data support is mostly negative. Results are for lag-one dynamics; no deployed generator was found to use the construction as shipped. We propose no new algorithm.

Keywords: causal discovery, sortability, time series, benchmark simulation, formal verification, finite-sample theory

1 Introduction

Simulated benchmarks drive much of causal discovery, and several of their regularities can be exploited without any causal reasoning. Marginal variances tend to grow along

the causal order (*var-sortability*; Reisach, Seiler, & Weichwald, 2021), so a method that sorts by variance does well on data from common generators, and it explains the strong benchmark performance of some continuous-optimisation methods, a family that includes NOTEARS (X. Zheng, Aragam, Ravikumar, & Xing, 2018). The coefficient of determination R^2 has the same property and is scale invariant (Reisach, Tami, Seiler, Chambaz, & Weichwald, 2023); internally standardised structural models are not var-sortable by construction and, empirically, mostly not R^2-sortable (Ormaniec, Sussex, Lorch, Schölkopf, & Krause, 2025); and Herman, Wahl, Ninad, and Runge (2025) relate the R^2 pattern to colliders and confounders and propose a generator that avoids it. Most recently, Reisach, Chambaz, Blanchard, and Weichwald (2026) defined *rel-sortability*, which orders a node by the number of its *relatives* (nodes joined to it by an open path, marginally dependent on it under faithfulness). Relatives are a property of the graph alone, so the criterion cannot be removed by reweighting. For random DAGs it is high; for *time-unrolled* DAGs built from a summary graph with self-loops, Reisach et al. (2026, Thm. 4) show that it tends to 1/2 if every weakly connected component of the summary graph is strongly connected, and they propose such graphs as benchmarks. Only sufficiency is shown in the arXiv version; they do not analyse the value beyond this case. The preprint fixes neither a horizon nor an initial law for the unrolled construction, and its data code is not public.

The thesis. Removing the time labels from an unrolled process does not remove the time order. The order leaves a footprint in label-free data: which variables are relatives of which, a property of the unrolled graph alone. *Time order survives label removal as a graph-determined footprint; the footprint is exactly computable, depends on how a simulation starts, is only partly visible at realistic sample sizes, and measurably changes what label-free learners appear to do.* The paper follows this one object through four questions, and they give it its order. **Q1:** what is the footprint exactly, and what does it identify when the labels are hidden? **Q2:** how does the start of a simulation shape it? **Q3:** how much of it can an estimator see in one finite, autocorrelated record? **Q4:** what does it do to label-free learners and to real data? The questions matter for real time series, whose summary graphs are often feed-forward (rivers, supply chains, signalling cascades), and for anyone who builds a benchmark from an unrolled graph.

Contributions. *(Q1) The footprint.* The exact rel-sortability of every edge of the time-unrolled DAG G_T, for every summary graph with self-loops and every T (Theorem 1); its limit $\frac{1}{2} + |E_\times|/(2|E_\mathrm{S}|)$ with the rate $\min(\Delta, T-1)/(T-1)$ (Theorem 2); and the necessity of the condition of Reisach et al. (2026, Thm. 4) (Corollary 3). These statements are machine-checked in Lean 4 (Section 3.5). When the labels are hidden, the slice of every column is a graph invariant of the unrolled DAG, its Markov equivalence class is explicit, and under a fixed start and faithfulness the unrolled DAG is identified from second moments up to the reversal of isolated units, with three exact non-identifiability results and one open conjecture (Section 3.4). These are paper proofs, conditional on faithfulness, and no human has read them yet. *(Q2) The start.* The excess for chains and out-trees is a boundary effect of the deterministic start (Proposition 4). After a latent burn-in of $B \geq \Delta$ slices, the value is fixed by the *source* strongly connected components for every T (Theorem 5 and Corollary 6).

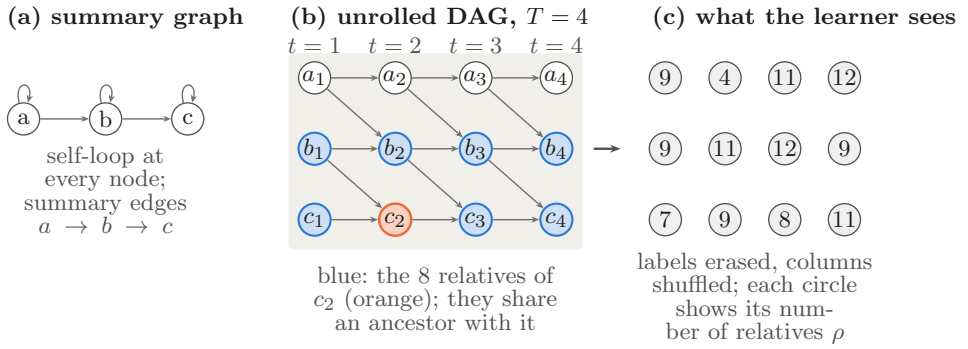


Fig. 1 The idea in one picture. (a) A summary graph with a self-loop at every node. (b) Its time-unrolled DAG G_T for $T = 4$: one column per slice, one row per variable. The *relatives* of a node are the nodes that share an ancestor with it; the blue nodes are the eight relatives of c_2. (c) A label-free learner receives only the columns of the data, in unknown order. Counting the partners of each column (here, exactly, the number of relatives ρ) still carries the time order: along every edge of (b) the count either grows or stays equal (for example $c_1 \to c_2$: $4 \to 8$). The rest of the paper asks how exactly this footprint can be computed, how a simulation's start changes it, and how much of it a learner sees.

This burn-in is sufficient. It is exactly minimal for the one-edge graph (Lean) and, by an exact argument on paper and exact enumeration, for directed chains at every length computed (Proposition 19); it is not minimal in general: in an enumeration of 726 random summary graphs, $B = \Delta - 1$ was insufficient in 445 and sufficed in 281, and in a second exact computation on 396 random graphs the minimal burn-in was Δ in 51% of the cases and up to four slices shorter in the rest. We have no closed form for the minimal burn-in of a given graph (Remark 20). Strongly connected summary graphs are not protected at short horizons: a ring has $s_\rho(G_T) = 1$ for all $T \le p$. *(Q3) Visibility.* An effective sample size for the relatives test, the saturation of relatives counts by a latent driver, and the effect of block means and subsamples (Section 4). The proofs are on paper, the non-asymptotic bounds are loose, and a fixed truncation lag controls the size only away from a unit root. *(Q4) Learners and real data* (Section 5). With the time labels hidden, the fixed start raised the order accuracy of a relatives-sorting baseline by 0.13 in simulations, and the gain replicated under nonlinear dynamics. All 42 simulated conditions have a positive exact gap and a positive gain, so agreement in sign does not test the theory. A designed test (Section 5.2) separates two readings. Data with the population relatives of a burn-in but the detectable correlations of the fixed start kept 0.96 of the gain, and the share of the exact gap that is realised grew with the sample size (0.18 at $n = 200$, 0.76 at $n = 25{,}000$) and with the spectral radius. The exact gap is therefore the large-sample ceiling, and the effect realised at a finite sample size is its detectable image. On raw data, tuned DAGMA, NOTEARS and GOLEM-NV overtook FLOP on in-trees under the fixed start and lost to it after a burn-in, which our plan had predicted would not happen. Matching the marginal variances across the two starts removed the advantage, which points to variance, but imposing the fixed-start variances on burn-in data restored less than half of it, so the cause is not fully explained and our theory does not cover it. The real-data studies ask a different question (can relatives be estimated from one record?) and mostly failed

or did not discriminate; the positive results come from one device, and at the level of independent runs or stations they cannot reach conventional significance.

Scope and status. The results are for lag-one unrollings with a self-loop at every node, and the simulations use $n = 1000$. The unrolled construction is a proposal of Reisach et al. (2026); our survey of common generators found none that hands slices to a method that ignores time order as shipped (Section 5.1). We propose no new causal-discovery method. The graph results of Q1 are short consequences of a walk-length lemma and of Reisach et al. (2026, Lemma 1); their value is that they settle the value for every summary graph, horizon and start. The finite-sample results use textbook tools. Table 1 lists the status of each claim.

Organisation. Section 2 places the work among studies of benchmark regularities, time-series causal discovery and identification. Section 3 answers Q1 and gives the theory part of Q2: the exact values, the effect of a burn-in, the Lean verification and the identifiability statements for hidden labels. Section 4 answers Q3 and treats block means and subsamples. Section 5 answers the empirical part of Q2 and Q4 with the benchmark experiments, the controlled tests of the mechanism and the real-data studies, and Section 6 gives recommendations, limitations and the statement on the use of AI tools.

2 Related Work and Background

The question of this section is where the four questions of Section 1 stand in the literature. The paper touches three bodies of work: studies of what simulated benchmarks reward, time-series causal discovery together with the generators that supply its test data, and the identification and finite-sample theory that the later sections use. We discuss each in turn and then state what is new here and what is not.

2.1 Sortability and the evaluation of continuous-optimisation learners

Continuous-optimisation learners cast structure learning as numerical optimisation over weighted adjacency matrices. NOTEARS (X. Zheng et al., 2018) replaces the combinatorial acyclicity constraint by a smooth function of the matrix, and later methods change the constraint or the objective: DAGMA (Bello, Aragam, & Ravikumar, 2022) uses a log-determinant characterisation of acyclicity, and GOLEM (Ng, Ghassami, & Zhang, 2020) maximises a likelihood with soft penalties. Constraint-based search (Spirtes, Glymour, & Scheines, 2000) and greedy score-based search (Chickering, 2002) remain the classical alternatives, and FLOP (Wienöbst, Henckel, & Weichwald, 2026), a recent discrete-search method, is the strongest method without time labels in our experiments (Section 5.1). All of these are developed and compared mostly on simulated linear or additive-noise models with random graphs. A regularity of the generator that tracks the causal order can therefore change a ranking without any causal reasoning on the part of the learner, and a line of work has identified such regularities one by one.

Reisach et al. (2021) showed that in common generators the marginal variance tends to grow along the causal order (var-sortability). A baseline that sorts by variance

and then regresses is competitive on such data, and continuous-optimisation methods lose accuracy once the data are standardised. Kaiser and Sipos (2022) reach a related conclusion from the side of the method: NOTEARS is not scale invariant, it aims at a parsimonious DAG that explains the residual variance, and the authors conclude that it is unsuited to identifying causal relationships from data with dimensional quantities. Reisach et al. (2023) introduced R^2-sortability, the same kind of regularity measured with a scale-invariant quantity, and a corresponding sorting baseline. Ormaniec et al. (2025) study internally standardised structural causal models, which are not var-sortable by construction and, empirically, mostly not R^2-sortable. Herman et al. (2025) relate the R^2 pattern to colliders and confounders and propose a unit-free generator that avoids it and whose time-series extension starts from the exact stationary distribution. Reisach et al. (2026) then defined rel-sortability, which counts the relatives of a node, that is, the nodes joined to it by an open path. Relatives depend on the graph alone, so reweighting cannot remove the criterion. For time-unrolled DAGs built from a summary graph with self-loops they proved that the criterion tends to $1/2$ when every weak component of the summary graph is strongly connected, and they proposed such graphs for benchmarks. Their statement is a sufficient condition. This paper completes it: the exact value of every edge for every summary graph and horizon, the exact limit, the effect of a burn-in, and the question of whether the value matters for what learners recover.

For time series, Lohse and Wahl (2025) adapt var- and R^2-sortability to summary graphs, observe that cyclic pairs push the signal towards $1/2$, and measure both on river-flow and Causal Chambers data. They give no theorem and do not study relatives or the way a record is observed. Robustness of score-matching learners to violated model assumptions has been studied by Montagna et al. (2023), which is a different mechanism from the generator regularities above but leads to the same practical caution about reading benchmark rankings.

Several studies question how far simulated benchmarks resemble the data that methods meet in practice. Gentzel, Garant, and Jensen (2019) argue for evaluating causal models with interventional measures and empirical data, and Brouillard et al. (2025) review the data sets used in causal discovery and how they relate to real applications. On the side of simulation design, Andrews and Kummerfeld (2026) propose a DAG adaptation of the onion method, and Nagarajan and Scutari (2025) show for Bayesian structure learning that the topology of the network changes benchmark outcomes; Vowels, Camgoz, and Bowden (2022) survey the structure-learning methods that these comparisons concern. Our subject is narrower. We take one proposed generator, the time-unrolled construction, and ask what its start does to the quantities that sorting baselines exploit.

Our SPADE paper (Dang, Le, & Dinh, 2026) introduces a B-spline additive-noise DAG learner and compares it with DAGMA, GraN-DAG, GOLEM, NOTEARS, SCORE, NoGAM, PCMCI and VAR-LiNGAM on simulated nonlinear DAGs, with mixed results against NoGAM. It also reports a graph learned from a financial panel in levels that collapsed on stationary returns, a reminder that a learned graph can reflect shared non-stationarity as well as mechanism, which is the situation that the fixed start creates in simulation.

2.2 Time-series causal discovery and its generators

Methods that use time order go back to Granger causality (Granger, 1969). VAR-LiNGAM (Hyvärinen, Zhang, Shimizu, & Hoyer, 2010) fits a structural vector autoregression and uses non-Gaussianity to orient the instantaneous effects. PCMCI and PCMCI+ (Runge, 2020; Runge, Nowack, Kretschmer, Flaxman, & Sejdinovic, 2019) combine lagged conditional-independence tests with a momentary-conditional-independence step, the latter also handling contemporaneous links, and are available in the tigramite package (Runge, 2026). DYNOTEARS (Pamfil et al., 2020) carries the NOTEARS formulation over to structural VAR models, and neural Granger methods (Tank, Covert, Foti, Shojaie, & Fox, 2021) replace the linear lag effect by a network. All of them are given the time index. Only a static learner applied to a time-unrolled DAG, as in Reisach et al. (2026), sees the slices as unlabelled variables, and that is the setting of this paper. Our learner extension in Section 5.1 therefore compares label-free static methods with PCMCI+ as a labelled reference.

Benchmarks for the time-indexed methods include the CauseMe platform (Runge, Bathiany, et al., 2019), the Causality 4 Climate competition (Weichwald et al., 2020) and CausalTime (Cheng et al., 2024), which builds its data from fitted models of real series. A data-generating process for time-series causal discovery with explicit lag structure is described by Lawrence, Kaiser, Sampaio, and Sipos (2021). Real records with partial ground truth are rarer; we use CausalRivers (Stein, Shadaydeh, Blunk, Penzel, & Denzler, 2025) and the Causal Chambers (Gamella, Peters, & Bühlmann, 2025) in Section 5.3. Assaad, Devijver, and Gaussier (2022) survey and evaluate time-series discovery methods. Reading a dynamical system as an unrolled causal graph has its own foundations: Mooij, Janzing, and Schölkopf (2013) relate ordinary differential equations to structural causal models in the deterministic case.

How a simulation starts a vector autoregression is old practice in econometrics. Textbook treatments discuss the initial condition and the choice between conditioning on it and drawing it from the stationary distribution (Hamilton, 1994; Kilian & Lütkepohl, 2017; Lütkepohl, 2005), and the effect of the initial condition on inference for persistent processes has a literature of its own (Müller & Elliott, 2003). The fixed start of the unrolled construction is the case of a zero initial state. What this paper adds is a statement about the dependence structure: a burn-in of Δ slices equalises the relatives of the nodes, whereas making the covariance stationary needs a burn-in that matches the spectral radius (Section 5.1). We did not find a treatment of this distinction for causal-discovery benchmarks.

How these generators start a trajectory matters for the artefact studied here. We read the source code of the simulators in common use (Appendix T). Some return the series from a fixed start with no burn-in; some discard between a few steps and several hundred, or a fraction of the series length; and the unit-free generator of Herman et al. (2025) draws the initial lags from the stationary distribution. Whether the start affects a result also depends on what is done with the data. A single record passed to a time-series method is outside the threat model, whereas the time-unrolled construction of Reisach et al. (2026) gives every slice, including those right after the start, to a method that ignores time order. Section 5.1 states the two conditions and applies them to each generator.

A related issue is the observation scheme. Block means and subsamples of a faster process change the model that the analyst sees (Breitung & Swanson, 2002; A.A. Weiss, 1984) and its identifiability (Gong, Zhang, Schölkopf, Glymour, & Tao, 2017; Gong, Zhang, Schölkopf, Tao, & Geiger, 2015; Hyttinen, Plis, Järvisalo, Eberhardt, & Danks, 2016). We ask what they do to the three sortability criteria in Section 4.2.

2.3 Tied unrolled models, time labels and identification

A time-unrolled DAG with a single transition applied at every step is a dynamic Bayesian network: a start network and a transition network that is repeated (Friedman, Murphy, & Russell, 1998). When the time labels are known, learning such a model is standard, because the slices and the tying are given. Hiding the labels turns the slice structure into something that has to be identified, which is the question of Section 3.4. The tied parameters also bear on whether the usual genericity argument for faithfulness applies. Coloured Gaussian DAG models, which tie vertex variances or edge weights across colour classes, are the closest static analogue. Boege, Kubjas, Misra, and Solus (2025) show that a generic covariance matrix of an edge-coloured DAG is faithful to the graph, and the same holds for vertex-coloured DAGs (their Propositions 5.3 and 5.4). They give an example with vertex and edge colours together that has no faithful distribution (Example 5.7) and leave open which coloured DAGs admit a faithful distribution (Question 5.8). Their common-head condition defines the subfamily on which they prove structural identifiability and run their search algorithm, and is not part of the faithfulness results. On our reading, the unrolled model ties the transition weights across slices (an edge colouring) and, when every slice has the same noise variances, the noise variances (a vertex colouring), so it is the mixed case for which faithfulness is open. We have not shown that this reading is right or that the model falls in a class covered by their positive results.

Two classical facts about time direction enter the non-identifiability results of Section 3.4. Time reversal can leave the covariance of a stationary linear process unchanged (G. Weiss, 1975), and the direction of a time series can be recovered under non-Gaussian or additive-noise assumptions that we do not make (Peters, Janzing, Gretton, & Schölkopf, 2009).

Section 3.3 asks whether the changing covariance during the transient after a fixed start carries information about the mechanism. The closest precedents are identification through heteroskedasticity (Lanne & Lütkepohl, 2008; Lewis, 2021; Rigobon, 2003), which needs the covariance of the structural shocks to change across regimes or over time, and identification of linear stochastic differential equations from temporal marginals (Guan et al., 2026; Wang, Geng, Huang, Huang, & Gong, 2023). Our remark is a short consequence of the Markov property and is stated to separate the artefact for label-free methods from information available to methods with labels.

2.4 Finite-sample behaviour of dependence tests under autocorrelation

Estimating relatives from one record means running marginal independence tests between series. The distortion of a test for the correlation of two series that uses the

nominal sample size, when the series are autocorrelated, is classical (Bartlett, 1935; Clifford, Richardson, & Hémon, 1989; Roy, 1989). In causal discovery it is usually handled by effective sample sizes (Afyouni, Smith, & Nichols, 2019; Runge, Nowack, et al., 2019) or by residual-based tests; testing independence of two series through lagged cross-correlations goes back to Haugh (1976). The marginal tests we analyse are Fisher-z tests of the kind used by the PC algorithm (Spirtes et al., 2000). For the variance of the sample correlation we use the Bartlett form, and for tail and limit statements we rely on standard tools for quadratic forms of Gaussian processes and for Toeplitz-type functionals (Arcones, 1994; Ginovyan & Taqqu, 2021; Hsu, Kakade, & Zhang, 2012; Laurent & Massart, 2000; Nourdin, Peccati, & Podolskij, 2011). Residual confounding by an unreliably measured covariate, which limits what a partial-correlation test can remove, is documented in Westfall and Yarkoni (2016). Our contribution in Section 4.1 is to combine these ingredients for the specific statistic that defines relatives, to make explicit when the plug-in estimate of the effective sample size is consistent, and to state where a fixed truncation lag breaks the size guarantee.

2.5 Formal verification

The graph results are checked in Lean 4 (de Moura & Ullrich, 2021) with Mathlib (The mathlib Community, 2020). A formally grounded agentic framework for automated research in causal inference has been described (Tan & Syrgkanis, 2026). We are not aware of a Lean formalisation of rel-sortability or of the unrolled-graph results here, and our development covers only the deterministic graph and algebra parts (Section 3.5). No human has yet checked that the Lean statements match the statements in the paper.

2.6 Positioning

New here are four things, in the order of the questions of Section 1. First, the exact per-edge value of rel-sortability of a time-unrolled DAG for every summary graph with self-loops and every horizon, its limit, and a sufficient burn-in condition, with the graph statements machine-checked; the burn-in is shown to be needed only for the one-edge graph. Second, a measurement of what the fixed start does to benchmark results, with time-stamped plans, replications and every failed prediction reported. Third, finite-sample statements about estimating relatives from one autocorrelated record, with their limits. Fourth, identifiability statements for hidden time labels, conditional on faithfulness. We do not propose a new causal-discovery method, and the finite-sample and identifiability statements rest on standard tools and on paper proofs that no human has read line by line.

Takeaway. The nearest work is the sortability line, which supplies the object (relatives) and the proposed benchmark; econometric practice and time-series generators supply the notion of a start; the tied-model and heteroskedasticity literatures frame what labels add. None of them gives the exact values or measures the effect of the start on label-free learners.

3 Time-Unrolled DAGs: Exact Values, Burn-in and Identifiability

Question. What is the footprint of time in the relatives structure of an unrolled DAG, how does the start of the simulation shape it, and what does it identify when the labels are hidden? This section collects the graph-theoretic answers. Section 3.1 fixes notation, Section 3.2 gives exact values, limits and the effect of a burn-in, Sections 3.3 and 3.4 ask what the start law and the time labels add, and Section 3.5 states what is machine-checked.

3.1 Setting and notation

For a directed graph, $\mathrm{anc}(y)$ and $\mathrm{desc}(y)$ contain y; the relatives of y are $\mathrm{rel}(y) = \mathrm{desc}(\mathrm{anc}(y)) = \{z : \mathrm{anc}(y) \cap \mathrm{anc}(z) \neq \emptyset\}$ and $\rho(y) = |\mathrm{rel}(y)|$. For an edge (x, y), $s_\rho(x, y) = 1, \frac{1}{2}, 0$ if $\rho(x) <, =, > \rho(y)$, and $s_\rho(G)$ is the mean over *edges*, as in Reisach et al. (2026, Eq. (4)) (quoted verbatim in Appendix B). A value $1/2$ means the criterion carries no orientation information.

A *summary graph* $G_\mathrm{S} = (V, E_\mathrm{S})$ has $p \geq 2$ nodes and a self-loop at every node. The *time-unrolled DAG* G_T has nodes $v(t)$, $t \in \{1, \ldots, T\}$, and an edge $(i(t), j(t{+}1))$ for each $(i, j) \in E_\mathrm{S}$ and $t < T$, so every summary edge has $T - 1$ copies, all at lag one (a VAR(1)-shaped unrolling; other lags are discussed after Theorem 5). $d(u, v)$ is the shortest-path length in G_S (self-loops ignored; $d(v, v) = 0$; ∞ if unreachable), and anc_S, rel_S are taken in G_S. SCC means strongly connected component. $E_\times \subseteq E_\mathrm{S}$ is the set of edges between different SCCs and $E_\circ = E_\mathrm{S} \setminus E_\times$. A *source* SCC has no incoming edge from another SCC, and $\mathrm{Src}(v)$ is the set of source SCCs that reach v.

3.2 Exact values, limit and burn-in

Theorem 1 (every edge, every T) *For $(i, j) \in E_\mathrm{S}$ and $1 \leq t < T$, the edge $(i(t), j(t{+}1))$ of G_T has $s_\rho = 1$ if some $u \in V$ satisfies $d(u, j) \leq t \leq d(u, i)$, and $s_\rho = 1/2$ otherwise. The value does not depend on T.*

Lean: edge_score_dist, edge_score, strictCond_iff_dist.

Sketch. Every edge of G_T advances time by one, so $u(s)$ is an ancestor of $v(t)$ iff $d(u, v) \leq t - s$ (pad a shortest path with self-loops). Hence the roots of $v(t)$ are $\{u(1) : d(u, v) \leq t - 1\}$. Since $i(t)$ is a parent of $j(t{+}1)$, $\mathrm{rel}(i(t)) \subseteq \mathrm{rel}(j(t{+}1))$, and the inclusion is strict iff $j(t{+}1)$ has a root that $i(t)$ lacks (Reisach et al., 2026, Thm. 1, Lemma 1). Full proof in Appendix B. (The Lean proof characterises rel in G_T directly and does not use their Lemma 1.)

Theorem 2 (exact limit and rate) *For every $T \geq 2$, with $L(G_\mathrm{S}) = \frac{1}{2} + \frac{|E_\times|}{2|E_\mathrm{S}|}$ and $\Delta \leq p - 1$ the largest finite shortest-path distance in G_S,*

$$0 \leq s_\rho(G_T) - L(G_\mathrm{S}) \leq \frac{\min(\Delta, T{-}1)\,|E_\circ|}{2(T-1)|E_\mathrm{S}|} \leq \frac{p-1}{2(T-1)},$$

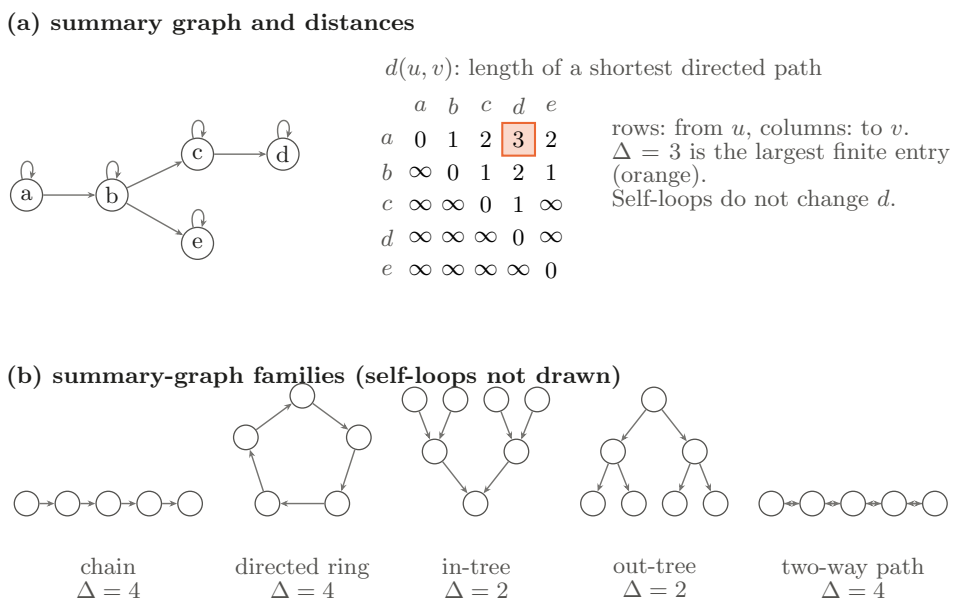


Fig. 2 Notation. (a) A summary graph with a self-loop at every node, and its table of distances $d(u, v)$ (shortest directed path; ∞ if v is not reachable from u). The largest finite entry is Δ (here 3, from a to d); it is the burn-in that Theorem 5 asks for. (b) The summary-graph families used in the paper, with their Δ: a directed chain, a directed ring (the smallest strongly connected example), in- and out-trees (a river network is an in-tree), and a two-way path.

so $\lim_{T\to\infty} s_\rho(G_T) = L(G_S)$.
Lean: sRho_unrolled_exact, sRho_unrolled_bounds_maxFinDist, tendsto_sRho_GT.

Sketch. For a crossing edge, j does not reach i; take $u = j$ in Theorem 1: all $T-1$ copies score 1. For an edge inside an SCC, i and j are reached by the same nodes, which forces $t \le d(u, i) \le \Delta$, so at most Δ copies score 1.

Corollary 3 (necessary and sufficient condition) $\lim_T s_\rho(G_T) = 1/2$ *iff every weakly connected component of* G_S *is strongly connected. Otherwise* $s_\rho(G_T) \ge \frac{1}{2} + \frac{1}{2|E_S|}$ *for every* T. *The "if" half is Reisach et al. (2026, Thm. 4). If* G_S *is acyclic apart from self-loops, with* m *other edges,* $L = \frac{1}{2} + \frac{m}{2(m+p)}$ *whatever its shape (chain, in-tree, out-tree, collider).*
Lean: tendsto_half_iff, crossEdges_eq_empty_tfae, sRho_unrolled_ge_of_cross, reisach_theorem4, limitValue_dag.

Proposition 4 (the excess is a boundary effect) *For* $t \ge p$, $\rho(v(t)) = T\,|\mathrm{rel}_S(v)| - \sum_{w\in\mathrm{rel}_S(v)} d(\mathrm{anc}_S(v), w)$.
Lean: rho_unrolled_large_t, rho_unrolled_eq.

So for a crossing edge with $\mathrm{rel}_S(i) = \mathrm{rel}_S(j)$ (root to child in an out-tree, a chain), the strict order rests on $O(1)$ extra relatives out of $\Theta(T)$, all located in the first p slices.

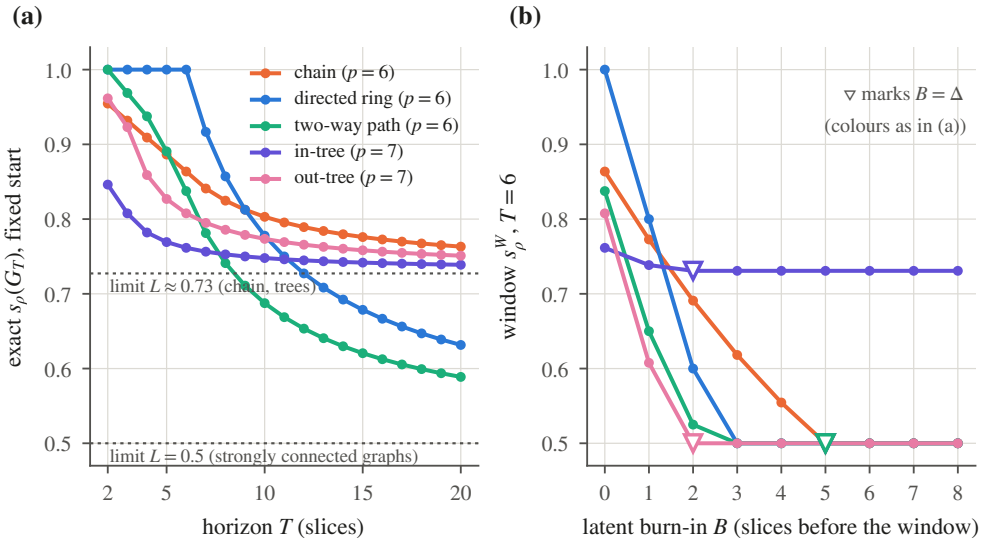


Fig. 3 Exact rel-sortability of unrolled DAGs (computed by enumerating ancestors and descendants of G_T, not from the theorems, which they confirm). (a) $s_\rho(G_T)$ under the fixed start against the horizon T. All curves tend to the limit of Theorem 2 (dotted: 0.73 for the chain and the trees, 0.5 for strongly connected graphs). The strongly connected ring has $s_\rho = 1$ for all $T \le p$ and the two-way path starts at 1 and decays slowly, so strongly connected summary graphs are not protected at short horizons. (b) The value of the last $T = 6$ slices after a latent burn-in of B slices: it falls to the stationary value of Theorem 5 at or before $B = \Delta$ (triangles).

Figure 4 shows a chain of three units under both starts. A sample-based relatives count on a long record cannot see such a difference, which motivates the stationary version.

Theorem 5 (stationary start) *Observe the window of the last $T \ge 2$ slices of G_{B+T} with a latent burn-in $B \ge \Delta$ (or of the infinite-past graph on $V \times \mathbb{Z}$), with relatives computed in the full graph. Then every window node has $\rho_W(v(t)) = T|\mathrm{rel}_S(v)|$, and for every T*

$$s_\rho^W = \tfrac{1}{2} + \frac{|E_\uparrow|}{2|E_S|}, \quad E_\uparrow = \{(i,j) \in E_S : \mathrm{Src}(i) \ne \mathrm{Src}(j)\}.$$

Lean: rhoBurn_eq_maxFinDist, sBurn_eq_maxFinDist, rhoZ_eq, sZ_eq, relSet_eq_iff_src; sharpness: edgeLoops_burn_sharp.

Sketch. After Δ slices every finite distance is absorbed, so $w(s) \in \mathrm{rel}(v(t))$ iff $\mathrm{anc}_S(v) \cap \mathrm{anc}_S(w) \ne \emptyset$. For an edge (i,j), $\mathrm{rel}_S(i) \subseteq \mathrm{rel}_S(j)$ with equality iff $\mathrm{Src}(i) = \mathrm{Src}(j)$. This burn-in is sufficient. For one edge with self-loops ($\Delta = 1$), $B = 0$ leaves $s_\rho^W \ge 2/3$ for every T (Lean), and for the directed chain with self-loops the window relatives at $B = \Delta - 1$ differ from the stationary ones at every length (Proposition 19); for other graphs a shorter burn-in can suffice (Remark 20).

What is and is not claimed about sharpness. Theorem 5 is a sufficient condition, valid for every summary graph with self-loops. That $\Delta - 1$ does not suffice is proved, and checked in Lean, for the one-edge graph ($\Delta = 1$); it is proved on paper

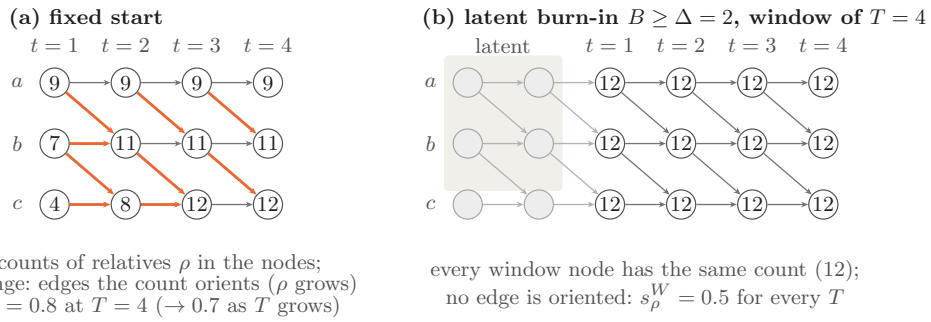


Fig. 4 The start artefact on the summary chain $a \to b \to c$. (a) Under the fixed start the first slice is independent noise, so the relatives counts ρ differ between a node and its child; every orange edge is oriented by the count. (b) With a latent burn-in $B \geq \Delta$ (grey, not observed) all window nodes of a unit have the same count and the stationary value is $s_\rho^W = 1/2$, because a chain has a single source component (Corollary 6). Relatives counts are computed in the full graph and counted inside the window.

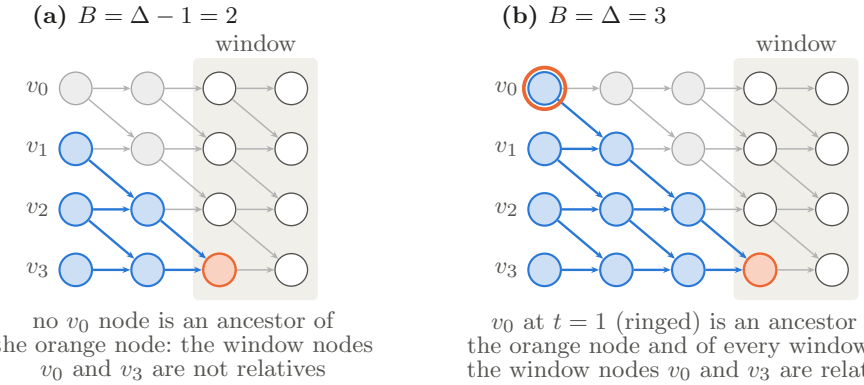


Fig. 5 Why $B = \Delta - 1$ is not enough for a chain (picture of the proof of Proposition 19; chain of four units, $\Delta = 3$, window of two slices). The orange node is the last unit at the first window slice; blue nodes are its ancestors. (a) With $B = 2$ latent slices no copy of v_0 is an ancestor, so the first-slice node of v_0 and the orange node are not relatives, whereas they are under the stationary start. (b) With $B = 3$ the copy of v_0 at slice 1 (ringed) is a common ancestor.

for the relatives of the directed chain at every length (Proposition 19, not checked in Lean), and exact computation gives a window s_ρ different from the stationary value at $B = \Delta - 1$ for chains with $p \leq 8$ and $T \leq 5$. For directed rings the minimal burn-in found by computation is $\lceil \Delta/2 \rceil$. For other graphs the bound is not sharp: in 726 random summary graphs $B = \Delta$ gave the stationary value in all of them and $B = \Delta - 1$ failed in 445, so it sufficed in 281; in a second exact computation on 396 random graphs (792 cases) none needed more than Δ, the minimal burn-in equalled Δ in 51% of the cases and was shorter by one to four slices in the rest. We have no closed form for the minimal burn-in of a given graph.

Scope: unit lags. Theorems 1, 2 and 5 are stated, and machine-checked, for VAR(1)-shaped unrollings. If each summary edge (i, j) instead acts at a lag $\tau_{ij} \in$

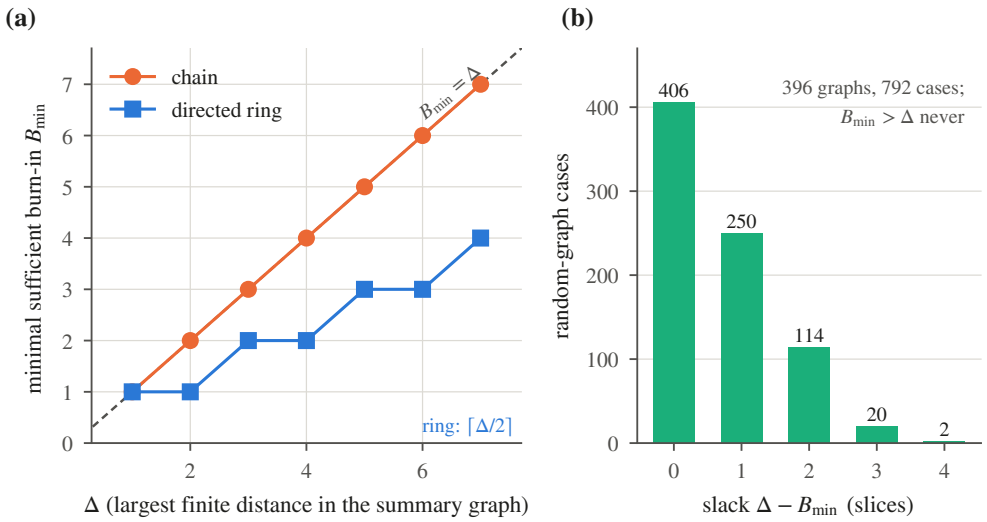


Fig. 6 The minimal sufficient burn-in $B_{\min}$ (smallest B after which the window relatives counts equal their stationary values) from exact computation (`b1y/q6_burnin_exact.py`, independent of the theorems; chains and rings $T = 2, \ldots, 5$, random graphs $T \in \{3, 5\}$). (a) Directed chains need $B_{\min} = \Delta$ for every Δ computed $(1, \ldots, 7)$; directed rings need only $\lceil \Delta/2 \rceil$. (b) For 396 random summary graphs (792 cases) $B_{\min}$ never exceeded Δ; it equalled Δ in 51% of the cases and was shorter by one to four slices in the rest.

$\{1, \ldots, \tau_{\max}\}$ (self-loops at lag one), the same proofs go through with d replaced by the lag-weighted distance d_τ: the copy $(i(t), j(t+\tau_{ij}))$ is strict iff some u has $d_\tau(u, j) - \tau_{ij} + 1 \le t \le d_\tau(u, i)$, the limit is unchanged, and the stationary result needs a burn-in $B \ge \Delta\tau_{\max}$ (paper proof, Appendix O; not machine-checked). Contemporaneous (lag-0) edges are not covered.

Corollary 6 (stationary condition; graph classes) $s_\rho^W = 1/2$ *iff every weakly connected component of* G_S *has exactly one source SCC. Strongly connected graphs, out-forests, chains and forks give* $1/2$. *In an in-forest (a river network),* $s_\rho^W = \frac{1}{2} + \#\{\text{edges into confluences}\}/(2|E_S|)$, *which exceeds* $1/2$ *iff there is a confluence.*

Lean: sBurn_eq_half_iff_maxFinDist, sZ_eq_half_iff, stat_half_of_outForest, statValue_inForest, statValue_inForest_gt_iff, collider_stat, fork_stat.

Table 2 shows the two versions for 3-node graphs: chain and fork have limit 0.7 and stationary value 0.5, the collider 0.7 in both. In the finite generator every summary DAG with the same m/p has the same value, so a benchmark built this way rewards relatives-sorting on chains and forks, where the stationary process gives no signal; for long stationary records the stationary value is the relevant one. All statements of Section 3.2 were checked by exhaustive enumeration ($p \le 4$ and 10^4 random graphs with $p \le 7$; 46,724 cases, 0 failures; Appendix C).

Without self-loops (Appendix N). Every edge of G_T still scores $1/2$ or 1, independently of T, by an explicit walk-length criterion (machine-checked), and Reisach

et al. (2026, Thm. 4) still holds, but the closed form and the converse of Corollary 3 fail: for the single edge $0 \to 1$ without loops $s_\rho(G_T) = 1/2$ for every T (with loops the limit is $2/3$). The loop assumption is therefore necessary for the closed forms.

What the value means. Show an edge in random order with the two relatives counts; the rule "orient towards the larger count, flip a coin on ties" has accuracy exactly s_ρ, never errs on an edge it decides, and no pairwise rule on the two counts does better (a few lines from Reisach et al.'s Theorem 1, Appendix Q). So $s_\rho = 0.7$ caps such a method at 70% orientation accuracy. Rules that also see node degrees can exceed it.

Lean: accRule_sortRule, accRule_le_sRho, edgeAcc_sortRule_strict.

Remark 7 (random summary graphs) For the directed Erdős–Rényi summary graph $G(p, c/p)$ with a self-loop at every node, the limits of $s_\rho(G_T)$ (as $T \to \infty$) and of s_ρ^W (as $p \to \infty$) are explicit functions of the mean degree c (Proposition 45). With x the largest root of $x = 1 - e^{-cx}$, the finite limit is $L(c) = \frac{1}{2} + \frac{c(1-x^2)}{2(1+c)}$; $L(1) = 3/4$, and L peaks at 0.754 near $c = 1.07$ and then decays like e^{-c}. The stationary value peaks at 0.665 near $c = 1.09$. At $c = 1$ convergence in p is slow: $S = 0.65574 \pm 0.00126$ at $p = 2 \cdot 10^4$, 0.65784 ± 0.00043 at $2 \cdot 10^5$ and 0.65796 at 10^6, against the formula value 0.65803. Random summary graphs are therefore rel-sortable only near or below criticality. The statement is proved on paper (Appendix P.1) and has not been checked by a human probabilist; Lean covers only the deterministic witness lemmas. The statement, the figure and the numerics are in Appendix P.

3.3 With time labels the start identifies nothing more

The experiments of Section 5.1 hide the time labels. One may ask whether the fixed start is also a source of information for a method that has them, for example through the changing covariance of the transient. It is not.

Proposition 8 (no gain from the transient) *Let $X_t = \mathsf{B}X_t + AX_{t-1} + e_t$ for $t = 1, \dots, T$, with $I - \mathsf{B}$ invertible, e_t i.i.d. and independent of the start $X_0 \sim \nu$ (ν may be a point mass). Put $M = (I - \mathsf{B})^{-1}A$ and $\eta_t = (I - \mathsf{B})^{-1}e_t$. Then the law of $(X_0, \dots, X_T)$ depends on $(\mathsf{B}, A,$ law of $e)$ only through $(M,$ law of $\eta)$. Two structures with the same reduced form are therefore observationally equivalent for every start law, every T and every number of independent replicates. If $\Sigma_t = \mathrm{Cov}(X_t)$ is nonsingular for some $t \le T - 1$, then $M = \mathrm{Cov}(X_{t+1}, X_t)\Sigma_t^{-1}$ and $\Omega = \mathrm{Cov}(\eta_t) = \Sigma_{t+1} - M\Sigma_t M^\top$ are identified.*

Lean: *not machine-checked (paper proof).*

Proof $X_t = MX_{t-1} + \eta_t$ is a Markov chain whose transition kernel depends on $(M,$ law of $\eta)$ only and whose initial law is ν. The two formulas follow from $\mathrm{Cov}(X_{t+1}, X_t) = M\Sigma_t$ and $\Sigma_{t+1} = M\Sigma_t M^\top + \Omega$. □

The covariance trajectory Σ_t is itself a function of (Σ_0, M, Ω). Identification through heteroskedasticity (Rigobon, 2003) needs the *innovation* covariance to change between regimes; in a transient with fixed mechanisms only the state covariance changes. In the lag-one model of this paper ($\mathsf{B} = 0$) the transition matrix is identified

under either start as soon as one Σ_t is nonsingular, in agreement with the labelled reference of Section 5.1 (PCMCI+, SHD/$|E|$ = 0.02 under both generators). The fixed start is therefore an artefact for label-free methods only. When each replicate is observed at a single time, so that only the marginal laws are seen, identification of the reduced form from these marginals has been established for linear stochastic differential equations (Guan et al., 2026; Wang et al., 2023); this is consistent with Proposition 8 and is not our contribution.

3.4 Is the time order identifiable without labels?

Reisach et al. (2026) end by raising the question to what extent knowledge of the time order can and should be used for causal discovery on independent trajectories of short time-unrolled DAGs. A prior question is how much of the time order is hidden when the labels are removed. The statements below are population statements with paper proofs (Appendix W); none is machine-checked, and we do not propose a learning method.

Setting. The columns of the data are the pT cells (unit j, slice t) of $X_{t+1} = AX_t + \varepsilon_t$ ($A_{ji} \neq 0$ iff $(i,j) \in E_S$, ε_t with diagonal covariance D, the same A and D for all t) in an unknown order; p and T are known. The *column DAG* is G_T on the columns. The *fixed-start law* has a diagonal covariance of the first slice (not necessarily equal to D), the *free law* an arbitrary one, which covers a burn-in window and an exact stationary start. The height of a node is one plus the number of edges of a longest directed path ending in it. A unit is *loop-only* if its only parent in G_S is itself, and *isolated* if in addition it has no child other than itself; m_0 is the number of loop-only units. A *representation* of a covariance Σ is a parameter of the same model class (an assignment of columns to cells, A' with non-zero diagonal, diagonal D', a start covariance of the given law) that reproduces Σ; it is *sparsest* if no representation has fewer non-zero entries in A'. Markov equivalence is equality of skeleton and unshielded colliders (Spirtes et al., 2000).

Proposition 9 (slices are a graph invariant) *(a) In G_T, the height of $v(t)$ is t. Hence T, p and the partition of the columns into slices are determined by the column DAG. (b) The partition of the columns into units is not: for the directed ring with self-loops, relabelling the cell (i,t) as unit $i - t \pmod p$ gives the reversed ring with self-loops and the same column DAG. (c) Without self-loops (a) fails: for the single edge $0 \to 1$ the unrolled graph is a set of disjoint edges, and the slice of an edge is not determined.*

Lean: not machine-checked (paper proof; (a) uses the walk-length argument behind edge_score_dist).

Proposition 10 (Markov equivalence class of G_T) *An edge of G_T is reversible within its Markov equivalence class iff it is an edge $j(t) \to j(t+1)$ of the self-chain of a loop-only unit j. The class has exactly T^{m_0} members.*

Lean: not machine-checked (paper proof).

So under the fixed start a perfect label-free learner that uses conditional independence only is left with T^{m_0} graphs: $m_0 = 0$ for a ring, 1 for a chain or an out-tree, and the number of leaves for an in-tree.

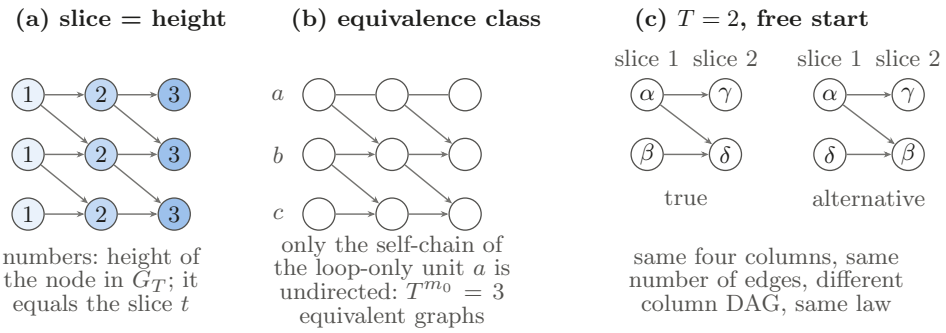


Fig. 7 What is identified when the labels are hidden. (a) The height of a node in G_T equals its slice, so the partition of the columns into slices is a graph invariant (Proposition 9). (b) The Markov equivalence class of G_T (Proposition 10): only the self-chain of the loop-only unit a is undirected, giving $T^{m_0} = 3$ graphs; the other edges are oriented by the colliders at b and c. (c) Without a restriction on the start ($T = 2$, free start law) the same four columns admit two labellings with the same number of edges and the same law (Proposition 12(a)): the time order is not identified.

Proposition 11 (fixed start, given faithfulness) *Let the columns be Gaussian with covariance Σ from the fixed-start law, $T \geq 2$, and let their distribution be faithful to G_T. (a) Every sparsest representation of Σ has the true column DAG, except that the self-chain of an isolated unit may be reversed as a whole. (b) For $T \geq 3$ the chain of an isolated unit j with coefficient $a = A_{jj}$, noise variance d and start variance v_1 can be reversed iff its start is stationary, $v_1(1 - a^2) = d$. For $T = 2$ it can always be reversed. (c) Hence for $T \geq 3$ the column DAG is identified from Σ iff no isolated unit has a stationary start.*

Lean: not machine-checked (paper proof; faithfulness is assumed, not shown to be generic).

Proposition 12 (three exact non-identifiability results) *(a) $T = 2$,* free law. *For the summary edge $a \to b$ with self-loops and a free start covariance, the labelling with slices $\{a_1, b_2\}, \{a_2, b_1\}$ is a representation with the same number of edges and a different column DAG ($a_1 \to a_2$, $a_1 \to b_1$, $b_2 \to b_1$). (b)* Isolated units, stationary start. *Reversing the chain of an isolated unit with a stationary start leaves Σ unchanged, for every T (Peters et al., 2009; G. Weiss, 1975); in particular $p = 1$ is never identified from a stationary window. (c)* Symmetric dynamics. *If $A = A^\top$, $D = \sigma^2 I$ and the start is stationary, reversing all slices leaves Σ unchanged for every T, and the reversed column DAG is a sparsest representation.*

Lean: not machine-checked (paper proof).

Conjecture 13 (burn-in window) *Under the free law with $T \geq 3$, $p \geq 2$, a weakly connected summary graph, faithfulness to the graph that adds a complete DAG on the first slice, and a time-reversed innovation covariance that is not diagonal, every sparsest representation has the true column DAG.*

We can reduce Conjecture 13 to a finite combinatorial statement (the skeleton admits only one partition into T layers of size p that is compatible with an unrolling) and have checked it by hand only for the summary edge $a \to b$ with $T = 3$ (Appendix W).

Reading. Under the fixed start, hiding the labels hides little in the population: the slices are a graph invariant, and the column DAG is identified up to isolated units. This is a third regularity of the time-unrolled construction, next to rel- and var-sortability, and unlike them it comes from the shape constraint (the graph is an unrolling), which a burn-in does not remove; whether it also survives a burn-in is Conjecture 13. Three cautions apply. The statements assume faithfulness, and the tied parameters of an unrolled model form a null set among all linear Gaussian models on G_T, so the usual genericity argument does not apply directly. For coloured Gaussian DAG models Boege et al. (2025) prove that a generic covariance matrix is faithful when only edges or only vertices are coloured, and give a mixed example that has no faithful distribution; which coloured DAGs admit one is open (Section 2.3). Our model ties both transition weights and, if the noise variances are equal across slices, noise variances, so on our reading it is the mixed case, and we have not established faithfulness for it. The statements concern the global optimum of a sparsity criterion, not a computable procedure. And they say nothing about finite samples: at $n = 1000$ FLOP is far from the labelled reference (Table 4). Dynamic Bayesian networks with a start network and a tied transition network are classical when the labels are known (Friedman et al., 1998).

3.5 Machine-checked proofs

The formalisation (Lean 4, de Moura & Ullrich, 2021, with Mathlib, The mathlib Community, 2020; 16 files, 4,159 lines, clean rebuild of the final state) certifies the exact and stationary results of Section 3.2 (Theorems 1, 2 and 5, Corollaries 3 and 6, Proposition 4), including the Δ forms of the bound and the burn-in and the sharpness example of Remark 20 (one edge, $\Delta = 1$), the loop-free exact edge value and the accuracy reading, the algebra of the real-data estimand $\tilde{s}$ (Appendix J), and the deterministic algebra of F1–F3. It does *not* certify the ER limits, the loop-free limit and periodicity, the lag-weighted extension, the statements of Sections 3.3 and 3.4, or any probability bound or limit theorem. The sources contain 259 theorem and lemma declarations and no `sorry`, `admit` or `axiom`. The axiom report, produced with `#print axioms`, lists 163 of them (33 headline theorems, 53 graph lemmas, 65 algebraic facts, many elementary, and 12 examples: concrete graphs, the one-edge sharpness example and one numerical instance), each depending only on `propext`, `Classical.choice` and `Quot.sound`; the other 97 theorem and lemma declarations are not in the report. Relatives, s_ρ, G_T and the windows are defined from scratch following Reisach et al. (2026). No human has yet checked that each Lean statement matches its paper statement, and *machine-checked* in this paper means a theorem proved in Lean for the statement as formalised. Appendix H maps every statement to its Lean names and status.

Takeaway (Q1, Q2). The footprint is a function of the summary graph alone. Its exact value, its limit and the burn-in after which it no longer depends on the start are settled for lag-one unrollings with self-loops, with the graph statements checked in Lean. Without labels the slices remain identifiable under a fixed start, conditional on faithfulness. The distributional side of the start (variances, correlations) is not covered by these results.

4 Finite Samples: How Much of the Footprint an Estimator Sees

Question (Q3). How much of the footprint can an estimator see? The exact values are population statements about a graph. Section 4.1 asks whether the relatives behind them can be estimated from one finite, autocorrelated record, and Section 4.2 asks what block means and subsamples of a faster process do to the three sortability criteria.

4.1 Estimating relatives from one record

We now ask whether the population relatives can be estimated. X is a zero-mean stationary Gaussian process with absolutely summable auto- and cross-covariances $\gamma_{ab}(h)$, standardised so that $\gamma_{aa}(0) = 1$, observed at n times. Two nodes are declared relatives if the Fisher-z test of the sample correlation $\hat{r}_{ab}$ rejects at level α with Bonferroni over $m = d(d-1)/2$ pairs, i.e. $|\hat{r}_{ab}| > \tanh(z_c/\sqrt{n'-3})$, $z_c = \Phi^{-1}(1-\alpha/2m)$; $\hat{\rho}(v) = 1 + \#\text{declared partners}$. The *nominal* test uses $n' = n$. Let $B_{ab} = \sum_h |\gamma_{aa}(h)\gamma_{bb}(h) + \gamma_{ab}(h)\gamma_{ab}(-h)|$ (for an independent pair $B_{ab} = \sum_h |\rho_a(h)\rho_b(h)|$, at least the Bartlett sum $W_{ab} = \sum_h \rho_a(h)\rho_b(h)$, with equality iff all products are nonnegative; e.g. 4.56 against 0.22 for AR(1) coefficients 0.8 and -0.8).

F1: recovery and effective sample size. That the nominal-n test has the wrong size under autocorrelation is classical: effective degrees of freedom go back to Bartlett (1935), with spatial and multivariate versions in Clifford et al. (1989) and Roy (1989), and the variance with exactly our summand B_{ab} is in Afyouni et al. (2019). We add a non-asymptotic exact-recovery bound for the relatives set (Theorem 28, Appendix F): if the thresholds separate the population correlations by more than explicit deviations, the declared set and hence $\hat{s} = s_\rho^{\text{pop}}$ are recovered with failure probability decaying as $\exp(-n\min(\tau, \rho_{\min} - \tau)^2/(2(\sqrt{B} + \rho_{\max}\sqrt{D_\gamma})^2))$, so n/B acts as an effective sample size. For i.i.d. snapshots $B = 1 + \rho^2$; under strong persistence B can be in the tens. This bound is for a fixed threshold; for the Bonferroni Fisher-z threshold $\tau_n \asymp z_c/\sqrt{n'}$ that we use, the failure probability it certifies tends to $d(d-1)e^{-z_c^2/(2B)} > 0$ (0.089 for i.i.d. data, $d = 5$, $\alpha = 0.01$; Proposition 32), and recovery is instead guaranteed asymptotically (Proposition 31).

Proposition 14 (size distortion of the nominal test) *In this setting, a non-relative pair is declared by the nominal test with probability* $\to 2\bar{\Phi}(z_c/\sqrt{W_{ab}})$, $W_{ab} = \sum_h \gamma_{aa}(h)\gamma_{bb}(h)$ *(*$\le B_{ab}$*, with equality iff all products are nonnegative), above the nominal* α/m *when* $W_{ab} > 1$. *With* $n_{\text{eff}} = n(1-\rho^2)^2/W_{ab}$ *(*W *the delta-method Bartlett variance) the false-positive rate tends to* α/m.

Lean: not machine-checked (uses the multivariate Breuer–Major theorem).

The joint asymptotic normality used here follows from the multivariate Breuer–Major theorem (Arcones, 1994; Nourdin et al., 2011) (Appendix F). On real data n_{eff} is estimated by $n/\hat{B}_{ab}$, $\hat{B}_{ab} = \max(1, 1 + 2\sum_{h=1}^{K} \hat{\rho}_a(h)\hat{\rho}_b(h))$. If $K = K_n \to \infty$ with $K_n = o(n)$, $\hat{B}_{ab}$ is consistent, the plug-in test has asymptotic size at most α/m per pair, and the relatives set is recovered with probability at least $1 - \alpha$ in the limit

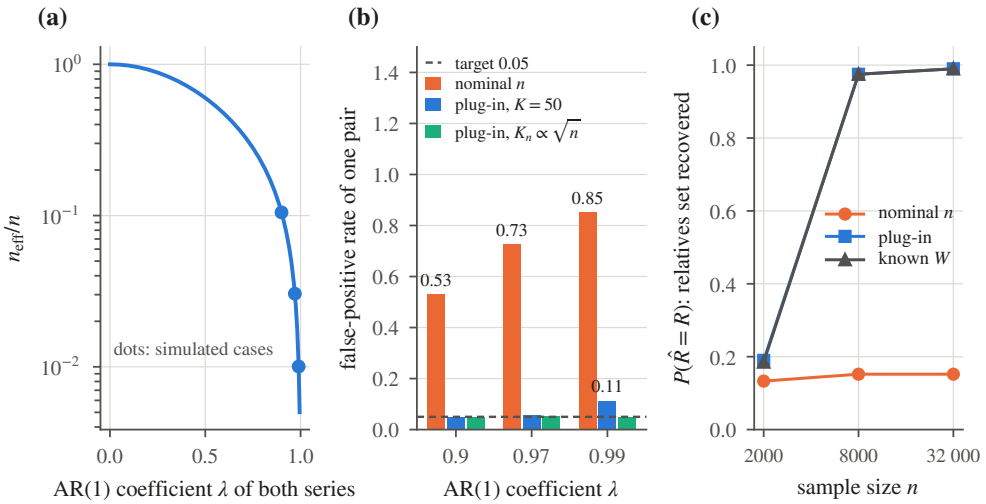


Fig. 8 Estimating relatives from one autocorrelated record (simulations of `finite2/sim_finite2.py`, fixed seeds). (a) The effective sample size $n_{\text{eff}} = n/W$ of the correlation of two independent AR(1) series with the same coefficient λ, with $W = (1+\lambda^2)/(1-\lambda^2)$; dots are the simulated cases. (b) False-positive rate of one non-relative pair at nominal level 0.05: the test that uses n rejects far too often, the plug-in test that estimates W is close to the target, except that a fixed truncation lag $K = 50$ is too short at $\lambda = 0.99$ (0.11). (c) Probability that the Bonferroni Fisher-z procedure recovers the relatives set of a five-node graph ($\lambda = 0.9$, $\alpha = 0.05$): the nominal test does not, the plug-in test does, and it behaves like the test that knows W.

(Proposition 31). Our code uses a fixed K (20–100), which controls size only if the truncated sum captures the Bartlett sum: at AR coefficient 0.99 and $K = 50$ the size is 0.113 against 0.05, while on the B3 data K up to 300 changes none of 2,730 pair decisions (Appendix F). We did not run this check on the river data, which are the most persistent, or on study 3 and E1. For linear processes with non-Gaussian i.i.d. innovations the null variance, and hence the size results, are unchanged, because independent series share no fourth cumulant; for relatives the variance changes, which affects power but not the limits (Remark 33).

F2: saturation by a latent driver. Let $X_v = Y_v + b_v C$ with C a stationary unit-variance process independent of the VAR Y, and $\ell_v = b_v^2/\mathrm{Var}X_v$ the driver's share of variance.

Proposition 15 (saturation) *For summary non-relatives, $\rho_{uv} = \mathrm{sgn}(b_u b_v)\sqrt{\ell_u \ell_v} =: \kappa_{uv}$. If all $b_v \neq 0$, every pair is a population relative and the population rel-sortability is exactly $1/2$; $P(\hat{s}_n \neq \frac{1}{2}) \leq (d^2+3d)e^{-cn}$ with $c \approx \kappa_{\min}^2/(18(\sqrt{B} + \sqrt{D_\gamma})^2)$. Asymptotically, a pair is detected with probability $> 1/2$ once $n > n_{1/2} = 3 + (z_c/\operatorname{atanh}\kappa)^2 \approx z_c^2/\ell^2$ (nominal test; for the effective test, the same bound on n_{eff}).*

Lean: prop21_corr_std, sOf_saturated, abs_sOf_sub_half_le, location of $n_{1/2}$: ncrit_eq_iff; the probability statements are not machine-checked.

$n_{1/2}$ does not depend on autocorrelation, which only sets the width of the transition. A driver with share $\ell = 0.1$ saturates at $n \approx 10^3$, and one with $\ell = 0.64$ at $n \approx 22$ (both for $z_c = 3.29$). A relatives count on a long record can recover the summary relatives only while $\max_{\text{non-rel}} |\kappa| < \tau_n < \min_{\text{rel}} |\rho|$, a window that closes at n of order $n_{1/2}$.

F3: conditioning on a proxy. A natural remedy is to condition on a measured proxy Z of the driver. It works only if the proxy is almost perfectly reliable.

Proposition 16 (conditioning) *Let $Z = \sqrt{r}\, C + \sqrt{1-r}\, e$ with reliability $r = \operatorname{corr}(Z, C)^2$, π_{uv} the partial correlation given Z, and ρ^Y_{uv} the driver-free correlation. Then*

$$\pi_{uv} = \frac{\rho^Y_{uv}\sqrt{(1-\ell_u)(1-\ell_v)} + (1-r)\operatorname{sgn}(b_u b_v)\sqrt{\ell_u \ell_v}}{\sqrt{(1-r\ell_u)(1-r\ell_v)}}.$$

At $r = 1$ the conditional relatives are the summary relatives and Theorem 5 is restored. For non-relatives with $\ell_u, \ell_v \in (0,1)$ and $b_u b_v \neq 0$, $|\pi_{uv}|$ is strictly decreasing in r and vanishes only at $r = 1$; fewer than half of the non-relative pairs are declared only if $r > r_{\text{crit}} = 1 - \frac{\tau'(1-\ell)}{\ell(1-\tau')}$, $\tau' = \tanh(z_c/\sqrt{n'-4})$ (equal shares). If the driver reaches nodes through different filters, a scalar proxy leaves a nonzero residual even at $r = 1$. The reliability is identified by $\rho_{uZ}\rho_{vZ}/\rho_{uv} = r$, the identity behind the three-indicator identification rule for one-factor models (Bollen, 1989, Ch. 7).

Lean: prop31_std, cor32, absPi_strictAntiOn, absPi_eq_zero_iff, rcrit_iff, thm35_bracket, pc_lipschitz, prop36_ar1_ne_zero, prop36_rank_one, prop38_ratio.

So spurious conditional relatives appear for every $r < 1$ once n is large; $r_{\text{crit}} = 0.95$ for $\ell = 0.64$ and $\tau' = 0.082$. Formalisation found four missing nondegeneracy hypotheses in our earlier statements (e.g. $b_u b_v \neq 0$; Appendix H). The algebra of F2 and F3 is machine-checked; the probability statements are not. The non-asymptotic bounds are loose (factors 3–13 in n, Appendix G); the asymptotic statements matched every simulation to within 0.03 and are the ones used in Section 5.3.

Detectability as the second factor. The statements above separate two things that the exact values of Section 3 run together. The graph decides which pairs are relatives; the sample size, the autocorrelation and any latent driver decide which relatives a test can detect. An estimator sees the footprint only through the second. In the benchmark experiment of Section 5.1 this shows up as follows: the exact gap between the two starts orders the simulated conditions, and a post hoc calculation of which correlations are detectable at $n = 1000$ reproduces the size of the observed gain. That calculation is the analytic expectation of the algorithm, so it checks consistency of size and does not separate the graph mechanism from a general change in correlation strength. A designed test does (Section 5.2): the share of the exact gap that is realised rises from 0.18 at $n = 200$ to 0.76 at $n = 25{,}000$, and from 0.09 to 0.51 when the spectral-radius cap rises from 0.5 to 0.95; and a start that has the population relatives of a burn-in but the detectable correlations of the fixed start keeps 0.96 of the gain.

4.2 Block means and subsamples

Real records are often block means or subsamples of a faster process. For a stable linear SVAR with independent innovations, nodes with disjoint summary ancestor sets are

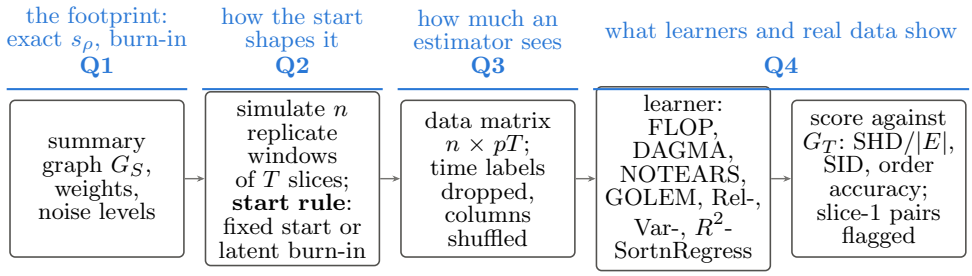


Fig. 9 The experimental pipeline and where the four questions enter. A summary graph is simulated under a start rule (fixed start or latent burn-in), the time labels are dropped and the columns shuffled, a learner returns a graph, and the result is scored against the window DAG G_T. Edges that leave the first slice are scored separately in the controlled tests, because under a burn-in their confounder lies in the unobserved past.

independent at all lags, and so are any per-node, time-invariant transforms of them, block means and subsamples included; otherwise the lag-0 covariance is nonzero for almost every parameter (Proposition 22). So the population relatives of contemporaneous values are rel_S for every observation scheme, and the population relatives sort takes the stationary value of Theorem 5 (on non-self-loop edges), not the finite-G_T value. Var- and R^2-sortability behave differently (Theorem 23): block means converge to the long-run structural model, whereas subsamples keep the snapshot covariance, so for large blocks block-mean R^2-sortability is 0 on out-stars and 1 on in-stars with at least two arms. Subsampling also changes which structures are identifiable (Gong et al., 2015; Hyttinen et al., 2016); we ask only about the three criteria. The statements, the numerical observations and the proofs are in Appendices D and E.

Takeaway (Q3). Relatives can be estimated from one record only while the effective sample size is above the saturation threshold of a latent driver and the true relatives are detectable at lag zero; under those conditions the population statement is recovered with probability of at least $1 - \alpha$ in the limit (Proposition 31), and outside them the count carries little or no information. The bounds are loose, the proofs are on paper, and block means and subsamples change var- and R^2-sortability while leaving the population relatives unchanged. In the benchmark simulations, the effect of the start on a relatives sort is the detectable image of the exact gap (Section 5.2).

5 Experiments and Real-Data Evidence

Question (Q2 empirically, and Q4). What does the start of a simulation do to what label-free learners appear to achieve, and does a relatives count orient edges in real records? Section 5.1 measures what the start does to benchmark results when the time labels are hidden from static learners. Section 5.2 reports controlled tests of two readings that Section 5.1 leaves open: what carries the effect of the start, and whether the learner reversal survives tuning and a variance control. Section 5.3 applies the finite-sample theory and the criterion to real records, which is a different question: whether relatives can be estimated from one record at all.

5.1 The start artefact in simulated benchmarks

Design (B1). We unrolled 420 summary graphs (30 per family and $p \in \{5, 10\}$; families in Table 3; details in Appendix R) over $T \in \{3, 5, 10\}$ and drew $n = 1000$ linear-Gaussian trajectories (positive weights, heterogeneous noise) either with a fixed start (G1: the first slice is pure noise, which is what unrolling from $t = 1$ gives; the preprint of Reisach et al. 2026 fixes no initial law) or after a latent burn-in of 100 steps (G2; an exact stationary start reproduced G2 to within 0.013). Methods saw the pT columns shuffled, with time labels hidden: rel-, R^2- and Var-SortnRegress (Reisach et al., 2026, 2023), FLOP (Wienöbst et al., 2026), PC (Spirtes et al., 2000) in causal-learn (Y. Zheng et al., 2024) and a random order; SID by gadjid (Henckel, Würtzen, & Weichwald, 2024). The analysis plan was committed before any algorithm was run (Appendix R). Weights were drawn uniformly from $[0.5, 1]$ on every summary edge and rescaled to a spectral radius of at most 0.8; the realised radii were not tabulated, and persistent dynamics near a unit root are not covered. CIs below resample whole summary graphs.

Rel-sortability: what the exact values do and do not explain. The fixed start raised rel-SortnRegress order accuracy in all five non-strongly-connected families (pooled over these five $+0.125$, 95% CI $[0.116, 0.134]$; Holm $p \leq 1.1 \cdot 10^{-10}$ per family; $+0.133$ over all seven). Because this order accuracy is essentially the estimated rel-sortability, the hypothesis tests whether estimated counts follow the exact values, not algorithm impact. Agreement in sign is not a discriminating test: all 42 (family, p, T) cells have a positive exact gap $s_\rho(G_T) - s_\rho^W$ and a positive observed gain. What the exact gap adds is a ranking and a rough scale: Spearman 0.80 across the 42 cells (graph-cluster permutation $p \leq 10^{-4}$; 0.62 over graph–horizon units) and slope 0.53, so the observed gain is about half the exact gap. A post hoc calculation (Figure 10, Appendix T) reproduces the size: the expected outcome of the same thresholded-correlation sort under the exact start-dependent window covariances, with the probability that each correlation is detected at $n = 1000$, matches the gain with cell $R^2 = 0.97$ and slope 1.03. This is the analytic expectation of the algorithm, so it checks that the simulated gain has the size its covariances imply. It does not separate the structural mechanism of Section 3.2 (pairs that are relatives under one start only) from a general change of correlation strength during the transient. In a joint regression the exact gap adds nothing to it (coefficient -0.063 over cells and -0.039 over graph–horizon units). For rings with $p = T = 10$, for example, the gap is 0.50 but only 40% of the relatives have a detectable correlation; the calculation gives $+0.094$ and we observed $+0.079$. The designed test that separates these readings is reported in Section 5.2. The gain does not come from the latently confounded first window slice: without edges leaving it, it is $+0.089$ $[0.079, 0.100]$. Strong connectivity did not protect short horizons: on rings and symmetrised ER graphs the gain was $+0.13$ to $+0.31$ in all six cells with $T < p$ and $+0.08$ to $+0.22$ in four of the six cells with $T \geq p$, consistent with the exact values (rings: $s_\rho(G_T) = 1$ for $T \leq p$). A separately planned replication (B1-NL, Appendix S; new seeds) with nonlinear additive dynamics, signed weights, and signed weights with uniform noise reproduced the gain in all 126 cells (pooled $+0.108$, $+0.137$, $+0.137$; 5/5 families significant in each; Spearman 0.68–0.82, slopes 0.34–0.56).

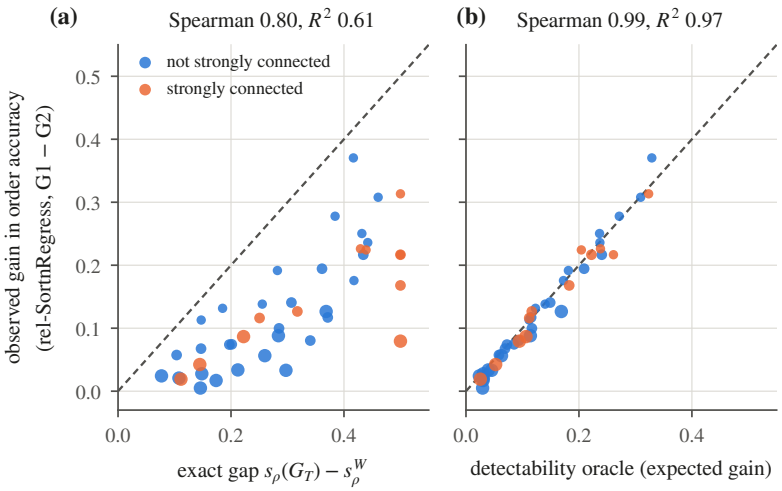


Fig. 10 B1: observed G1−G2 gain in rel-SortnRegress order accuracy per (family, p, T) cell (marker size grows with T; squares: strongly connected) against the exact population gap (left) and against the post hoc expectation of the algorithm under the exact window covariances (right). Dotted: identity.

Sort-based baselines. The start changed the size of most method differences (Holm-significant G1×G2 interactions in 294 of 595 pair×cell SHD tests) and flipped the sign of 11 (Holm over 1,190 tests), all among sort-based baselines: e.g. on chains rel-SortnRegress beat R^2-SortnRegress under G1 and lost under G2. These baselines are no better than the empty graph on SHD under G2 (SHD/$|E|$ = 1.39–2.04; FLOP 0.46 and PC 0.64 never exceed 1), so SID is the more informative metric for them. In B1-NL the swap hypothesis was supported by the plan's rule under nonlinear dynamics, but only through one swap between two weak baselines, and not under signed weights; we discount it. The interactions themselves replicate (in-tree Var-SortnRegress vs FLOP on SHD, SID and orientation in all three parametrisations; chain rel- vs R^2-SortnRegress on SID and orientation, and on SHD in the short-horizon cells); sign reversals on SHD did not recur because the B1-NL level differences are larger.

Competitive learners: the second artefact. The fixed start also makes the window var-sortable, because variance accumulates from the noise-only first slice (Var-SortnRegress +0.20 raw, +0.002 standardised; rel-SortnRegress is scale invariant). A third planned experiment (Appendix T; 384 new data sets, families ring, chain, in-tree, ER(c=2), $T \in \{3, 5\}$) ran DAGMA (Bello et al., 2022), NOTEARS (X. Zheng et al., 2018) and GOLEM (Ng et al., 2020) with default settings on raw and standardised data (Table 4). On raw data all four variants improved sharply under G1 (order accuracy +0.16 to +0.23, SHD/$|E|$ −0.41 to −0.55; Holm $\leq 4 \cdot 10^{-8}$). This changed their ranking against FLOP: on in-trees DAGMA and NOTEARS beat FLOP under G1 and lost under G2 (both reversals Holm-significant and repeated on SID and order accuracy; DAGMA's also with first-slice pairs excluded; GOLEM only drew level with FLOP under G1), and a continuous method had the lowest mean SHD/$|E|$ among label-free methods in 3 of 8 (family, p) cells under G1 (one of them by 0.09 against 0.10, untested) against 0 of 8 under G2. The plan had predicted no rank

change against FLOP; that prediction failed. It had predicted that standardisation would reduce the gain without removing it. By the plan's rule this held, except that DAGMA's SHD gain was removed; but the order-accuracy effect of the start became negative for DAGMA and GOLEM (-0.06 to -0.10), which the plan had not expected, and NOTEARS kept a small significant gain ($+0.036$; SHD/$|E|$ -0.077). At the time of this experiment we could only suspect var-sortability: standardised DAGMA and GOLEM are below chance in order accuracy under both starts (0.31–0.42), so they cannot show whether a scale-free part remains; standardised NOTEARS shows that some does; the fixed start also raises R^2-sortability, which is scale free; and the var-sortability gain is the same on rings and chains while the learners' gains are not. The variance-matched control and the tuned re-run were planned after this experiment and are reported in Section 5.2. In any case Theorems 2 and 5 do not predict these ranking changes; they cover the scale-free gain of the sort baselines (rel-SortnRegress $+0.19$ again here). PCMCI+ (Runge, 2020), which uses the time labels, recovers the window under both generators (SHD/$|E| = 0.02$). Holm across the 21 claim-level tests of B1, B1-NL, B3 and this experiment keeps every claim except the B1-NL swap claim under signed weights ($p = 0.108$). The claim-level p-values were computed after the fact, the sequence of studies was adaptive, and the correction does not cover the real-data tests of Section 5.3.

Which generators are exposed. Two conditions are needed. First, the slices must be given to a method that does not use time order; a single record given to a time-series method is outside this threat model. Second, the start must matter: for the graph artefact of Section 3.2 this means a burn-in shorter than $\Delta\tau_{\max}$ steps, and for the variance artefact it means a start that has not mixed in distribution, which depends on the spectral radius and which we did not compute for any deployed generator. We read the source code of common time-series generators (Appendix T). The time-unrolled construction of Reisach et al. (2026) meets both conditions: it starts at $t = 1$ with no burn-in and gives every slice to the method as a separate variable (its data code is not public). The others produce one record for time-series methods, so they fail the first condition as shipped. On the second: an exact stationary start (Herman et al., 2025) removes both artefacts; burn-ins of 100 to 1000 steps (causalnex, the Neural-GC simulators (Tank et al., 2021)) remove the graph artefact and the variance artefact only if the process mixes within them; tigramite's `structural_causal_process` (Runge, 2026) drops $\lfloor 0.2T \rfloor$ steps, its older `var_process` none, the CDML generator (Lawrence et al., 2021) $2\tau_{\max}$, and CausalTime's demo rolls out 20 steps (Cheng et al., 2024), so windows cut from their first steps would be exposed. We could not check how the CauseMe data were initialised.

5.2 Controlled tests of the two-factor reading and of the learner reversal

Question. Which part of the start artefact is carried by the population relatives pattern and which by the correlations that a test can detect, and does the learner reversal survive tuning and a variance control? **Plan.** A time-stamped plan (Appendix U) fixed designs, thresholds, seeds, a table that maps outcomes to the paper's reading, and a multiplicity family before any data were simulated. Its two addenda (the second

restricts NOTEARS to $d \leq 30$ after fit times were measured) were committed before the scored runs they govern. In this subsection all metrics drop the edges that leave the first slice (SHD_{x1}, OA_{x1}): under a latent burn-in the first window slice is confounded by the last burn-in slice and its pairs are scored differently (reference SHD 0.044 under the fixed start and 0.273 after a burn-in, but 0.041 against 0.048 without first-slice pairs). Units are summary graphs, and intervals are graph-cluster bootstrap intervals.

Matched-correlation counterfactual. For the same graphs, weights and noise (300 graphs that are not strongly connected, $T \in \{3, 5, 10\}$) only the covariance of the first slice was changed. The counterfactual CF1 keeps the population relatives of the fixed start and rescales the first slice by a factor $c \in \{0.25, \ldots, 4\}$. The counterfactual CF2 has the population relatives of the burn-in, hence the same exact s_ρ, but in its smallest setting every new correlation is below half the detection threshold, so that the pairs a test can detect are those of the fixed start; larger settings make the burn-in pattern detectable. A second independent draw of the fixed start is the negative control. The control was valid (-0.004 $[-0.011, +0.002]$; limit ± 0.01), and the burn-in was mixed (exact stationary start within 0.001). The fixed start raised the order accuracy of rel-SortnRegress by $+0.090$ $[0.080, 0.101]$. CF2 with undetectable burn-in correlations kept 0.955 $[0.884, 1.030]$ of this gain; where the burn-in pattern became detectable the retention fell to 0.92 and 0.79. Changing the strength at the fixed start's own pattern moved the gain by 0.008 $[0.001, 0.016]$ between the weakest and the strongest setting, which the plan's rule (a quarter of the gain) calls flat. By the plan's rule the first result reads "detectable strength, not population pattern": the order accuracy of the relatives sort responds to which pairs clear the detection threshold, not to the population relatives pattern itself.

Dose. On rings, chains and ER(c=2) graphs the share $Q(n)$ of the exact gap that is realised was 0.18, 0.40, 0.57 and 0.76 at n = 200, 1000, 5000 and 25,000 ($+0.36$ $[0.29, 0.43]$ between 1000 and 25,000), and 0.09 and 0.51 at spectral-radius caps 0.5 and 0.95 (difference $+0.42$ $[0.35, 0.49]$). In the main design (graphs that are not strongly connected, $n = 1000$) the within-family slope of the gain on the exact gap was 0.57 $[0.40, 0.73]$, against 0.53 pooled in B1. Read together: the exact gap is the large-sample ceiling of what the start does to the relatives sort, and the effect realised at a finite n is its detectable image, which grows with n and with the correlation strength that the spectral radius sets. The share had not reached one at n = 25,000.

Learner reversal. We used chains, in-trees, rings and ER(c=2) graphs in six (p, T) cells with d from 15 to 60 and 12 graphs per (family, p). DAGMA and NOTEARS hyper-parameters were chosen on separate validation seeds, pooled over both starts so that tuning could not favour the artefact; GOLEM-EV and GOLEM-NV ran at their defaults ($d \leq 50$), and NOTEARS only for $d \leq 30$ because fits on burn-in data were too slow in larger cells. Two further arms rescaled each data set so that its marginal variances equal those of the other start, leaving correlations unchanged; the manipulation check passed (order accuracy of Var-SortnRegress within 0.01 of the target start). On in-trees, tuned DAGMA, NOTEARS and GOLEM-NV beat FLOP under the fixed start and lost to it under the burn-in (SHD_{x1} of FLOP minus method: $+0.200$ and -0.365 for DAGMA, $+0.234$ and -0.393 for NOTEARS, $+0.123$ and -0.325 for GOLEM-NV; all intervals exclude 0). GOLEM-EV only tied FLOP under

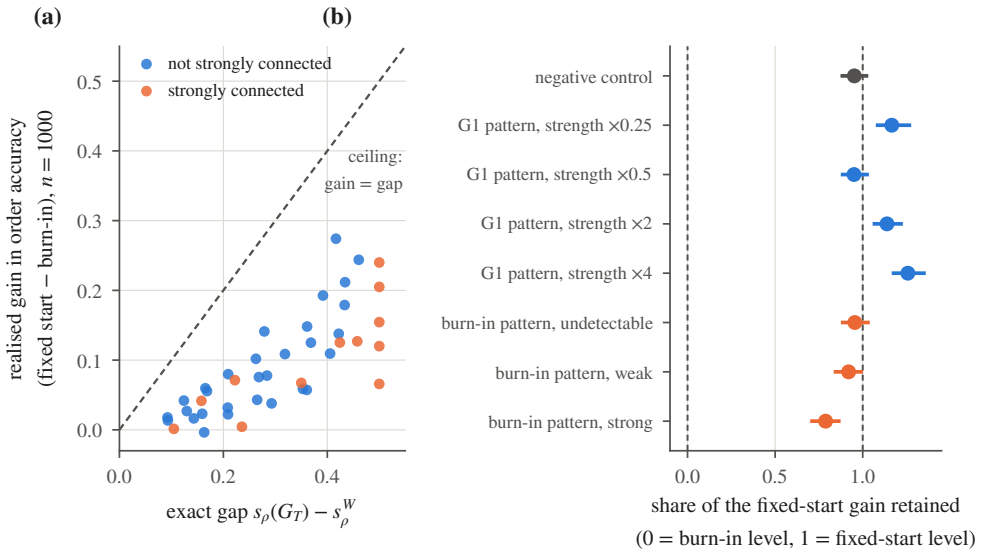


Fig. 11 Matched-correlation counterfactual (`b1y/results/q1.csv.gz`; graphs that are not strongly connected unless marked; order accuracy without first-slice edges). (a) Each point is one (family, p, T) cell: the realised gain of the fixed start over a burn-in, $n = 1000$, against the exact gap of Theorems 2 and 5. All cells lie below the line gain = gap. (b) Share of the fixed-start gain retained when only the first-slice covariance is changed (graph-cluster 95% intervals). Data with the fixed start's population pattern keep the gain at every strength. Data with the burn-in's population pattern keep it (0.96) while the new correlations are below the detection threshold and keep less of it (0.92, 0.79) as they become detectable.

the fixed start, so the sign pattern did not hold for it. FLOP's own SHD_{x1} rose by 0.099 $[0.080, 0.119]$ on in-trees under the burn-in. By the pre-specified rule the channel is variance for all four methods: giving fixed-start data the burn-in variances removed the advantage and made the methods worse than under the burn-in itself (the ratio is −1.25 for DAGMA, −0.44 NOTEARS, −0.41 GOLEM-NV and −1.50 GOLEM-EV, where 1 means the advantage is untouched and 0 that it falls to its burn-in level), an over-correction that we do not interpret. Imposing the fixed-start variances on burn-in data restored only 0.43 (DAGMA) and 0.49 (NOTEARS) of the advantage, below the 0.60 that the plan required for variance alone to be sufficient. So, in the plan's sense, variance is involved in the reversal and is not sufficient for it. Rings and ER(c=2) graphs showed no reversal, and on chains the continuous methods lost to FLOP under both starts. The seven claims of the plan's multiplicity family (the strength dose, the two dose tests, the slope, the DAGMA and NOTEARS preconditions and FLOP's shift) all reject under Holm at 0.05 (adjusted p from 0.0007 to 0.016; the weakest is the strength dose, whose effect-size verdict is "flat"). This family is separate from the 21-claim family above, and the sequence of studies is adaptive.

First-slice-excluded scoring. Dropping the first-slice pairs removes nearly a third of the rel-SortnRegress order-accuracy gain (+0.125 to +0.089) and about half of FLOP's SHD shift (+0.197 to +0.093) (Table 5); the artefact survives, and the rank changes against FLOP on in-trees remain and grow slightly. For B1 and the learner

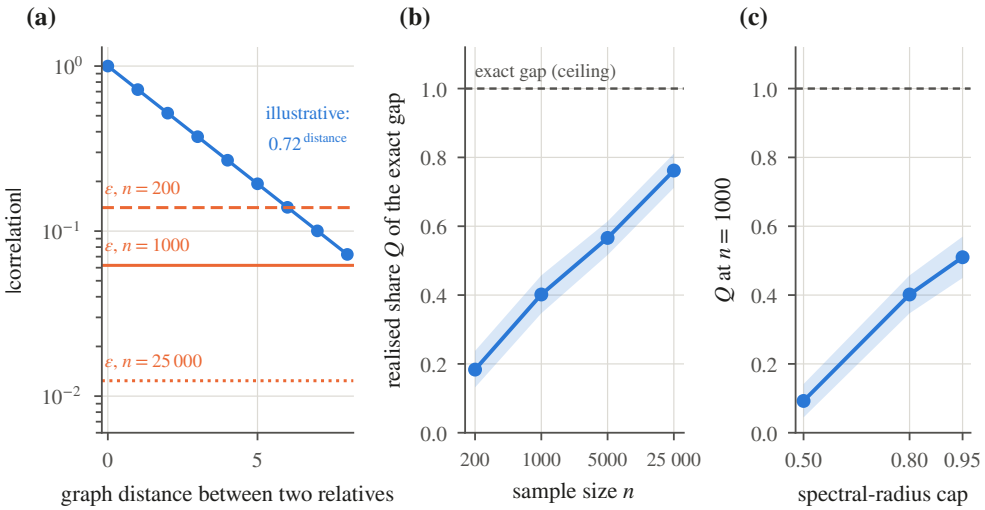


Fig. 12 Two factors behind the effect of the start. (a) Illustration (not data): the correlation between two relatives decays with their graph distance, and a pair is detected only if it clears the threshold $\varepsilon(n)$, which falls as n grows, so more of the exact footprint is visible at large n. (b) The share Q of the exact gap that is realised in the order accuracy of the relatives sort, against n (rings, chains, ER(c=2); graph-cluster 95% bands). (c) The same share at $n = 1000$ against the spectral-radius cap of the weights: stronger dynamics make more relatives detectable. Q stays below one at $n = 25\,000$.

extension the full-window numbers stay the primary results of their own plans, and the excluded-slice numbers are shown beside them.

Further baselines. On 384 new data sets (rings, chains, in-trees and ER(c=2); $p \in \{5, 10\}$; $T \in \{3, 5\}$) FLOP had the lowest mean SHD_{x1} among the label-free methods (FLOP, GES from Tetrad, GRaSP and BOSS in causal-learn, DirectLiNGAM, rel- and Var-SortnRegress) in 7 of 8 (family, start) groups; the exception is in-trees under the fixed start, where Var-SortnRegress (0.136) beat FLOP (0.251), which fits the variance account. Mean $\text{SHD}_{x1}/|E|$ under the fixed start and after a burn-in was 0.161 and 0.353 for FLOP, 0.358 and 0.449 for BOSS, 0.344 and 0.522 for GRaSP, 0.596 and 0.944 for GES, 1.417 and 1.688 for DirectLiNGAM, 0.699 and 1.392 for rel-SortnRegress and 0.487 and 1.139 for Var-SortnRegress; the empty graph scores 1. Methods that use the time labels are much closer to the truth and do not depend on the start: PCMCI+ (Runge, 2020) 0.021 and 0.022, DYNOTEARS (Pamfil et al., 2020) 0.025 and 0.047, pooled lag-1 least squares 0.000 and 0.071, and a pooled variant of VAR-LiNGAM (Hyvärinen et al., 2010) for replicated short trajectories 0.007 and 0.080 (our noise is Gaussian, so its identifiability condition does not hold). With the labels the gap to FLOP is 0.14 under the fixed start and 0.33 after a burn-in. Every label-free method got worse under the burn-in; PCMCI+ did not change.

Takeaway (Q2, Q4). The start changes what label-free learners appear to do in two ways that the theory only partly covers. For the relatives sort, the exact gap is the large-sample ceiling and the effect at a given sample size is the detectable part of it. For tuned continuous-optimisation learners, a fixed start produces a ranking change

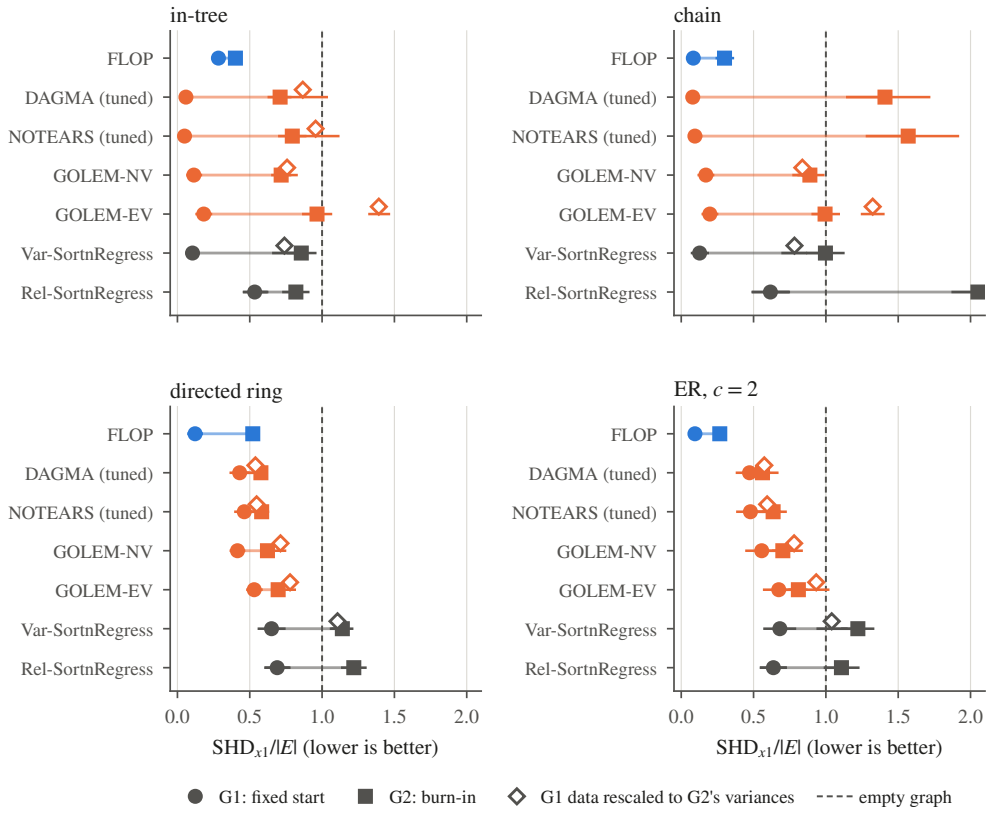


Fig. 13 Learners under the fixed start (G1, circles) and a burn-in (G2, squares) on four summary-graph families (SHD per true edge without first-slice edges; lower is better; 95% graph-cluster intervals; dashed: the empty graph). On in-trees the tuned continuous-optimisation methods (orange) beat FLOP (blue) under the fixed start and lose to it under the burn-in. Diamonds: the fixed-start data rescaled to the marginal variances of the burn-in, which removes the advantage and over-corrects (see text). The panels use only the cells every method has $(p, T \in \{(5,3),(5,5),(10,3)\}$, see `figs/fig_learners.py`); the tables pool each method over its own cells.

against FLOP on in-trees that variance accounts for in part and that our theory does not predict.

5.3 Real-data evidence: can relatives be estimated from one record?

What this part tests. The benchmark experiments give a method many short replicate windows. A real record is one long, autocorrelated series, and the question becomes whether the relatives of the unrolled structure can be estimated from it at all (Q3), and whether the resulting count orients edges (Q4). This differs from the effect of the start, and it is tested on few independent units, because the units are segments of single records or stations that recur across subgraphs (Table 6). The overall result is negative or does not discriminate between hypotheses.

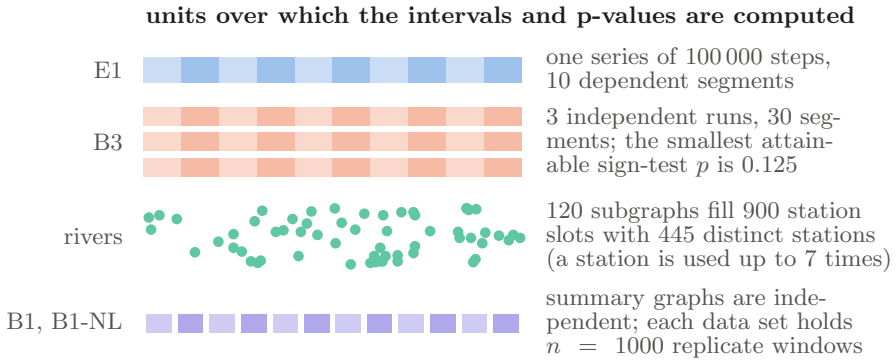


Fig. 14 The units behind the inference of each study (Table 6). Segments of one series, runs, and stations that recur across subgraphs are not independent replicates, so the real-data results are effect sizes, not tests with nominal error rates.

Protocol and its limits. We ran four studies (P1–P9, then E1/E2), each with an analysis plan (predictions, thresholds, estimator, seeds, stop conditions) committed to version control before any dependence statistic of its data was computed (timeline in Appendix K). Nine plans were written in all, counting the plans of B1, B1-NL, B3, the learner extension and the controls of Section 5.2, and the hosting service's push records were retrieved for three of them. The plans, analyses and results were produced by AI agents within about six hours of one day, commit times are set by the client and no external registry was used, so we call them *time-stamped analysis plans*, not pre-registrations. An earlier exploratory project in the same repository had already found unrelated CausalRivers stations about as correlated as river siblings, so our account of that failure is explanatory, not predictive; and the graph refinement used by E1 was motivated by study-3 observations on the same device (Appendix K).

Estimand. For an observed node set O of a known summary graph, the predicted value $\tilde{s}$ is the edge average of the scores of the counts $|\mathrm{rel}_S(v) \cap O|$: population lag-0 relatives among the observed nodes, as a perfect lag-0 dependence oracle would give them (Proposition 22, Definition 34). $\tilde{s} > 1/2$ forces an up-edge among the induced edges, but not conversely (Proposition 35, machine-checked). The estimator is the relatives count of Section 4.1 at $\alpha = 0.01$ with Bonferroni; orientation accuracy is the mean edge score of $\hat{\rho}$; CIs are cluster bootstraps. Lohse and Wahl (2025) measured var- and R^2-sortability on the same families; we know of no earlier test of rel-sortability on real data.

CausalRivers: a failed test, explained post hoc. River networks are in-forests, so Corollary 6 predicts $\tilde{s} > 1/2$ for subgraphs with confluences (mean 0.82). The estimated relatives sort carried no orientation information (P2, P3 fail; Table 7) because 98% of non-relative station pairs tested dependent. Read through F2 (post hoc), non-relative correlations $\kappa \approx 0.64$–0.74 give $n_{1/2} \approx 20$–27, while the records have $n = 377$–$3{,}652$ steps, so the study sat deep in the saturated regime (Figure 15, Appendix I), where Proposition 15 predicts $\hat{s} = 1/2$; regional precipitation acts as a latent driver. The F3 remedy fixed in plan 2 (conditioning on a leave-out regional mean) also failed (P4: 0.74), consistently with F3: the triad estimate of the proxy's

reliability, $\hat{r} \approx 0.62$, is below the required $r_{\text{crit}} \approx 0.87$, and the in-sample plug-in prediction of the dependent fraction, 0.743, matches the observed 0.741. The 120 subgraphs reuse stations (Table 6), so subgraph-level intervals understate the variance; wider intervals would leave P2 failed.

Causal Chambers: designed independent sources. In the wind-tunnel random-walk runs (Gamella et al., 2025) five actuators are independent random walks, hence separate source SCCs, and 25 of 26 edges are predicted strictly sorted. Saturation is absent (P7, 0.039) and the relatives count orients confluence edges above chance (P8, 0.656), but the graded agreement P9 missed its threshold ($0.267 < 0.3$), so by the plan the study is *not* a positive demonstration. The reason is F1: 16 of 26 edges are essentially undetectable at lag 0 at $n \approx 10^3$. The light-tunnel replication passes all three criteria, and F-theory predictions from the first half of each unit match the second half to within 0.02 (Table 8, Appendix I). The wind-tunnel graph is bipartite, so the prediction coincides with "sensors have more relatives than actuators".

E1 and B3: the theorem against naive rules. Plan 4 tested the edge values of Corollary 6 against eight naive predictors (sensor indicator, ancestor count, in-degree, depth, "every edge 1", ...) on the light-tunnel series `citris_1` (10 units). On the primary graph G_{lin} it passed 8/8 ($\Delta = -0.50$ to -0.53 in absolute loss, 10/10 units, Holm $p = 0.0078$). This is weaker than it looks (Appendix L): G_{lin} removes six real edges that carry no linear dependence under an ideal-sensor assumption, and on the published graph the sensor indicator wins on strictly predicted edges; each weak component of G_{lin} has a single source, so the prediction coincides with the rule "tie iff same weak component, else compare component sizes"; and most of the separation comes from ties. B3 (Appendix V) therefore used held-out light-tunnel data whose components have several independent sources (30 units of three AR(1)-driven series; 20 units of an actuator collider), with constant 1/2 and the weak-component rule as competitors. On the 2,361 instances where the theorem and the weak-component rule disagree, the counts were sorted as predicted in 99.8% ($\Delta = -0.50$ against both tie rules, 30/30 units); in the collider, R→G was sorted whenever B was observed (264/264) and tied otherwise (123/123). On these instances the theorem predicts strict order, so E1's naive sort rules agree with it there and are separated only on the 1,251 predicted ties whose own edge is detected ($\Delta = -0.43$, 30/30 units). The test was low-risk (same device; colour edges known to be detectable; results 3–4 minutes after the plan), and it failed where the graph is unfaithful: with an almost constant source the tie rule won ($\Delta = +0.12$, 0/10 units), and in one series the data contradicted G_{lin} (the design graph matched 0.96 against 0.04 where they differ). The p-values of E1 and B3 rest on units that are not independent. The ten E1 units are segments of one series, and a sign-flip test over ten units cannot give a Holm-adjusted p below 0.0078 against eight competitors, which is the value reported. The 30 B3 units come from three runs, and at the level of runs a sign-flip test cannot give p below 1/8. We read E1 and B3 as descriptions of effect sizes on one device, and not as tests with their nominal error rates.

E2: lagged relatives (Appendix M). A Bonferroni max-over-lags cross-correlation test (Haugh, 1976) on 64 held-out wind-tunnel regimes raised detection from 0.328 to

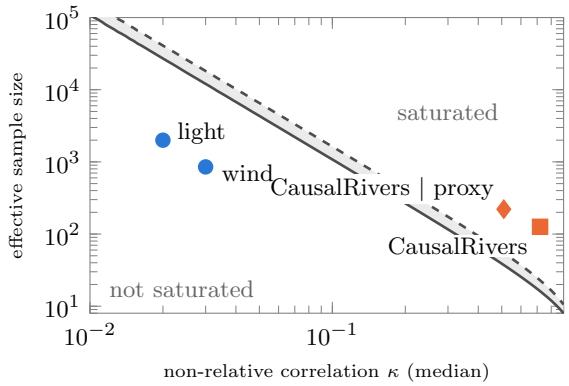


Fig. 15 Saturation regime of Proposition 15. Curves: onset $n_{1/2}$ for $z_c = 3.29$ ($d = 5$, solid) and 4.04 ($d = 20$, dashed); above them a non-relative pair is detected with probability $> 1/2$. Points: medians from the real-data studies (CausalRivers unconditioned and given the regional proxy, κ replaced by $|\pi|$; Causal Chambers wind and light tunnel). Values are in Table 7 and Appendix I.

0.398 without more false positives, but failed the graded criterion ($0.243 < 0.3$), so E2 does not succeed, and neither did the lag-0 estimator on these regimes.

Run- and station-level inference. A later plan (Appendix U) recomputed the real-data evidence at the level of independent units from committed results, with no new data. B3: the theorem sorted as predicted in 3 of 3 independent runs (H1) and 2 of 2 (H2), with 10 of 10 segments agreeing in each run; with so few runs the smallest attainable one-sided sign-test p is 1/8 and 1/4, so no run-level test can reach 0.05, and the segment-level p-values above are not run-level evidence. CausalRivers: the analysed subgraph set has 900 station slots and 445 distinct stations, one station is reused up to 7 times, and 112 of its 120 subgraphs share a station with another (25 connected components, the largest of 28 subgraphs); the 457 stations and reuse of 8 quoted in earlier drafts belong to a different, earlier subgraph set whose results are not analysed here. For the orientation accuracy of the conditioned relatives sort (P5, 12 h) a station bootstrap gives 0.539 $[0.527, 0.551]$, but in random station-disjoint subsets of about 42 subgraphs the lower bound of the interval exceeded 0.5 in only 27% of 200 orderings (median lower bound 0.491). The pass of P5 is therefore narrow and does not hold in most station-disjoint subsets. Sensitivity of the plug-in test to its truncation lag ($K \in \{20, 50, 100, 300\}$): at most 4% of decisions changed on the primary wind-tunnel analysis and on the river data, and none of 2,730 on B3; the E1 levels analysis is less stable (up to 13% changed). The re-fetched data matched every recorded SHA-256 value, and the rebuilt river caches were byte-identical. We conclude that the real-data tests cannot reach conventional significance at the level of independent runs or stations, and we read them as effect sizes.

Data selection, preprocessing and secondary analyses are in Appendix I.

Takeaway (Q2, Q4). Under a fixed start, the relatives-sorting baseline does better in every simulated condition we tried. The controlled tests of Section 5.2 show that its gain follows the correlations a test can detect, with the exact gap as the large-sample ceiling. Tuned competitive learners change rank against FLOP on in-trees;

variance is involved and does not explain the effect alone. In real records the relatives count orients edges only where the sources are designed to be independent and the relatives are detectable, which was the case on one device. The real-data results are too few, and too dependent, to say more.

6 Discussion, Limitations and Conclusions

Question. What should someone who builds a benchmark from an unrolled graph do, and what do our results not show?

6.1 Implications for benchmark design and real data

For benchmark design. These recommendations apply when slices of a simulated record are handed to a method that ignores time order; we found no deployed generator that does so as shipped (Section 5.1). The finite time-unrolled generator is rel-sortable above 1/2 for every summary graph that is not a union of strongly connected components (Corollary 3), strong connectivity alone does not help at short horizons (Section 5.1), and part of the excess comes only from the start at $t = 1$. We recommend three things. (a) Start trajectories from the stationary distribution (an exact Lyapunov or Cholesky start, as in Herman et al., 2025), or use a burn-in that is checked for distributional mixing. A burn-in of Δ slices ($\Delta\tau_{\max}$ with lags up to $\tau_{\max}$) removes the structural artefact (Theorem 5); it is sufficient, it is the minimal burn-in for chains (relatives: proved; value: computed for $p \leq 8$, $T \leq 5$) and for the one-edge graph, and a shorter burn-in suffices for many other graphs (Remark 20). It does not remove the distributional transient, which decays with the spectral radius: B1 needed 100 steps to pass its mixing check. As a rule of thumb, burn in for at least Δ slices and long enough for the covariance to be near its stationary value, or draw the start from the stationary law. (b) Report standardised columns as well as raw ones, since the fixed start also inflates var-sortability and the raw-data advantage of DAGMA and NOTEARS over FLOP under the fixed start was reduced, and for DAGMA's SHD removed, on standardised data. Standardisation alone is not a remedy: standardised DAGMA and GOLEM are below chance in order accuracy. Tuning the hyper-parameters does not remove the effect either: tuned DAGMA, NOTEARS and GOLEM-NV still beat FLOP on in-trees under the fixed start (Section 5.2). (c) Score the window without first-slice pairs or as an ADMG: a latent burn-in makes the last burn-in slice a set of latent common causes of the first window slice, so a labels-using reference that has normalised SHD 0.044 under the fixed start has 0.273 after a burn-in, but 0.041 against 0.048 without first-slice pairs; dropping them reduced the rel-SortnRegress gain from +0.125 to +0.089 and FLOP's SHD shift from +0.197 to +0.093 (Table 5). An algorithm that sees time stamps orients every cross-time edge for free, so the artefact matters for methods that do not use time order. Random summary graphs are sortable only near criticality. Removing the time labels does not hide the time order in the population under the fixed start (Section 3.4), and for a method that has the labels the start carries no extra information (Proposition 8).

For real data. Population rel-sortability above 1/2 does not imply that a relatives count can be estimated. Two conditions are needed: no latent common

driver with share ℓ such that $n \gtrsim z_c^2/\ell^2$ (or a near-perfect, rank-one proxy), and relatives detectable at lag 0 at the available n_{eff}. Diagnostic: compare n_{eff} with $3 + (z_c/\operatorname{atanh}\hat{\kappa})^2$, where $\hat{\kappa}$ is the typical correlation of pairs believed unrelated; if n_{eff} is far above it, the count is saturated whatever the graph.

6.2 Limitations

The closed forms assume self-loops at every node and unit lags (Appendices N and O). The burn-in bound is machine-checked sharp only for one edge ($\Delta = 1$), proved on paper to be sharp for the relatives of chains, and not sharp in general; we have enumeration and no closed form for the minimal burn-in. No human has compared the Lean statements with the paper's statements. The spectral radii of the transition matrices of B1 were bounded by 0.8 by construction and not tabulated. The ER limits are paper proofs not yet checked by a human probabilist; they are iterated limits ($T \to \infty$ at fixed p, then $p \to \infty$; the stationary value needs no T-limit), with no rates and no result for fixed T. Section 4.2 covers linear SVARs only. The finite-sample bounds are loose; Proposition 22 and the probabilistic parts of F1–F3 are proved on paper only, for Gaussian processes (the non-Gaussian CLT step is cited, not re-proved), and our code's fixed K controls size only away from a unit root. The analysis plans are internally time-stamped and agent-executed (push order cannot show that nothing was computed before a plan), and the CausalRivers plan was not blind to an earlier analysis of the same data. The positive real-data evidence comes from one device per chamber, and the units behind its p-values are not independent (Table 6); E1 coincides with a weak-component-size rule on its primary graph and is carried mainly by ties; B3 separates the theorem from that rule but on the same light tunnel, with no external registration. The graded criteria P2, P3, P4, P6, P9, Q9 and one B3 secondary set failed. The benchmark experiments are simulations at $n = 1000$ (the dose study goes to 25,000 for three families); the learner extension uses one linear-Gaussian parametrisation, $T \leq 5$ and default hyperparameters, and the tuned control covers DAGMA, NOTEARS ($d \leq 30$) and GOLEM ($d \leq 50$) with $T \leq 10$ and $p \leq 20$. Our theory does not predict the ranking changes among competitive methods. They survive tuning and point to a variance channel by the plan's rule, but imposing the fixed-start variances on burn-in data restored less than half of the advantage, and we have no theory or experiment that explains the rest; the plan's prediction of no rank change failed. The detectability calculation and the re-analyses with graph-cluster inference are post hoc. The designed counterfactual separates the population pattern from detectable correlations, but only for one estimator (rel-SortnRegress) at $n = 1000$ and in the dose study; we have not shown that the exact gap itself, as opposed to its detectable image, drives the effect for any learner. At the level of independent runs or stations the real-data tests cannot reach conventional significance (B3 has three runs; the P5 interval excludes 0.5 in 27% of station-disjoint subsets). Exposure of deployed generators to the variance artefact was not measured (no mixing times). The identifiability statements of Section 3.4 are paper proofs that assume faithfulness, which we have not shown to hold for almost all tied parameters; Conjecture 13 is open; and they give no finite-sample guarantee and no algorithm.

6.3 Conclusions

Time order survives label removal as a footprint that the unrolled graph determines. For lag-one unrollings with self-loops we computed the footprint exactly: the value of every edge, its limit, and a burn-in after which it no longer depends on the start, with the graph statements checked in Lean. The burn-in is sufficient, minimal for chains (proved for the relatives, computed for the value with $p \leq 8$ and $T \leq 5$) and the one-edge graph, and longer than needed for many other graphs. Under a fixed start the footprint is inflated, strongly connected summary graphs are not protected at short horizons, and a relatives-sorting baseline gained 0.13 in order accuracy in simulation. Controlled tests show that this gain follows the correlations a test can detect, not the population relatives pattern, and that the exact gap is the large-sample ceiling of the effect, reached in part (0.76) at $n = 25{,}000$. How much of the footprint an estimator sees depends on the sample size, the autocorrelation and latent drivers; the real-data support was mostly negative and, at the level of independent runs or stations, cannot reach conventional significance. Tuned competitive learners change rank against FLOP on in-trees under the fixed start; variance is involved but does not explain the effect on its own, and our theory does not cover it. The finite-sample and identifiability statements are paper proofs, with the limits listed above. A builder who hands the slices of a simulated record to static learners should start from the stationary distribution or use a burn-in that has been checked for mixing, and should score the window without first-slice pairs.

6.4 Use of AI tools

Most of this project was carried out with AI coding agents (large language models, Claude-based) working under the author's direction; no AI system is an author. The paper proofs (including Appendices P.1 and W), all Lean code, all experiment and analysis code, the analysis plans, the simulated pre-submission reviews that motivated several revisions, and the manuscript text were generated by these agents. What is machine-checked: the Lean theorems of Section 3.5 (163 declarations in the axiom report, standard axioms only, out of 259 theorem and lemma declarations in the sources), which cover the exact and stationary results of Section 3.2 including the Δ forms of Theorems 2 and 5 and the sharpness example for one edge, the estimand algebra and the deterministic algebra of F1–F3. No human has checked that the Lean statements say what the paper statements say; the map in Appendix H is the agents' reading. What is checked only by computation: the enumeration of Appendices C and N, the ER numerics and the simulations of Appendix G. What no human has yet read line by line: the paper-only proofs, namely Appendix P.1, Proposition 22, Theorem 23, the probabilistic parts of Theorem 28 and Propositions 14 to 16 and 31, Proposition 32, Remark 33, the loop-free limit results of Appendix N, the lag-weighted extension (Appendix O), and Propositions 8 to 12. The author takes responsibility for all content.

Supplementary information. Code, a pinned environment (`requirements.txt`; Lean 4.35.0, Mathlib v4.35.0), dataset sources with SHA-256 checksums (`DATA.md`), the analysis plans, run logs and a reproduction script (`reproduce.sh`, with stages for

the Lean build, the graph checks, the simulations including the benchmark experiments and the learner extension, the post hoc re-analyses, and the real-data studies; `REPRODUCE.md`) accompany the article. The full stages of the script have not been rerun end to end, because of the CPU budget, and not every committed output can be regenerated by one command: the seeds behind some random-graph tables were not recorded, the sweeps of the C++ checks are not scripted, and the single random graph with $p = 10^6$ is not scripted (`REPRODUCE.md` lists them).

Statements and Declarations

- **Funding.** The author received no funding for this work.
- **Competing interests.** The author declares no competing interests.
- **Ethics approval and consent to participate.** Not applicable. The study uses simulations and public sensor data sets and involves no human participants, human data or animals.
- **Consent for publication.** Not applicable.
- **Data availability.** No new data were collected. The real-data studies use the public CausalRivers benchmark (Stein et al., 2025) and the public Causal Chambers data sets (Gamella et al., 2025) (CC BY 4.0); sources, versions and SHA-256 checksums are listed in `DATA.md` of the supplementary repository, and the raw data are not redistributed. All simulated data are regenerated from fixed seeds by the reproduction script.
- **Code availability.** All code, including the Lean development, is in the supplementary repository, https://github.com/vinhqdang/KAN_time_series (directory `research/theory2`).
- **Materials availability.** Not applicable.
- **Author contributions.** Q.-V. Dang is the sole author: he conceived and directed the study, supervised the AI agents described in Section 6.4, and takes responsibility for all content.
- **Use of AI tools.** See Section 6.4.

Appendix A Overview and status labels

Appendix B gives full proofs of Section 3.2 (including Remark 20), Appendix C the exhaustive enumeration check, Appendices D and E the statements and proofs for Section 4.2, Appendix F the proofs of Section 4.1 (including the plug-in effective sample size and the scope of the recovery bound), Appendix G the simulation checks of F1–F3, Appendix H the map from statements to Lean names, Appendix I data selection and preprocessing for the real-data studies, Appendix J the real-data estimand $\tilde{s}$, Appendix K the analysis-plan timeline for all studies and experiments, prior exposure and deviations, Appendices L and M the experiments E1 and E2 in full, Appendices N to Q the loop-free graphs, lagged unrollings, the random summary graphs (with the full proof of Proposition 45 in Appendix P.1) and the accuracy interpretation, Appendices R to T and V the benchmark experiment B1, its replication B1-NL, the extension to competitive learners with the post hoc re-analysis, detectability calculation and generator survey, and the discriminating test B3, Appendix U the plan,

design and verdicts of the controlled tests of Section 5.2, and Appendix W the proofs of Sections 3.3 and 3.4.

Status labels. *Machine-checked*: a Lean theorem with no `sorry` and only the standard axioms. *Proved*: a complete paper proof given here. *Asymptotic*: a limit statement proved here from the multivariate Breuer–Major theorem (Arcones, 1994; Nourdin et al., 2011) (cited, not re-proved). *Sketched*: a derivation with one analytic step left out, stated as such (used only for statements outside the main claims: the windowed loop-free stationary value, the layered-DAG stationary value and the constants of the lagged recovery bound). *Enumerated*: checked by exact computation on many instances, not proved.

Appendix B Proofs for Section 3.2

Definition of s_ρ, verbatim. Reisach et al. (2026, Sec. 2.2) (arXiv v1) write: "For a graph $G = (V, E)$ with $|E| \geq 1$ we define sortability by the number of relatives, or rel-sortability for short, as $s_\rho(G) := \frac{\sum_{(x,y)\in E} s_\rho(x,y)}{|E|}$ with $s_\rho(x, y) := 0$ if $\rho(x) > \rho(y)$, $\frac{1}{2}$ if $\rho(x) = \rho(y)$, 1 if $\rho(x) < \rho(y)$, for $(x, y) \in E$." (their Eq. (4); the cases are typeset as a brace there), and add: "Note that our definition differs from the sortabilities in these works in that we count edges instead of paths." Our s_ρ is this edge average; in G_T the edges are the copies $(i(t), j(t+1))$.

Throughout, $G_S = (V, E_S)$ has a self-loop at every node, $p = |V| \geq 2$, and $d(u, v)$ is as in Section 3.1. We use two facts of Reisach et al. (2026), proved there for DAGs: (Thm. 1) if $x \in \mathrm{anc}(y)$ then $\mathrm{rel}(x) \subseteq \mathrm{rel}(y)$; (Lemma 1, eq. (10)) if $x \in \mathrm{anc}(y)$ then $\mathrm{rel}(x) \subsetneq \mathrm{rel}(y)$ iff $\mathrm{roots}(y) \setminus \mathrm{roots}(x) \neq \emptyset$, where $\mathrm{roots}(y) = \{x \in \mathrm{anc}(y) : \mathrm{anc}(x) = \{x\}\}$. Both are also machine-checked independently (relSet_subset_of_edge, and the direct characterisation isRel_unrolled_iff).

Lemma 17 (ancestors in G_T) *$u(s) \in \mathrm{anc}(v(t))$ iff $s \leq t$ and $d(u, v) \leq t - s$.*
Lean: unrolled_reach_iff, unrolled_reach_iff_dist.

Proof Every edge of G_T increases time by one, so a directed path from $u(s)$ to $v(t)$ has $t - s$ edges, and its projection to V is a walk of length $t - s$ from u to v in G_S (self-loops allowed as steps). Such a walk exists iff $d(u, v) \leq t - s$: a walk contains a path, so its length is at least $d(u, v)$; conversely, pad a shortest path with $t - s - d(u, v)$ self-loop steps at u. Every such walk lifts to a directed path in G_T. □

Lemma 18 (roots) *$\mathrm{roots}(v(t)) = \{u(1) : d(u, v) \leq t-1\}$, and $\mathrm{desc}(u(1)) = \{w(s) : d(u, w) \leq s - 1\}$.*
Lean: roots_unrolled_iff, isRoot_unrolled_iff, desc_root_unrolled_iff.

Proof A node $u(s)$ with $s \geq 2$ has the parent $u(s - 1)$ through the self-loop, so it is not a root; $u(1)$ has no parent. Apply Lemma 17 with $s = 1$, in both directions. □

Proof of Theorem 1 $i(t)$ is a parent of $j(t{+}1)$, so $\mathrm{rel}(i(t)) \subseteq \mathrm{rel}(j(t{+}1))$ and $s_\rho(e) \in \{1/2, 1\}$, with value 1 iff the inclusion is strict, iff $\mathrm{roots}(j(t{+}1)) \setminus \mathrm{roots}(i(t)) \neq \emptyset$. By Lemma 18 this set is $\{u(1) : d(u,j) \leq t,\ d(u,i) > t-1\}$, and since d is integer-valued or ∞, $d(u,i) > t-1$ iff $d(u,i) \geq t$. □

Proof of Proposition 4 (and the general count) Since every ancestor of y descends from a root of y that is itself an ancestor, $\mathrm{rel}(y) = \mathrm{desc}(\mathrm{roots}(y))$. By Lemma 18, $w(s) \in \mathrm{rel}(v(t))$ iff some u has $d(u,v) \leq t-1$ and $d(u,w) \leq s-1$, i.e. iff $s \geq 1 + m_t(v,w)$ with $m_t(v,w) = \min\{d(u,w) : d(u,v) \leq t-1\}$. Counting $s \in \{1,\dots,T\}$ gives $\rho(v(t)) = \sum_w \max(T - m_t(v,w), 0)$ (rho_unrolled_eq). For $t \geq p$ every finite $d(u,v)$ is at most $p-1 \leq t-1$, so the admissible u are $\mathrm{anc}_S(v)$ and $m_t(v,w) = d(\mathrm{anc}_S(v), w)$, finite iff $w \in \mathrm{rel}_S(v)$ and then at most $p-1 < t \leq T$. (The Lean statement rho_unrolled_large_t needs only $t \geq p$; an earlier hypothesis $T \geq p$ is redundant.) □

Proof of Theorem 2 Crossing edge $(i,j) \in E_\times$: j does not reach i, so $d(j,i) = \infty$, and $u = j$ gives $d(j,j) = 0 \leq t \leq \infty$ for every t: all $T-1$ copies score 1. Edge in $E_\circ$: i and j reach each other (or $i = j$), so for every u, $d(u,i) < \infty$ iff $d(u,j) < \infty$. If the condition of Theorem 1 held with $d(u,i) = \infty$ then $d(u,j) = \infty > t$, a contradiction; hence it forces $t \leq d(u,i) \leq \Delta$, and at most $N_T(i,j) \leq \min(\Delta, T-1)$ copies score 1 ($\Delta \leq p-1$, maxFinDist_le_card_sub_one). Summing over the $(T-1)|E_S|$ edges,

$$s_\rho(G_T) = \tfrac{1}{2} + \frac{1}{2(T-1)|E_S|}\Big[(T-1)|E_\times| + \sum_{(i,j)\in E_\circ} N_T(i,j)\Big],$$

which gives both bounds; let $T \to \infty$. Lean: sRho_unrolled_bounds_maxFinDist (the Δ form), with the $p-1$ form sRho_unrolled_bounds as a corollary. □

Proof of Corollary 3 $E_\times = \emptyset$ iff no edge joins two SCCs. A weakly connected component made of two or more SCCs is connected by undirected edges, one of which joins two different SCCs; conversely a strongly connected component has no crossing edge. The uniform bound is Theorem 2 with $|E_\times| \geq 1$. For the DAG case, every SCC is a singleton, $E_\times$ consists of the m non-loop edges, and $|E_S| = m + p$. □

Proof of Theorem 5 In G_{B+T}, by Lemma 18 shifted by B, $w(s) \in \mathrm{rel}(v(t))$ iff some u has $d(u,v) \leq t-1$ and $d(u,w) \leq s-1$. For window nodes $t-1, s-1 \geq B \geq \Delta$, and every finite distance is at most Δ, so the condition is $u \in \mathrm{anc}_S(v) \cap \mathrm{anc}_S(w)$, independent of s,t. Hence $\mathrm{rel}_W(v(t)) = \{w(s) \in W : w \in \mathrm{rel}_S(v)\}$, of size $T|\mathrm{rel}_S(v)|$. On $V \times \mathbb{Z}$, Lemma 17 holds verbatim and the common ancestor can be taken arbitrarily early, giving the same set. Each copy of (i,j) therefore compares $|\mathrm{rel}_S(i)|$ with $|\mathrm{rel}_S(j)|$; as $i \in \mathrm{anc}_S(j)$, $\mathrm{rel}_S(i) \subseteq \mathrm{rel}_S(j)$, so the score is 1 iff $\mathrm{rel}_S(i) \neq \mathrm{rel}_S(j)$, and 1/2 for self-loops. Source SCCs: $\mathrm{rel}_S(j) \setminus \mathrm{rel}_S(i) \neq \emptyset$ iff $\mathrm{anc}_S(j) \not\subseteq \mathrm{rel}_S(i)$ (if $z \in \mathrm{rel}_S(i)$ then $\mathrm{desc}_S(z) \subseteq \mathrm{rel}_S(i)$). Every node of $\mathrm{anc}_S(j)$ is reached from some $\sigma \in \mathrm{Src}(j)$, and $\mathrm{desc}_S(\sigma) \subseteq \mathrm{rel}_S(i)$ whenever $\sigma \subseteq \mathrm{rel}_S(i)$. A source SCC meets $\mathrm{rel}_S(i)$ iff it reaches i (its only ancestors are its own nodes), and then it is contained in $\mathrm{rel}_S(i)$. So the condition is $\mathrm{Src}(j) \not\subseteq \mathrm{Src}(i)$, which, as $\mathrm{Src}(i) \subseteq \mathrm{Src}(j)$, is $\mathrm{Src}(i) \neq \mathrm{Src}(j)$. Lean: rhoBurn_eq_maxFinDist, sBurn_eq_maxFinDist (with $B \geq \Delta$), and the $B \geq p-1$ forms rhoBurn_eq, sBurn_eq as corollaries. □

Proposition 19 (sharpness of $B \geq \Delta$ for the relatives of a chain) *Let the summary graph be the directed chain $0 \to 1 \to \cdots \to p-1$ with a self-loop at every node, $p \geq 2$, so $\Delta = p-1$. Take a latent burn-in of $B = \Delta - 1 = p-2$ slices before a window of any length $T \geq 1$. Then the window nodes $(0,1)$ and $(p-1,1)$ are not relatives, whereas they are relatives under the stationary start. Hence the relatives set of $(0,1)$ differs from its stationary value, and so does the relatives count ρ.*

Proof Number the window slices $1, \ldots, T$ and the burn-in slices $-B+1, \ldots, 0$. The parents of (v,t) are $(v,t-1)$ and $(u,t-1)$ for each summary edge (u,v). The ancestors of $(0,1)$ are the copies $(0,s)$, $s \leq 1$. A copy (u,s) is an ancestor of $(p-1,1)$ iff $1-s \geq p-1-u$, because each step of the chain takes one slice and self-loops only allow waiting. A common ancestor must be some $(0,s)$, which requires $1-s \geq p-1$, that is $s \leq 2-p$. The earliest slice present with $B = p-2$ burn-in slices is $3-p > 2-p$, so there is no common ancestor. Neither node is an ancestor of the other: $(0,1)$ would need $p-1 \geq 1$ slices to reach node $p-1$, and edges do not point backwards along the chain. With the stationary start (a burn-in of $B \geq \Delta = p-1$ slices, or an infinite past) the copy $(0, 2-p)$ exists and is a common ancestor, so the two nodes are relatives. The relatives count ρ of $(0,1)$ therefore differs by at least one. □

This shows that the relatives, and so the counts, differ at $B = \Delta - 1$; it does not by itself show that the edge-level value s_ρ^W differs, which we computed exactly for $p \leq 8$ and $T \leq 5$ (below). The proposition is a paper proof; it is not checked in Lean.

Remark 20 (sharpness of $B \geq \Delta$) For the single edge $0 \to 1$ with self-loops, $\Delta = 1$; with $B = 1$ the window value is the stationary value $1/2$, and with $B = \Delta - 1 = 0$ it is at least $2/3$ for every $T \geq 2$ (edgeLoops_burn_sharp). So $B \geq \Delta$ cannot be lowered for this graph; this is the part checked in Lean. Proposition 19 extends the statement for the relatives of the chain to every length, on paper. An exact computation (`b1y/q6_burnin_exact.py`, independent of the theorems; Appendix U) took the minimal burn-in $B_{\min}$ to be the smallest B such that the window relatives counts, and separately s_ρ^W, equal their stationary values for every larger B. It gave: chains, $p = 2, \ldots, 8$ and $T = 2, \ldots, 5$ (28 cases): $B_{\min} = \Delta$ for both measures in every case; directed rings in the same range: $B_{\min} = \lceil \Delta/2 \rceil$; 396 random summary graphs ($p = 3, \ldots, 7$, $T \in \{3, 5\}$; 792 cases): never $B_{\min} > \Delta$, $B_{\min} = \Delta$ in 406 cases (51%), and slack $\Delta - B_{\min}$ of 1, 2, 3 and 4 in 250, 114, 20 and 2 cases; counts and s_ρ gave the same $B_{\min}$ in every case. So $B \geq \Delta$ is the minimal burn-in for the one-edge graph and, in this range, for chains, and it is longer than needed for many other graphs. The enumeration of 726 random summary graphs (Appendix R) agrees: stationary at $B = \Delta$ in all of them and at $B = \Delta - 1$ in 281. We have no closed form for the minimal burn-in of a given graph, and no proof that s_ρ^W differs at $B = \Delta - 1$ for chains of every length.

Proof of Corollary 6 If each weak component C has a unique source SCC σ_C, then $\mathrm{Src}(v) = \{\sigma_C\}$ on C and $E_\uparrow = \emptyset$. Conversely, if $E_\uparrow = \emptyset$, Src is constant along edges and hence on each weak component; a node v of a source SCC σ has $\mathrm{Src}(v) = \{\sigma\}$, so there is only one. Out-forests (at most one parent besides itself; acyclicity is not needed, stat_half_of_outForest) and strongly connected graphs (stat_half_of_strong) have one source SCC per component. In an in-forest (acyclic apart from self-loops, out-degree at most one), if $u \in \mathrm{anc}(v)$ then $\mathrm{desc}(u)$ is the unique downstream path through v, so $\mathrm{rel}(v) = \mathrm{anc}(v) \cup \mathrm{desc}(v)$ (inForest_isRel_iff); the

source SCCs are the headwaters, and $\mathrm{Src}(j) = \mathrm{Src}(i) \cup \bigcup_k \mathrm{Src}(k)$ over the other parents k of j, with disjoint headwater sets on different branches, so $\mathrm{Src}(j) \neq \mathrm{Src}(i)$ iff j is a confluence (inForest_up_iff). Out-degree at most one alone is not enough: for $a \rightleftarrows b$ with $c \to a$, a has two parents but the only source SCC is $\{c\}$. □

Remark 21 (self-loops) Without self-loops the lemmas fail: for the single edge $0 \to 1$ without loops, $s_\rho(G_T) = 1/2$ for every $T \geq 2$ (edgeNoLoops_sRho); with loops the limit is $2/3$ (edgeLoops_limit).

Appendix C Exhaustive Enumeration Check

Before formalisation, every statement of Section 3.2 was checked by brute force (bitset relatives with the definitions of Reisach et al., 2026) on all $4 + 64 + 4096$ summary graphs with $p = 2, 3, 4$ (all $T \leq 9$) and on 10,000 random graphs ($p = 2$–7; Erdős–Rényi, DAG, in-forest, out-forest, chain with back edge and block-SCC families; three values of $T \leq 14$ each, fixed seed). In total 46,724 (graph, T) cases and 1,895,609 edges were checked with zero failures of Theorem 1, Proposition 4, the bounds of Theorem 2, Theorem 5 (with $B = p - 1$ and $B = 2p + 1$) and the source-SCC characterisation. The conjecture fixed in plan 1 ("an edge is strictly sorted in the limit iff $|\mathrm{rel}_S(j)| > |\mathrm{rel}_S(i)|$") disagreed with the exact G_T limit in 9,603 random graphs and never with the stationary value; the conjecture is Theorem 5, not the limit of G_T.

Appendix D Block Means and Subsamples: Statements

For a stable linear SVAR with independent innovations, nodes with disjoint summary ancestor sets are independent at all lags, and so are any per-node, time-invariant transforms, including block means and subsamples; otherwise the lag-0 covariance is nonzero for almost every parameter (Proposition 22). So the population relatives of contemporaneous values are rel_S for every observation scheme, and the population relatives sort is the stationary value of Theorem 5 (on non-self-loop edges), not the finite-G_T value. Var- and R^2-sortability behave differently (Theorem 23): block means converge, at the textbook Fejér rate for averages of a stationary series (Brockwell & Davis, 1991, Ch. 7), to the long-run structural model of Gong et al. (2017), whereas subsamples keep the snapshot covariance. As a consequence, for large block size, block-mean R^2-sortability is 0 on out-stars and 1 on in-stars with at least two arms for all weights; on out-stars subsampled data are R^2-sorted correctly on average (0.84) while block means drift to 0. We claim only these sortability consequences.

Assumptions: $X(t) = \sum_{\tau=1}^{L} A_\tau X(t-\tau) + \varepsilon(t)$, $t \in \mathbb{Z}$, i.i.d. innovations with diagonal positive-definite covariance Σ, stable (all roots of $\det(I - \sum_\tau A_\tau z^\tau)$ outside the unit disc), stationary solution; the summary graph has $i \to j$ iff $(A_\tau)_{ji} \neq 0$ for some τ. $\Gamma(h)$ is the autocovariance, $A(1) = \sum_\tau A_\tau$, and for a block size D the block means are $\bar{Y}_s = \frac{1}{D} \sum_{u<D} X(sD+u)$ and the subsamples $Z_s = X(sD)$.

Proposition 22 (relatives survive any per-node filter) *If* $\mathrm{anc}_S(i) \cap \mathrm{anc}_S(j) = \emptyset$ *then* X_i *and* X_j *are independent at all lags, and so are any per-node, time-invariant (possibly nonlinear) transforms of them, including* $\bar{Y}$ *and* Z *for every* D*. If the ancestor sets intersect,* $\Gamma(0)_{ij}$, $\mathrm{Cov}(\bar{Y})_{ij}$ *and the long-run covariance are nonzero for Lebesgue-almost every parameter.*

Lean: not machine-checked (needs the SVAR moving-average representation and real-analytic genericity).

Theorem 23 (block means vs subsampling) *(a)* $\|D\,\mathrm{Cov}(\bar{Y}) - S_{LR}\| \le \frac{1}{D}\sum_h |h|\|\Gamma(h)\|$*, with* $S_{LR} = (I - A(1))^{-1}\Sigma(I - A(1))^{-\top}$*. (b) If* G_{S} *is acyclic apart from self-loops,* S_{LR} *is the covariance of a static SEM (this long-run SEM is due to Gong et al., 2017, Eqs. (5), (10)) on* G_{S} *with weights* $A(1)_{ji}/(1-\lambda_j)$ *and noise variances* $\sigma_j^2/(1-\lambda_j)^2$*,* $\lambda_j = A(1)_{jj} < 1$*. (c)* $\mathrm{Cov}(Z_s) = \Gamma(0)$ *for every* D*. (d) Hence, for large* D*, block-mean* R^2*-sortability is* 0 *on out-stars and* 1 *on in-stars with* ≥ 2 *arms, for all weights; under subsampling every contemporaneous criterion keeps its* $\Gamma(0)$ *value.*

Lean: (a) counting identity fejer_identity; (b) longrun_factor, selfloop_sum_lt_one; the probabilistic steps, the rate and (c)–(d) are not machine-checked.

Part (a) is the standard Fejér/Cesàro computation for the variance of the mean of a stationary series (Brockwell & Davis, 1991, Ch. 7), written for block means of a vector process; (b) is due to Gong et al. (2017); (c) is stationarity. We claim only the sortability consequences (d).

The R^2 statement for stars strengthens the three-node analysis in Herman et al. (2025, App. D), whose qualitative message (that R^2 reflects connectivity, with an asymmetry between colliders and confounders) we credit. Numerical observations (exact population moments, not proved): mean block-mean R^2-sortability over 150 in-trees moves from 0.853 ($D = 1$) to 0.834 ($D = 4$–8) and 0.880 ($D = 256$), so it is not monotone in D; on out-stars subsampled data are R^2-sorted at 0.842 on average while block means give 0.546 ($D = 4$), 0.101 ($D = 16$) and 0 in the limit.

Appendix E Proofs for the Block-Mean and Subsampling Statements

Proof of Proposition 22 By the MA(∞) representation $X(t) = \sum_k \Psi_k \varepsilon(t-k)$, and because Ψ_k is supported on the reachability relation, $X_i(t)$ is a measurable function of $\{\varepsilon_k(s) : k \in \mathrm{anc}_S(i)\}$. Disjoint ancestor sets mean independent innovation sets, so the processes X_i and X_j are independent, and so are any per-node transforms. For the generic part, each covariance entry is real-analytic in (A, Σ) on the connected open set of stable parameters with the given support, hence either identically zero or nonzero almost everywhere. With all coefficients positive, including the self-loop at a common ancestor k, all Ψ entries are nonnegative and $\Psi_{ik}(m), \Psi_{jk}(m) > 0$ for all large m (self-loops pad the walks), so $\Gamma(0)_{ij} > 0$; the same holds for $\Gamma(h)$, $h \ge 0$, and therefore for the Fejér sums and S_{LR}. "Generic" does not mean "large": with little autocorrelation, $\Gamma(0)_{ij}$ along a lagged edge can be tiny (down to 10^{-14} in our exact computations; exactly 0 without autocorrelation). □

Proof of Theorem 23 (a) $D^2\mathrm{Cov}(\bar{Y}) = \sum_{u,w<D} \Gamma(u-w) = D\Gamma(0) + \sum_{h=1}^{D-1}(D-h)(\Gamma(h) + \Gamma(-h))$ (the counting identity is fejer_identity). The stable VAR has $\|\Gamma(h)\| \le Cr^{|h|}$, $r < 1$,

so $\sum_h |h|\|\Gamma(h)\| < \infty$ and $S_{LR} = \sum_h \Gamma(h) = \Psi(1)\Sigma\Psi(1)^\top$ with $\Psi(1) = (I - A(1))^{-1}$. Then $S_{LR} - D\,\mathrm{Cov}(\bar{Y}) = \sum_{|h|\geq D} \Gamma(h) + \sum_{|h|<D} \frac{|h|}{D}\Gamma(h)$, and $1 \leq |h|/D$ in the first sum. (b) Write $I - A(1) = (I - \Lambda) - N$ with Λ the diagonal and N the off-diagonal part; then $I - A(1) = (I - \Lambda)(I - \tilde{B})$ with $\tilde{B} = (I - \Lambda)^{-1}N$ strictly triangular in a topological order (longrun_factor, longrun_inv). Under acyclicity the AR polynomial factorises into univariate $\phi_j(z) = 1 - \sum_\tau (A_\tau)_{jj} z^\tau$; stability, $\phi_j(0) = 1$ and the intermediate value theorem give $\phi_j(1) = 1 - \lambda_j > 0$ (selfloop_sum_lt_one). (c) Stationarity. (d) For a linear SEM $X = BX + e$ with noise variances ω, $R_i^2 = 1 - 1/(\Sigma_{ii}(\Sigma^{-1})_{ii})$ and $(\Sigma^{-1})_{ii} = 1/\omega_i + \sum_{k\in\mathrm{ch}(i)} B_{ki}^2/\omega_k$. In an out-star with root 0 and leaves k, $1 - R_0^2 = 1/(1 + \sum_k r_k)$ and $1 - R_k^2 = 1/(1 + r_k)$ with $r_k = \omega_0 B_{k0}^2/\omega_k > 0$, so with ≥ 2 leaves the root has strictly larger R^2 than every leaf; the in-star is symmetric. Strict inequalities in the limit SEM of (b), plus (a) and continuity of $S \mapsto R^2$ at positive-definite S, give the claim for all large D. Under subsampling every contemporaneous statistic is a function of $\Gamma(0)$. □

Appendix F Proofs for Section 4.1

Notation. $\hat{C}_{ab} = \frac{1}{n}\sum_t (X_a(t) - \bar{X}_a)(X_b(t) - \bar{X}_b)$, $\hat{r}_{ab} = \hat{C}_{ab}/\sqrt{\hat{C}_{aa}\hat{C}_{bb}}$. Constants: B_{ab} as in the main text; $D_\gamma = \max_a 2\sum_h \gamma_{aa}(h)^2$; $\kappa_1 = \max_a \sum_h |\gamma_{aa}(h)|$; $\kappa_2 = \max_{a\neq b}(\sum_h |\gamma_{aa}(h)| + \sum_h |\gamma_{ab}(h)|)$. For i.i.d. sampling, $B_{ab} = 1 + \rho_{ab}^2$, $D_\gamma = 2$, $\kappa_1 = 1$.

Lemma 24 (Gaussian chaos tail) *Let $g \sim N(0, I_N)$, A symmetric, $S = g^\top Ag - \mathrm{tr}A$. For $x > 0$, $P(S \geq 2\|A\|_F\sqrt{x} + 2\|A\|_{op}x) \leq e^{-x}$, and the same for $-S$.*

Lean: algebraic steps phi_le, chernoff_logmgf_le, chernoff_choice, sqrt_two_v_x; the probability statement is not machine-checked.

Proof Diagonalise; S has the law of $\sum_i \lambda_i(g_i^2 - 1)$. For $2s\lambda < 1$, $\log Ee^{s\lambda(g^2-1)} = \varphi(s\lambda)$ with $\varphi(u) = -\frac{1}{2}\log(1 - 2u) - u = \sum_{k\geq 2}(2u)^k/(2k) \leq u^2/(1 - 2|u|)$ for $|u| < 1/2$. With $c = 2\|A\|_{op}$, $v = 2\|A\|_F^2$ and $0 \leq s < 1/c$, $\log Ee^{sS} \leq vs^2/(2(1 - cs))$. Markov's inequality with $t = \sqrt{2vx} + cx$, $q = \sqrt{2x/v}$ and $s = q/(1 + cq)$ gives exponent exactly x. This is the Gaussian case of Hsu et al. (2012, Prop. 1), who derive it from Laurent and Massart (2000); the proof here is self-contained. □

Lemma 25 (exact variance) *With $\tilde{\gamma}_{ab} = \frac{1}{n}\sum_t X_a(t)X_b(t) = g^\top Ag$ for $A = C^{1/2}MC^{1/2}$, $M = \frac{1}{2n}\left(\begin{smallmatrix} 0 & I \\ I & 0 \end{smallmatrix}\right)$,*

$$2n\|A\|_F^2 = \sum_{|h|<n} \left(1 - \tfrac{|h|}{n}\right)\left[\gamma_{aa}(h)\gamma_{bb}(h) + \gamma_{ab}(h)\gamma_{ab}(-h)\right] \leq B_{ab},$$

$\|A\|_{op} \leq \kappa_2/(2n)$; on the diagonal $2n\|A\|_F^2 \leq D_\gamma$ and $\|A\|_{op} \leq \kappa_1/n$; and $\bar{X}_a \sim N(0, \bar{\sigma}_a^2)$ with $\bar{\sigma}_a^2 \leq \kappa_1/n$.

Lean: trace and Toeplitz algebra quad_M, trace_MC, lemma12_trace, lemma12_fejer, lemma12_le_B, lemma12_diag_le_D; the step $\|C^{1/2}MC^{1/2}\|_F^2 = \mathrm{tr}(MCMC)$ and the operator-norm bounds are not machine-checked.

Proof $\|A\|_F^2 = \mathrm{tr}(MCMC)$ and, with Toeplitz blocks $(C_{ab})_{st} = \gamma_{ab}(s-t)$, $4n^2\mathrm{tr}(MCMC) = 2\sum_{s,t}[\gamma_{aa}(s-t)\gamma_{bb}(s-t) + \gamma_{ab}(s-t)\gamma_{ab}(t-s)]$ (using that γ_{aa}, γ_{bb} are even); count pairs with $s-t=h$. $\|A\|_{op} \le \|C\|_{op}\|M\|_{op}$ and $\|C\|_{op}$ is at most the largest absolute row sum. □

Lemma 26 (deviation event) *Let* $\varepsilon_{ab} = \sqrt{2B_{ab}x/n} + \kappa_2 x/n + 2\kappa_1 y/n$ *and* $\varepsilon_d = \sqrt{2D_\gamma y/n} + 4\kappa_1 y/n$. *With probability at least* $1 - d(d-1)e^{-x} - 4de^{-y}$, $|\hat{C}_{ab} - \rho_{ab}| \le \varepsilon_{ab}$ *for all* $a \ne b$ *and* $|\hat{C}_{aa} - 1| \le \varepsilon_d$ *for all* a.

Lean: deterministic steps centred_cov_eq, centred_dev_offdiag, centred_dev_diag, lemma13_eps, lemma13_count.

Proof $\hat{C}_{ab} = \tilde{\gamma}_{ab} - \bar{X}_a\bar{X}_b$. Apply Lemmas 24 and 25 to each off-diagonal and diagonal $\tilde{\gamma}$, bound $|\bar{X}_a| \le \sqrt{2\kappa_1 y/n}$ by the Gaussian Chernoff bound, and take a union bound. □

Lemma 27 (correlation perturbation) *If* $a', b' \in [1-\varepsilon_d, 1+\varepsilon_d]$, $\varepsilon_d < 1$ *and* $|c' - \rho| \le \varepsilon$, *then* $|c'/\sqrt{a'b'} - \rho| \le (\varepsilon + |\rho|\varepsilon_d)/(1-\varepsilon_d)$.

Lean: corr_perturb (the hypothesis $|\rho| \le 1$ *of our earlier statement is not needed).*

Proof $|c'/\sqrt{a'b'} - \rho| \le |c' - \rho|/\sqrt{a'b'} + |\rho||1/\sqrt{a'b'} - 1|$, with $1/\sqrt{a'b'} \le 1/(1-\varepsilon_d)$ and $|1/\sqrt{a'b'} - 1| \le \varepsilon_d/(1-\varepsilon_d)$. □

Theorem 28 (recovery bound) *For* $x, y > 0$ *let* ε_{ab}, $\varepsilon_d < 1$ *be as in Lemma 26 and* $\delta_{ab} = (\varepsilon_{ab} + |\rho_{ab}|\varepsilon_d)/(1-\varepsilon_d)$. *If thresholds separate* ($\delta_{ab} \le \tau_{ab}$ *on non-relatives and* $|\rho_{ab}| - \delta_{ab} > \tau_{ab}$ *on relatives), the declared set equals the true one, and hence* $\hat{s} = s_\rho^{\mathrm{pop}}$, *with probability at least* $1 - d(d-1)e^{-x} - 4de^{-y}$. *For a common threshold* $\tau < \rho_{\min}$ *the failure probability decays as* $\exp(-n\min(\tau, \rho_{\min} - \tau)^2/(2(\sqrt{B} + \rho_{\max}\sqrt{D_\gamma})^2))$.

Lean: deterministic core thm15_det, decide_eq_target, corr_perturb, sOf_cntDecl_congr; Chernoff and trace algebra chernoff_choice, phi_le, lemma12_trace, lemma12_le_B. The Gaussian tail bound itself is not machine-checked.

The proof combines a Gaussian-chaos tail bound (the Gaussian case of Hsu et al., 2012, Prop. 1; see also Laurent & Massart, 2000), an exact finite-n variance identity for $\hat{r}$, and a deterministic perturbation lemma.

Proof of Theorem 28 Set $\delta_{ab} = (\varepsilon_{ab} + |\rho_{ab}|\varepsilon_d)/(1-\varepsilon_d)$. On the event of Lemma 26, Lemma 27 gives $|\hat{r}_{ab} - \rho_{ab}| \le \delta_{ab}$, so non-relatives satisfy $|\hat{r}| \le \tau$ and relatives $|\hat{r}| > \tau$ (decide_eq_target, thm15_det). Equal declared sets give equal counts and hence $\hat{s} = s_\rho^{\mathrm{pop}}$ (sOf_cntDecl_congr). For a VAR observed as a long stationary series, s_ρ^{pop} is the stationary value of Theorem 5 on non-self-loop edges (Proposition 22). Explicit form: with $y = x + \log(4/(d-1))_+$ (explicit_y) and a common threshold $\tau < \rho_{\min}$, $P(\hat{R} \ne R) \le 2d(d-1)\exp\big(-n\min(\tau, \rho_{\min} - \tau)^2/(2(\sqrt{B} + \rho_{\max}\sqrt{D_\gamma})^2)(1 + o(1))\big)$. □

The multivariate Breuer–Major theorem. We use it in the form stated by Nourdin et al. (2011, Thm. 1.1), who attribute the vector case to Arcones (1994): let $X = (X_k)_{k\in\mathbb{Z}}$ be a centred stationary Gaussian process in $\mathbb{R}^d$ with covariances

$r^{(i,l)}(j) = E[X_1^{(i)} X_{1+j}^{(l)}]$, and $f : \mathbb{R}^d \to \mathbb{R}$ with $Ef(X_1)^2 < \infty$ and Hermite rank $q \geq 1$ with respect to X (orthogonal to all polynomials of degree $< q$, not to some of degree q). If $\sum_{j\in\mathbb{Z}} |r^{(i,l)}(j)|^q < \infty$ for all i, l, then $n^{-1/2} \sum_{k\leq n}(f(X_k) - Ef(X_k)) \Rightarrow N(0, \sigma^2)$ with $\sigma^2 = \mathrm{Var} f(X_1) + 2\sum_{k\geq 1} \mathrm{Cov}(f(X_1), f(X_{1+k})) \in [0, \infty)$.

Proof of Proposition 14 Take $d = 2$, $X_k = (X_a(k), X_b(k))$ and, for $c \in \mathbb{R}^3$, $f_c(x) = c_1(x_a^2 - 1) + c_2(x_b^2 - 1) + c_3(x_a x_b - \rho_{ab})$. f_c is a centred polynomial of degree 2 without linear part, so it is orthogonal to every polynomial of degree ≤ 1 (odd Gaussian moments vanish); if it were orthogonal to every polynomial of degree 2 it would be orthogonal to itself, so either $f_c(X_1) = 0$ a.s. (then the partial sums vanish) or its Hermite rank is 2. The condition $\sum_h |\gamma_{ab}(h)|^2 < \infty$ holds because the covariances are absolutely summable and bounded. By the theorem and the Cramér–Wold device, $\sqrt{n}(\tilde\gamma_{aa} - 1, \tilde\gamma_{bb} - 1, \tilde\gamma_{ab} - \rho_{ab}) \Rightarrow N(0, V)$, where, by Isserlis' theorem, $V(ab, cd) = \sum_h [\gamma_{ac}(h)\gamma_{bd}(h) + \gamma_{ad}(h)\gamma_{bc}(h)]$ (as in Lemma 25). Centring by the sample means changes each entry by $\bar X_a \bar X_b$, and $\sqrt{n}\bar X_a \bar X_b \to 0$ in probability since $E(\sqrt{n}\bar X_a)^2 \leq \kappa_1$ and $\bar X_b \to 0$; Slutsky's lemma applies. The delta method with gradient $(-\rho/2, -\rho/2, 1)$ gives $\sqrt{n}(\hat r - \rho) \Rightarrow N(0, W)$, $W = V(ab, ab) - \rho[V(ab, aa) + V(ab, bb)] + \frac{\rho^2}{4}[V(aa, aa) + 2V(aa, bb) + V(bb, bb)]$, which equals $\sum_h \gamma_{aa}(h)\gamma_{bb}(h)$ for an independent pair (this is B_{ab} when all products $\gamma_{aa}(h)\gamma_{bb}(h)$ are nonnegative, and at most B_{ab} in general). Then $\sqrt{n}\,\operatorname{atanh}\hat r \Rightarrow N(0, W_{ab})$ under the null and $P(|\operatorname{atanh}\hat r| > z_c/\sqrt{n-3}) \to 2\bar\Phi(z_c/\sqrt{W_{ab}})$, which exceeds the nominal α/m in the limit whenever $W_{ab} > 1$. □

Lemma 29 (sample autocovariances) *Let X_a be a standardised stationary Gaussian series with $\sum_h |\gamma_{aa}(h)| \leq \kappa_1$ and $\hat\gamma_a(h) = \frac{1}{n}\sum_{t\leq n-h}(X_a(t+h) - \bar X_a)(X_a(t) - \bar X_a)$. For $0 \leq h < n$, $E(\hat\gamma_a(h) - \gamma_{aa}(h))^2 \leq (8\kappa_1 + 108\kappa_1^2)/n + 2(h/n)^2\gamma_{aa}(h)^2$.*

Proof For $\tilde\gamma(h) = \frac{1}{n}\sum_{t\leq n-h} X(t+h)X(t)$, $E\tilde\gamma(h) = (1-h/n)\gamma(h)$ and, by Isserlis, $\mathrm{Var}\,\tilde\gamma(h) = n^{-2}\sum_{s,t\leq n-h}[\gamma(s-t)^2 + \gamma(s-t+h)\gamma(s-t-h)] \leq 2\kappa_1/n$, since $|\gamma| \leq 1$ gives $\sum_k \gamma(k)^2 \leq \kappa_1$ and $\sum_k |\gamma(k+h)\gamma(k-h)| \leq \kappa_1$. The centring term is $-\bar X(S_1 + S_2)/n + \frac{n-h}{n}\bar X^2$ with S_1, S_2 partial sums; $\bar X$, S_1/n, S_2/n are centred Gaussians with variance $\leq \kappa_1/n$ and $E(UV)^2 \leq 3\sigma^4$ for such pairs, so its second moment is $\leq 27\kappa_1^2/n^2$. Combine with $(u+v+w)^2 \leq 4u^2 + 4v^2 + 2w^2$ for $u + v$ and the bias $w = h\gamma(h)/n$. □

Lemma 30 (plug-in Bartlett factor) *Let X_a, X_b satisfy the assumptions of Section 4.1 (no condition on their cross-dependence), $\hat\rho_a(h) = \hat\gamma_a(h)/\hat\gamma_a(0)$, $\hat W_n = 1 + 2\sum_{h=1}^{K_n} \hat\rho_a(h)\hat\rho_b(h)$ and $W_{ab}^\circ = \sum_{h\in\mathbb{Z}} \rho_a(h)\rho_b(h)$. If $K_n \to \infty$ and $K_n/n \to 0$, then $|\hat W_n - W_{ab}^\circ| = O_P(n^{-1/2} + K_n/n) + 2\sum_{h>K_n} |\rho_a(h)\rho_b(h)| \to 0$; for fixed K, $\hat W \to 1 + 2\sum_{h\leq K} \rho_a(h)\rho_b(h)$ in probability.*

Proof With $\Delta_a = \hat\rho_a - \rho_a$, on $\{\hat\gamma_a(0), \hat\gamma_b(0) \geq \frac{1}{2}\}$ (probability $\to 1$), $|\Delta_a(h)| \leq 2|\hat\gamma_a(h) - \gamma_{aa}(h)| + 2|\gamma_{aa}(h)||\hat\gamma_a(0) - 1|$. Write $\hat\rho_a\hat\rho_b - \rho_a\rho_b = \Delta_a\rho_b + \rho_a\Delta_b + \Delta_a\Delta_b$. By Lemma 29, $E\sum_{h\leq K} |\hat\gamma_a(h) - \gamma_{aa}(h)||\rho_b(h)| \leq \kappa_1\sqrt{c_0/n} + \sqrt{2}\kappa_1 K/n$, so the linear terms are $O_P(n^{-1/2} + K/n)$; by Cauchy–Schwarz and $E\sum_{h\leq K}(\hat\gamma_a - \gamma_{aa})^2 \leq Kc_0/n + 2\kappa_1(K/n)^2$, $\sum_{h\leq K} |\Delta_a\Delta_b| = O_P(K/n)$; the tail is at most $\sum_{h>K} |\rho_a(h)| \to 0$. □

Only $K_n = o(n)$ is needed, because the target is a sum of products of two summable sequences; we do not rely on general HAC bandwidth conditions.

Proposition 31 (plug-in effective sample size) *Under the assumptions of Proposition 14, use the plug-in threshold* $\tanh(z_c/\sqrt{n/\hat{B}_{ab} - 3})$ *with* $\hat{B}_{ab} = \max(1, \hat{W}_n)$*, where* $1 \le K_n = o(n)$*. (i) For an independent pair with* $W_{ab} > 0$ *and* $K_n \to \infty$*, the false-positive probability tends to* $2\bar{\Phi}(z_c\sqrt{\max(1, W_{ab})/W_{ab}})$*, which is* α/m *if* $W_{ab} \ge 1$ *and smaller otherwise. For fixed* K *it tends to* $2\bar{\Phi}(z_c\sqrt{\max(1, W^{\circ}_{K,ab})/W_{ab}})$*,* $W^{\circ}_{K,ab} = 1 + 2\sum_{h \le K} \rho_a(h)\rho_b(h)$*. (ii) Every pair with* $\rho_{ab} \ne 0$ *is declared with probability* $\to 1$*. (iii) If the non-relatives are independent pairs (summary non-relatives of a Gaussian VAR) and every relative has* $\rho_{ab} \ne 0$*, then* $\liminf_n P(\hat{R} = R) \ge 1 - |N|\alpha/m \ge 1 - \alpha$*, and the same holds for* $P(\hat{s} = s^{\mathrm{pop}}_{\rho})$*, provided* $K_n \to \infty$*.*

Lean: not machine-checked (uses the multivariate Breuer–Major theorem through Proposition 14).

Proof (i) $T = |\sqrt{n}\,\operatorname{atanh}\hat{r}|\,(1/\hat{B} - 3/n)^{1/2}$. By Proposition 14, $\sqrt{n}\,\operatorname{atanh}\hat{r} \Rightarrow N(0, W)$, and $\hat{B} \to b = \max(1, W)$ (or $\max(1, W^{\circ}_K)$) in probability by Lemma 30. By Slutsky, $T \Rightarrow |N(0, W/b)|$, whose law has no atom at z_c. (ii) By Lemma 25, $\hat{r}_{ab} \to \rho_{ab}$ in probability. By Cauchy–Schwarz $|\hat{\rho}_a(h)| \le 1$, so $\hat{B} \le 1 + 2K_n = o(n)$ and $n/\hat{B} \to \infty$. (iii) Union bound over N with (i), and (ii) on R; equal declared sets give equal $\hat{s}$. □

Proposition 32 (scope of Theorem 28) *Let the non-relatives* $N \ne \emptyset$ *have deterministic thresholds with* $n\tau^2_{n,ab} \to t^2_{ab} \in (0, \infty)$*:* $t^2 = z_c^2$ *for the nominal Fisher-z test, and* $t^2 = z_c^2 b_{ab}$ *for the effective one with* $n' = n/b_{ab}$*. Then the smallest failure probability certified by Theorem 28 tends to* $d(d-1)\exp(-\min_N t^2_{ab}/(2B_{ab})) > 0$ *when the relatives have* $|\rho| \ge \rho_{\min} > 0$ *(and is at least this in general). For i.i.d. data and the nominal test this is* $d(d-1)e^{-z_c^2/2} \ge \sqrt{2\pi}\, z_c\alpha$ *(Mills), e.g.* 0.089 *for* $d = 5$*,* $\alpha = 0.01$*. If instead* $\alpha_n \to 0$ *with* $\log(1/\alpha_n) = o(n)$*, Theorem 28 gives* $P(\hat{R} \ne R) \le d(d-1)\exp(-(1-\eta)z^2_{c,n}r/2) + 4d/n \to 0$ *for any* $\eta > 0$ *and large* n*, with* $r = \min_N (n/n'_{ab})/B_{ab}$*.*

Proof For $\rho_{ab} = 0$ the separation condition $\delta_{ab} \le \tau_{n,ab}$ and $\varepsilon_{ab} \ge \sqrt{2B_{ab}x/n}$ force $x \le n\tau^2_{n,ab}/(2B_{ab})$, which gives the lower bound. For attainment, take $x_n = (1 - \eta)\min_N n\tau^2_{n,ab}/(2B_{ab})$ and $y_n = \log n$. Then all lower-order terms are $o(n^{-1/2})$ and $\varepsilon_d \to 0$, so the conditions hold for large n on N, and on the relatives because $\tau_n \to 0$. The vanishing-level statement is the same computation with $z_{c,n} \to \infty$. □

Remark 33 (non-Gaussian innovations) Let $X(t) = \sum_j \Psi_j \varepsilon(t - j)$ with $\sum_j \|\Psi_j\| < \infty$ and i.i.d. innovations with independent unit-variance coordinates and fourth cumulants $\kappa_i = E\varepsilon_i^4 - 3$. Put $G_i(a, b) = \sum_j \Psi_j[a, i]\Psi_j[b, i]$. Then $n\,\mathrm{Cov}(\tilde{\gamma}_{ab}, \tilde{\gamma}_{cd}) \to V(ab, cd) + \sum_i \kappa_i G_i(a, b)G_i(c, d)$ (moment–cumulant formula and dominated convergence), so $W^{\mathrm{ng}}_{ab} = W_{ab} + \sum_i \kappa_i g_i^2$ with $g_i = G_i(a, b) - \frac{\rho_{ab}}{2}(G_i(a, a) + G_i(b, b))$. Independent series share no innovation, so $g_i = 0$: the null variance $\sum_h \rho_a(h)\rho_b(h)$, and hence Propositions 14 and 31, are unchanged, given the CLT for linear processes, which we do not re-prove. For relatives, W changes by $\sum_i \kappa_i g_i^2$ (of either sign, since $\kappa_i \ge -2$). This affects power and transition widths,

not the population statements or the location of $n_{1/2}$. The non-asymptotic bounds use the Gaussian chaos tail. For sub-Gaussian innovations the Hanson–Wright inequality (Rudelson & Vershynin, 2013, Thm. 1.1) has the same form with unspecified constants.

Fixed K in practice (simulation and data check). Independent AR(1) pairs, 4,000 replications (2,000 at $n = 64{,}000$): at AR coefficient 0.9 and 0.97 the plug-in test with $K = 50$ has size 0.050 and 0.056 (level 0.05; limits 0.05 and 0.0556), but at 0.99 it has 0.113 (limit 0.118), while $K_n = n^{0.6}$ gives 0.0495 and the nominal test 0.853; for the opposite-sign pair the floor makes the test conservative (size 0). On the B3 primary data (30 segments of 3,000 rows, 2,730 unit-pairs) the declared status of no pair changes when $K = 50$ is replaced by 55, 122 or 300 (post hoc). This was not checked for study 3, E1 or the river data. Scripts: `finite2/`.

F2 and F3. The statements are Propositions 15 and 16 in Section 4.1.

Proof of Proposition 15 $\mathrm{Cov}(X_u, X_v) = \gamma^Y_{uv}(0) + b_u b_v$ and $\gamma^Y_{uv} = 0$ for non-relatives (Proposition 22); divide by the standard deviations (prop21_corr_std). For relatives the sum is nonzero except on a Lebesgue-null set. All counts equal d, so every edge scores 1/2 (sOf_saturated). The rate: apply Theorem 28 with target "all pairs" and $x = y = cn$, c small enough that $\delta \le \kappa_{\min}/3$ and $\tau_n \le \kappa_{\min}/3$ eventually; $|\hat{s} - 1/2| \le 1/2$ always (abs_sOf_sub_half_le). The onset: by the asymptotic normality of $\operatorname{atanh} \hat{r}$, the declaration probability crosses 1/2 where $\tanh(z_c/\sqrt{n-3}) = |\kappa|$, i.e. at $n_{1/2}$ (ncrit_eq_iff). □

Proof of Proposition 16 $\mathrm{Cov}(X_u, Z) = b_u\sqrt{r}$ and $\mathrm{Var}Z = 1$; the partial covariance is $\gamma^Y_{uv} + (1-r)b_u b_v$ and the partial variance $\mathrm{Var}Y_u + (1-r)b_u^2$; divide and standardise (prop31, prop31_std; positive $\mathrm{Var}Y_u$ is required). At $r = 1$, $\pi = \rho^Y$ (cor32). For non-relatives $|\pi_{uv}| = \sqrt{\ell_u \ell_v}\sqrt{f_u f_v}$ with $f_w(r) = (1-r)/(1-r\ell_w)$ and $f'_w(r) = (\ell_w - 1)/(1-r\ell_w)^2 < 0$ (absPi_strictAntiOn, absPi_eq_zero_iff; both need $b_u b_v \neq 0$). Bounds $(1-r)\sqrt{\ell_u\ell_v} \le |\pi| \le (1-r)\sqrt{\ell_u\ell_v}/(1-\ell_{\max})$ (absPi_lower, absPi_upper). r_{crit} solves $(1-r)\ell/(1-r\ell) = \tau'$ (rcrit_iff, rcrit_example). The finite-sample bracket uses the Lipschitz bound $|g(a',b',c') - g(a,b,c)| \le L(q)\max(|a'-a|, |b'-b|, |c'-c|)$ for $g(a,b,c) = (a-bc)/\sqrt{(1-b^2)(1-c^2)}$, $|b|, |c| \le q < 1$, $L(q) = 1/(1-q^2) + 4q/(1-q^2)^2$ (pc_lipschitz, thm35_bracket), combined with Lemma 26 for (X, Z). Multi-filter driver: for $D_v(t) = b_v C(t - k_v)$ with C an AR(1) with coefficient ϕ and $Z = C(t)$,

$$\pi_{uv} = \frac{b_u b_v(\phi^{|k_u - k_v|} - \phi^{k_u + k_v})}{\sqrt{(s_u^2 - b_u^2\phi^{2k_u})(s_v^2 - b_v^2\phi^{2k_v})}},$$

nonzero for $k_u, k_v > 0$, $\phi \in (0,1)$, $b_u b_v \neq 0$ (prop36_ar1, prop36_ar1_ne_zero); in general $\pi_{uv} \propto \mathrm{Cov}(D_u - P_Z D_u, D_v - P_Z D_v)$ vanishes for all non-relative pairs iff the residual driver components are uncorrelated, which holds for a rank-one driver with $Z = F$ (prop36_pcov, prop36_zero_iff, prop36_rank_one). A proxy that contains summary relatives of u, v adds the terms $-c_Z(b_u q_v + b_v q_u) - q_u q_v$ to the numerator (prop37_pcov, prop37_bound). Triad identity: $\rho_{uZ}\rho_{vZ}/\rho_{uv} = r$ and, for $r > 0$, $\ell_u = \rho^2_{uZ}/r$ (prop38_ratio, prop38_loading). □

Appendix G Simulation Checks of F1–F3

All with fixed seeds; full tables are in the supplementary material. **F1** (5 processes: in-forests with $d = 5$ and $d = 10$, self-loop persistence $\lambda \in \{0.5, 0.9, 0.97\}$, and a near-white-noise case; $n = 200$–$20{,}000$; 1,000 replications per cell; nominal, oracle-effective

and fixed thresholds). The bound was never violated in 105 cells. It is loose: with a fixed threshold $\rho_{\min}/2$, it reaches 0.9 at about 3–13 times the n at which the simulated probability does, and much later with Bonferroni thresholds, where the Chernoff step loses the Mills-ratio factor. The empirical $n\,\mathrm{Var}(\hat{r})/W$ lies in $[0.93, 1.11]$. With the nominal test, $P(\hat{R} = R)$ plateaus at 0.94, 0.22, 0.03 and 0.27, as Proposition 14 predicts (predicted vs simulated false-positive rates, e.g. 0.33 vs 0.32–0.33 at $\lambda = 0.9$). **F2** ($\ell \in \{0.02, 0.1, 0.3, 0.64\}$): the asymptotic detection curve matches the simulated fraction of declared non-relatives within 0.02 for $n \geq 200$; predicted $n_{1/2} = 1{,}079$ for $\ell = 0.1$, simulated 0.479 at $n = 1{,}000$. **F3**: predicted $n_{\mathrm{crit}} = 472$ ($\ell = 0.64$, $r = 0.9$) gives a simulated declared fraction 0.495 at $n = 471$; over 38 cells the asymptotic prediction is within 0.03. With travel-time lags $(0, 1, 2, 4, 8)$, $\phi = 0.9$, $r = 1$, the declared fraction still rises to 0.63 at $n = 10{,}000$, as the multi-filter part of Proposition 16 predicts.

Appendix H Lean Formalisation: Map and Findings

Lean 4.35.0 and Mathlib v4.35.0; 4,159 lines in 16 files. The build log (a clean rebuild of the final state, 9,064 jobs, including Observed.lean and the Δ refinements), the `#print axioms` output for each of the 163 declarations in the axiom report (all `[propext, Classical.choice, Quot.sound]`) and a grep showing no `sorry`/`admit`/`axiom` are in the supplementary material.

What is and is not certified. Certified: the exact and stationary results of Section 3.2, the loop-free exact edge value, the accuracy reading, the algebra of the estimand $\tilde{s}$ (Appendix J), and the deterministic and algebraic parts of Sections 4.1 and 4.2. Not certified: the ER limits (only the deterministic witness lemmas), the loop-free limit, periodicity and stationary value, the lag-weighted extension, and every probability bound, limit theorem and genericity statement.

Breakdown of the 163 declarations (counted from the Lean files): 33 headline theorems, i.e. Lean counterparts of numbered paper statements (29 for Section 3.2, the loop-free edge value and the accuracy reading, including the four Δ forms of Theorems 2 and 5 and Corollary 6; 4 for Proposition 35); 53 supporting graph lemmas (ancestry, roots, distance, the largest finite distance Δ (maxFinDist, 10 lemmas), source SCCs, 5 ER witness lemmas, 6 estimand lemmas); 65 algebraic identities and inequalities of the aggregation and finite-sample files (correlation perturbation, Chernoff algebra, trace and Toeplitz identities, partial correlation formulas, threshold locations; several are elementary, e.g. $|\hat{s} - \frac{1}{2}| \leq \frac{1}{2}$); 12 examples (10 concrete graphs, the sharpness example edgeLoops_burn_sharp, 1 numerical r_{crit} instance). The sources contain 259 theorem and lemma declarations in all; the axiom report lists the 163 above, and the other 97 are not in it. No human has checked, line by line, that each Lean statement says what the corresponding paper statement says; the map below is the authors' (AI-generated) reading.

statement	Lean names	status
Lemmas 17 and 18	unrolled_reach_iff, unrolled_reach_iff_dist, roots_unrolled_iff, desc_root_unrolled_iff, isRel_unrolled_iff	checked
Theorem 1	edge_score, edge_score_dist, strictCond_iff_dist	checked
Proposition 4	rho_unrolled_eq, mT_le_iff, rho_unrolled_large_t	checked
Theorem 2	sRho_unrolled_exact, sRho_unrolled_bounds_maxFinDist (Δ form), sRho_unrolled_bounds ($p-1$ form), tendsto_sRho_unrolled, tendsto_sRho_GT; Δ: maxFinDist, dist_le_maxFinDist, exists_dist_eq_maxFinDist, maxFinDist_le_card_sub_one	checked
Corollary 3	tendsto_half_iff, crossEdges_eq_empty_tfae, sRho_unrolled_ge_of_cross, reisach_theorem4, limitValue_dag	checked
Theorem 5	rhoBurn_eq_maxFinDist, sBurn_eq_maxFinDist, relBurn_iff_maxFinDist, walkN_of_reach_maxFinDist ($B \geq \Delta$); rhoBurn_eq, sBurn_eq ($B \geq p-1$); rhoZ_eq, sZ_eq, sOf_window, relSet_eq_iff_src, mem_upEdges_iff	checked
Remark 20 (sharpness of $B \geq \Delta$ for one edge)	edgeLoops_burn_sharp, sBurn_zero, relBurn_zero_iff, maxFinDist_edgeLoops	checked
Corollary 6	sBurn_eq_half_iff_maxFinDist, sBurn_eq_half_iff, sBurn_eq_half_iff_common_anc, sZ_eq_half_iff, upEdges_eq_empty_tfae, stat_half_of_strong, stat_half_of_outForest, inForest_isRel_iff, inForest_up_iff, statValue_inForest, statValue_inForest_gt_iff	checked
Table 2	collider_limit, fork_limit, chain3_limit, collider_stat, fork_stat, chain3_stat	checked
self-loop remark	edgeLoops_limit, edgeNoLoops_sRho	checked
Theorem 37 (ancestors, relatives, edge value)	isRel_unrolled_iff_NL, reduce_ancestor, relSet_unrolled_eq_iff_NL, edge_score_NL	checked
Corollary 38	sRho_unrolled_eq_NL, sRho_unrolled_ge_half_NL	checked
Proposition 41 (witness $u = j$, fresh parent)	strictNL_of_closed_walk, strictNL_of_fresh_parent, sRho_unrolled_eq_one_NL	checked
loop-free examples (Table N6)	colliderNL_sRho, edgeLoopHead_sRho, edgeNoLoops_sRho	checked
eventual periodicity, limit, same-SCC edges, Theorem 43	–	paper / enumerated

statement	Lean names	status
Proposition 45	witness lemmas upEdge_witness, mem_upEdges_of_witness, mem_upEdges_of_disjoint_anc, upEdges_subset_confluenceNL, card_upEdges_le_confluenceNL	paper proof (App. P.1); lemmas checked
Proposition 44	–	paper only
Proposition 35 (a)–(d)	relObs_subset_of_edge, score_obsEdge, sTilde_eq, sTilde_ge_half, upObs_subset_upEdges, sTilde_gt_half_iff, exists_upEdge_of_sTilde_gt_half, sTilde_eq_half_of_no_up, mem_upObs_of_witness, upObs_univ	checked
Proposition 55 (oracle sort accuracy; optimality)	accRule_sortRule, accRule_sortRule_rho, accRule_le_sOf, accRule_le_sRho, edgeAcc_sortRule_strict, rho_not_gt_of_edge	checked
Proposition 22	–	paper only
Theorem 23(a)	fejer_identity (counting identity)	partial
Theorem 23(b)	longrun_factor, longrun_inv, selfloop_sum_lt_one	partial
Theorem 23(c,d)	–	paper only
Lemma 24	phi_le, chernoff_logmgf_le, chernoff_choice, sqrt_two_v_x	algebra only
Lemma 25	quad_M, trace_MC, lemma12_trace, lemma12_fejer, avg_toeplitz_le_tsum, lemma12_le_B, lemma12_diag, lemma12_diag_le_D	algebra only
Lemma 26	centred_cov_eq, centred_dev_offdiag, centred_dev_diag, lemma13_eps, lemma13_count	deterministic part
Lemma 27	corr_perturb	checked
Theorem 28	decide_eq_target, decide_eq_support, thm15_det, explicit_y, sOf_cntDecl_congr	deterministic part
Proposition 14	–	asymptotic (cited CLT)
Proposition 15	prop21_corr, prop21_corr_std, sOf_saturated, sOf_const, abs_sOf_sub_half_le, thm22_const, ncrit_iff, ncrit_eq_iff	deterministic part
Proposition 16	pcorr_eq_pcov, pcorr_eq_corr_resid, prop31, prop31_std, cor32, prop33, prop33_abs, absPi_lower, absPi_upper, absPi_le_kappa, absPi_lt_kappa, absPi_eq_zero_iff, absPi_strictAntiOn, absPi_equal, rcrit_iff, rcrit_eq_iff, rcrit_example, pc_lipschitz, thm35_bracket, prop36_pcov, prop36_zero_iff, prop36_rank_one, prop36_ar1_pcov, prop36_ar1_pvar, prop36_ar1, prop36_ar1_ne_zero, prop37_pcov, prop37_bound, prop38_ratio, prop38_loading	population part checked

Modelling choices. Summary graphs are relations on a finite type with a self-loop hypothesis; anc and desc are reflexive-transitive closures ("a path exists iff a walk exists" is the one convention adopted without proof); G_T with $T = n+1$ slices lives on $V \times \mathrm{Fin}(n+1)$; source SCCs are represented by their member nodes (equality of SCC sets as equality of node unions is elementary and not formalised). In the finite-sample files, zero-mean variables are vectors of an arbitrary real inner-product space and lag-0 covariance is the inner product; independence enters only as vanishing inner products.

Findings. No statement was false. Formalisation added four missing nondegeneracy hypotheses to our earlier finite-sample statements: (i) the strict claims of Proposition 16 need $b_u b_v \neq 0$; (ii) $\ell_u = \rho_{uZ}^2/r$ needs $r > 0$; (iii) nonvanishing of the multi-filter residual needs $b_u b_v \neq 0$; (iv) the partial correlation needs positive partial variances ($\mathrm{Var} Y_u > 0$ suffices). It also showed that one hypothesis of Proposition 4 and one of Lemma 27 are redundant, and that the in-forest statement needs acyclicity, not only out-degree at most one.

Appendix I Real-Data Details

The account of the real-data studies, with Tables 7 and 8 and Figure 15, is in Section 5.3. This section gives data selection, preprocessing and secondary analyses.

CausalRivers (Stein et al., 2025). Study 1 (P1–P3): the benchmark's connected-subgraph sampler, seeds 7000–7029, East Germany and Bavaria, $d \in \{5, 10\}$, 120 subgraphs, 791 induced edges; block means at 12 h and 24 h; nominal Fisher-z. Study 2 (P4–P6): seeds 8000–8029, 120 subgraphs, 1,376 non-relative pairs; proxy Z is the mean of standardised stations of the region outside the subgraph and outside 2 undirected river hops (429–615 stations); primary test on partial correlations with a residual Bartlett n_{eff} ($K = 100$ steps at 12 h, 50 at 24 h); median n_{eff} 139/113 (raw) and 231/211 (residual). Before conditioning, 0.98 (nominal) and 0.87 (effective) of non-relative pairs tested dependent. Median proxy contamination by river relatives was 2.6% (maximum 72%), a case of the contaminated-proxy term in Proposition 16. Post hoc and without verdict: conditioning on $k = 1, 3, 5, 10$ principal components of the leave-out stations gave dependent fractions $0.745, 0.605, 0.576, 0.491$ (12 h), with accuracy 0.580 and Spearman 0.16 at $k = 10$.

Causal Chambers (Gamella et al., 2025). Candidate selection (recorded before analysis) compared Causal Chambers, CausalRivers with measured precipitation (rejected: the driver reaches each gauge through its own catchment filter, so by Proposition 16 a scalar proxy cannot remove it), a semi-synthetic cross-sectional manufacturing benchmark (rejected: not real, not time series) and fMRI with physiological regressors (rejected: no ground truth). Wind tunnel: runs `actuators_random_walk_2`–`16` (run 1 excluded before analysis: truncated); 16 random variables, 26 edges, 63 non-relative pairs per run; block means over 10 rows, first differences, standardisation, $n = 999$; effective-n Fisher-z with Bartlett $K = 20$. Light tunnel: `actuators_white`, 20 variables, 39 edges, ten segments of 2,000 steps, levels. Predicted values come from the package's published ground-truth graph with constant parameters removed. Secondary analyses (no verdicts): nominal n gives slightly more false positives (0.048), as Proposition 14 predicts; levels instead of increments gave 0.054, contradicting our own

secondary prediction that levels of random walks would produce many spurious dependences; without block means ($D = 1$) relative recovery drops to 0.18 and Spearman to 0.10. In the light tunnel, polariser angles act on two sensors through Malus' law, a nonlinearity with near-zero linear correlation; as stated in the plan, these relative pairs were never detected, without affecting orientation.

Appendix J The Real-Data Estimand $\tilde{s}$

The theorems of Section 3.2 concern s_ρ of an unrolled DAG. The real-data studies estimate something else: lag-0 dependences between the observed series of one long record. This section defines the population quantity they target and proves what it does and does not say.

Definition 34 (estimand) Let $G_{\mathrm{S}} = (V, E_{\mathrm{S}})$ be a summary graph (no loop assumption) and $O \subseteq V$ the observed nodes. For $v \in O$ let $\tilde{\rho}_O(v) = |\mathrm{rel}_S(v) \cap O|$ (relatives computed in the full graph G_{S}, so unobserved nodes act as latent common ancestors). Let $E_O = \{(i,j) \in E_{\mathrm{S}} : i, j \in O,\ i \neq j\}$ be the induced non-self-loop edges and, for $E_O \neq \emptyset$, $\tilde{s}(G_{\mathrm{S}}, O) = \frac{1}{|E_O|} \sum_{(i,j)\in E_O} s_\rho\big(\tilde{\rho}_O(i), \tilde{\rho}_O(j)\big)$, with the edge score of Section 3.1. Edges of the latent projection that are not edges of G_{S} (paths through unobserved nodes) are not scored.

Why this is the target. Under the SVAR model of Appendix E, Proposition 22 gives, for almost every parameter, $\Gamma(0)_{uv} \neq 0$ iff $\mathrm{anc}_S(u) \cap \mathrm{anc}_S(v) \neq \emptyset$ iff $v \in \mathrm{rel}_S(u)$, for every observation scheme. So a perfect lag-0 dependence oracle on the observed series declares exactly $\mathrm{rel}_S(v) \cap O$, and the relatives sort of the observed counts, averaged over E_O, is $\tilde{s}$. (Lag-0 faithfulness is the generic case of Proposition 22; it can fail for specific parameters, e.g. through Malus-law nonlinearities.) Relation to Theorem 5: with $O = V$, every window node of the stationary unrolled DAG has $\rho_W(v(t)) = T|\mathrm{rel}_S(v)|$, so $\tilde{s}(G_{\mathrm{S}}, V)$ is the stationary value computed over non-self-loop edges only.

Proposition 35 (sign relation) *Let $E_O \neq \emptyset$ and $E_O^{\uparrow} = \{(i,j) \in E_O : \mathrm{rel}_S(i) \cap O \neq \mathrm{rel}_S(j) \cap O\}$. (a) Every edge of E_O scores $1/2$ or 1, and $\tilde{s} = \frac{1}{2} + |E_O^{\uparrow}|/(2|E_O|) \geq \frac{1}{2}$. (b) $E_O^{\uparrow} \subseteq E_{\uparrow}$. Hence $\tilde{s} > \frac{1}{2}$ iff $E_O^{\uparrow} \neq \emptyset$, and then some induced edge is an up-edge of G_{S} ($\mathrm{Src}(i) \neq \mathrm{Src}(j)$); if no induced edge is an up-edge, $\tilde{s} = \frac{1}{2}$. (c) $(i,j) \in E_O$ belongs to $E_O^{\uparrow}$ iff some observed w lies in $\mathrm{rel}_S(j) \setminus \mathrm{rel}_S(i)$, i.e. shares an ancestor with j but none with i; in particular an observed other parent k of j with $\mathrm{anc}_S(k) \cap \mathrm{anc}_S(i) = \emptyset$ suffices. (d) With $O = V$, $E_V^{\uparrow} = E_{\uparrow} \setminus \{\text{self-loops}\}$, so $\tilde{s}(G_{\mathrm{S}}, V) - \frac{1}{2}$ and $s_\rho^W - \frac{1}{2}$ have the same sign. (e) The converse of (b) fails: for the collider $a \to c \leftarrow b$ with loops and $O = \{a, c\}$, the only induced edge (a, c) is an up-edge, but $\mathrm{rel}_S(a) \cap O = \mathrm{rel}_S(c) \cap O = \{a, c\}$, so $\tilde{s} = 1/2$.*

Lean: (a)–(d): relObs_subset_of_edge, score_obsEdge, sTilde_eq, sTilde_ge_half, upObs_subset_upEdges, sTilde_gt_half_iff, exists_upEdge_of_sTilde_gt_half, sTilde_eq_half_of_no_up, mem_upObs_of_witness, upObs_univ (file Observed.lean); (e) paper proof only (not machine-checked; written by the AI agents and not yet read by a human).

Proof (a) $i \in \mathrm{anc}_S(j)$ gives $\mathrm{rel}_S(i) \subseteq \mathrm{rel}_S(j)$ (Reisach et al., 2026, Thm. 1), hence the same after intersecting with O; nested sets have equal or strictly ordered sizes. (b) If the observed parts differ, the full sets differ, and by Theorem 5 that is $\mathrm{Src}(i) \neq \mathrm{Src}(j)$. (c) By (a) the observed sets are nested, so they differ iff the larger one has an observed element outside the smaller; $w \in \mathrm{rel}_S(j) \setminus \mathrm{rel}_S(i)$ means $\mathrm{anc}_S(w) \cap \mathrm{anc}_S(j) \neq \emptyset = \mathrm{anc}_S(w) \cap \mathrm{anc}_S(i)$; k itself is such a w. (d) is (b) with $O = V$ and $\mathrm{Src}(i) \neq \mathrm{Src}(j)$ iff $\mathrm{rel}_S(i) \neq \mathrm{rel}_S(j)$ (Theorem 5); $s_\rho^W > \frac{1}{2}$ iff $E_\uparrow \neq \emptyset$, and self-loops are never in $E_\uparrow$. (e) $\mathrm{rel}_S(a) = \{a, c\}$, $\mathrm{rel}_S(c) = \{a, b, c\}$; b is unobserved. □

What is and is not implied. For river subgraphs (in-forests) an induced edge (i, j) counts iff an observed station lies upstream of j on a branch other than i's (in an in-forest $\mathrm{rel}_S(v) = \mathrm{anc}_S(v) \cup \mathrm{desc}_S(v)$ and $\mathrm{desc}_S(j) \subseteq \mathrm{desc}_S(i)$), so a confluence in the induced subgraph whose other branch is observed is sufficient. $\tilde{s}$ is a statement about population lag-0 relatives of a stationary process; it is not s_ρ of the finite unrolled DAG, of the induced subgraph's unrolled DAG, or of the latent projection, and it says nothing about lagged dependence. A deviation of the estimate from $\tilde{s}$ is a failure of estimation (F1–F3) or of lag-0 faithfulness, not evidence against Theorems 1 and 2.

Appendix K Time-Stamped Analysis Plans: Process, Prior Exposure, Timeline, Deviations

Process. The four real-data analysis plans, all analyses and the results were produced by AI coding agents (Section 5.3, AI-use statement) within about six hours of one day; the later plans for B1, B1-NL, B3, the learner extension and the controls (b1y) are listed at the end of this section. Each plan was committed to version control before its study's data were loaded for any dependence statistic in this project. Commit timestamps are set by the committing client; no third-party registry (e.g. OSF) was used. The hosting service's push records are given below for the later plans and can be supplied for the others. The gap between plan and results was short: about 3 min 41 s for plan 3 and 3 min 36 s for plan 4 (each including a run of about two minutes). We therefore call these *internally time-stamped analysis plans*. The commits are in the history of the supplementary repository.

Prior exposure to CausalRivers (disclosure). One day before plan 1, a separate exploratory project in the same repository (a test of contemporaneous-edge discovery on aggregated CausalRivers; commits dated 2026-11-03) had analysed CausalRivers with the same regions (East Germany, Bavaria), at resolutions 6 h, 12 h and 24 h (the studies here use 12 h and 24 h), and with the same data loaders and subgraph sampler that the plan-1 code reuses (with new seeds). It ran correlation, VAR, PCMCI+, VAR-LiNGAM and DYNOTEARS baselines and recorded a hidden-driver diagnostic: at 24 h the VAR(1) residual correlation among sibling stations was 0.51, the same as among unrelated non-adjacent stations in the same subgraph (0.51), and below adjacent pairs (0.65); it concluded that shared regional forcing is large and contemporaneous dependence is not specific to the river edges. This was known to the agents (and recorded in the repository) before plan 1 was written. Plan 1 nevertheless predicted that the relatives sort would orient river edges above chance; it failed.

Because the strong shared dependence had been observed beforehand, the F2 saturation account of that failure (and the F3 account of study 2) is *explanatory*, not a prediction made blind to the data. The Causal Chambers data and the E1/E2 held-out sets had not been analysed before their plans.

Times below are commit timestamps relative to the first plan commit (T_0).

time	event
$T_0 + 0{:}00$	plan 1 committed: targets T1–T3, predictions P1–P3, stop condition
$T_0 + 0{:}14$	theory of Sections 3.2 and 4.2, enumeration checks, operationalisation of P1–P3 (before data)
$T_0 + 0{:}17$	study 1 results (P2, P3 fail; stop condition met)
$T_0 + 1{:}05$	plan 2 committed: F1–F3 targets, Lean target, predictions P4–P6, stop condition
$T_0 + 1{:}18$–$1{:}31$	Lean formalisation of Section 3.2 and algebraic parts of Section 4.2
$T_0 + 1{:}38$	finite-sample theory, simulations, operationalisation of P4–P6 (before data)
$T_0 + 1{:}44$	study 2 results (P4, P6 fail; stop condition met)
$T_0 + 4{:}58$	plan 3 committed: candidate scoping, graph-only predictions, P7–P9 (before any dependence statistic)
$T_0 + 5{:}02$	study 3 results (P7, P8 pass; P9 fails; replication passes)
$T_0 + 5{:}14$–$5{:}24$	Lean formalisation of the finite-sample lemmas; errata
$T_0 + 5{:}45$	manuscript draft committed (with placeholders for E1/E2)
$T_0 + 5{:}50$	plan 4 committed: E1 (test of the theorem against eight naive predictors), E2 (lagged relatives), theory addendum A1–A3 and B1–B3, graph-only predictions, design check; development runs on the study-3 data only; no dependence statistic of any held-out dataset computed
$T_0 + 5{:}54$	E1/E2 results committed as planned (E1 passes 8/8 on the primary graph; Q7, Q8, Q10 pass, Q9 fails); out-of-sample F checks; run logs

Deviations (all recorded at the time). (1) Plan 1 said "theorem-predicted s_ρ" without naming the version; the stationary version was fixed in writing before data were loaded (the finite version would predict $> 1/2$ for every subgraph, including chains). (2) Self-loops are excluded from the predicted average because the estimator cannot orient them; this does not change which side of $1/2$ the prediction is on. (3) Only 6 of 120 study-1 subgraphs lacked a confluence; the plan set no minimum size for that group, and its trivial pass is not interpreted. (4) Plan 2 did not say which test defines "dependent" in P4; the effective-n test was fixed as primary before data, and the nominal test fails P4 as well. (5) In study 3, a planned triad-reliability check was not applicable (the measured driver reaches only four mutually related barometers) and is not interpreted. (6) The P9 failure is close to the threshold and its CI contains 0.3; as fixed in the plan, the point estimate decides. (7) Plan 4 (E1/E2): no threshold, unit, seed, K, L, D, graph or node set changed after the results were seen. A stray cache file committed with plan 4 was removed in the results commit (no content

change). The F-check tolerance of ± 0.03 is the one fixed in plan 4; the citris "predicted accuracy" is undefined (no full-set edge of G_{lin} is predicted 1). The constant-1/2 comparison and the weak-component-size comparison are post hoc. The choice of G_{lin} as the primary graph was fixed in plan 4, together with the secondary analysis on the published graph G_D, which is the one on which the theorem fails against the sensor indicator on strictly ordered edges.

Later plans (benchmark experiments and B3). Times in the table are commit times (UTC, set by the committing client). Nine plans were written in all: the four real-data plans, and the plans of B1, B1-NL, B3, the learner extension and the controls (b1y). For three of the nine we also retrieved the hosting service's own push records (its activity log, UTC, 09 Nov): B1-NL plan 04:10:47; B3 plan 04:21:12 and B3 results 04:24:37; learner-extension plan 08:20:31 and results 12:44:27. Each is 2 to 28 seconds after the corresponding commit time. We did not retrieve push records for the B1 plan or the four earlier plans, and for the controls plan the activity feed did not list the pushes when queried, so only client commit times are recorded for it. Push records order pushes, not computations: a plan pushed three minutes before its results, with a run of about two minutes in between, is consistent with the stated process but cannot exclude that something was computed before the plan was written. Only registration with a third party before data access could, and no plan was posted to an external registry.

time (2026)	event
08 Nov 22:31	B1 plan committed (generators, methods, hypotheses H1–H3, decision rules)
08 Nov 22:57	B1 code; deviations D1 (burn-in 50→100 after the validation mixing check) and D2 recorded before the test run
09 Nov 01:26	B1 results (test run redone once after a container restart, D3)
09 Nov 04:10	B1-NL plan committed
09 Nov 04:20	B3 plan committed and pushed before any held-out byte was read
09 Nov 04:24	B3 results
09 Nov 04:34	B1-NL analysis code; validation notes N1, N2
09 Nov 06:49	B1-NL results
09 Nov 08:20	learner-extension plan (Appendix T) committed and pushed before any score
09 Nov 08:34	post hoc re-analysis of B1/B1-NL and detectability oracle (no new data)
09 Nov 12:44	learner-extension results (one run, 08:28–12:41)
10 Nov 05:15	controls plan (b1y) committed (`5c9f942`); exact burn-in computation (Q6) committed with it
10 Nov 06:15	Addendum 1 (environment, added learners, budget; `583ff30`) committed before any scored run
10 Nov 06:20	Q1 results (counterfactual, dose)
10 Nov 06:28	Q4 results (run- and station-level inference, truncation-lag sensitivity)
10 Nov 06:32	Addendum 2 (NOTEARS restricted to $d \leq 30$) before the Q2 test run
10 Nov 07:26	Q2 tuning results and tuned hyper-parameters, before the Q2 test run
10 Nov 07:40	Q5 results (further baselines)
10 Nov 11:14	Q2 results (variance-matched control, tuned learners); Holm table

Appendix L Experiment E1: Discriminating Test Against Naive Predictors

Status and provenance. Specified in plan 4 (Appendix K); the plan was committed before any dependence statistic of any held-out dataset was computed. Done before it: file hashes, column and row counts, regeneration of the `citris_1` input protocol from its published generator (it matches the recorded inputs exactly), graph-only predictions, and development runs on the study-3 data only. All numbers below come from the committed result files of the experiment folder (`results4.json`, `summary4.txt`); none is estimated or rounded beyond the printed digits.

Why `citris_1`. Every published Causal Chambers graph is bipartite except the pressure-control configuration, whose only experiment varies no independent input, so we rejected it. The light-tunnel `citris_1` series (100,000 steps about 1.8 s apart) is generated by a time-series structural model among the actuators: red is a reflected random walk, green moves halfway towards red and blue a quarter of the way towards

green (each plus a uniform step), pol_1 is a random walk and pol_2 moves towards or away from pol_1 according to sign(blue − red). With the sensors this is a chain of length 3, red → green → blue → sensors. The secondary data are snapshot (i.i.d.) structural models among the actuators (`buchholz_1_obs`; `scm_{1,2,4,5}_reference`, including a chain and a collider) and the wind-tunnel regime-jump runs (bipartite; disagreements arise only from latent-node ties).

Linear-zero lemma (the refinement). Write $\mathcal{C}$ for the colour block (red, green, blue, ir_1, vis_1, ir_2, vis_2, ir_3, vis_3, current) and $\mathcal{Q}$ for the polariser block (pol_1, pol_2, angle_1, angle_2). *B1 (proved for the actuators).* Conditionally on the flag and the colour paths, the law of the polariser paths is invariant under negation, because every update (uniform draw on a symmetric interval, move towards or away with a sign fixed by the colours, reflection at ±90, rounding to 0.1) commutes with negation. *B2 (assumed).* The first-two-layer sensors depend only on the colours and independent noise; ir_3 and vis_3 depend on the polarisers only through an even function of $\theta_1 - \theta_2$ (Malus' law with aligned zero marks); each angle sensor is affine in its polariser plus independent noise. This is a physical assumption that the design documents do not let us check. *B3.* Under B1 and B2 every variable of $\mathcal{C}$ has zero cross-covariance, at every pair of times, with every variable of $\mathcal{Q}$ (a product of an even and an odd function of the polariser path has conditional mean zero). The six real edges red → pol_2, blue → pol_2 and pol_k → ir_3, vis_3 therefore carry no linear dependence. The relatives that a correlation-based estimator targets are those of the *linear-dependence graph* G_{lin} (the design graph G_D without these six edges), which has two weak components, each with a single source (red; pol_1). By Corollary 6, every edge inside a block has predicted value 1/2 for every node set; the six cross-block edges are compared by block size: red → pol_2 and blue → pol_2 are predicted 0 (anti-sorted) and pol_k → ir_3, vis_3 are predicted 1 for the full set. Faithfulness inside the blocks is assumed, as in Proposition 22. A zero-mark offset or a non-affine angle sensor would break B2 and move the prediction towards G_D.

Metric and tests. For each edge instance (the full node set plus 20 subgraphs per unit) the observed value is $o \in \{0, \frac{1}{2}, 1\}$ from the estimated counts, and the loss is $|o - p|$ for the predicted value p. For each competitor take the instances where it disagrees with the theorem; Δ = loss(theory) − loss(competitor), pooled, with a unit-level exact sign-flip test (1,024 flips for 10 units) and Holm correction over eight competitors, and a unit-bootstrap CI. E1a (prerequisites): non-relatives declared < 0.5 and relatives detected ≥ 0.8. E1b: $\Delta \leq -0.10$ and Holm $p \leq 0.05$ for all eight. E1c: the same restricted to instances whose own pair is declared dependent ($\Delta < 0$). Separation was to be claimed iff E1a, E1b and E1c all pass on `citris_1`. Stop conditions: K1 (E1a fails: no claim), K2 (a competitor is significantly better: no claim against it), K3 (E1b passes and E1c fails: attributed to power-induced ties). None was triggered on the primary analysis.

Observed values given the theorem's prediction: predicted 1/2 (1,365 instances) observed exactly 1/2 in 99.2%; predicted 1 (187 instances, polariser → ir_3, vis_3) observed 1 in 99.5%; predicted 0 (110 instances, red and blue → pol_2) observed 0 in 99.1%. The detected-edge subset has 1,338 instances, all observed exactly 1/2. Prerequisites: non-relatives declared 0.000 [0, 0] (40 pairs per unit), relatives detected 0.969 [0.945, 0.992] (51 pairs per unit).

Strict subset and the power safeguard. On the 297 instances where the theorem predicts 0 or 1 it beats all eight ($\Delta = -0.539$ for the sensor indicator, -0.948 for ancestors and depth on the full graph, -0.991 for "every edge 1"; Holm $p = 0.0078$, 10/10 units). There, an estimator without power would favour the competitors (the "all counts tied" baseline is $+0.443$ to 0), so this part cannot come from power-induced ties.

Published graph G_D (secondary). Applying the theorem to the unmodified design graph, relatives detection is 0.659 [0.643, 0.675] (75 relative pairs per unit) and non-relatives declared 0.000. Overall the theorem still beats all eight ($\Delta = -0.323$ for the sensor indicator; -0.424 for ancestors, depth and "every edge 1"; -0.420 and -0.417 for in-degree; 10/10 units, Holm $p = 0.0078$), and its exact-match rate is 0.714. But on the 442 edges where G_D predicts strict order (1), the theorem matches 0.222 and the sensor indicator 0.319; the sensor indicator is significantly better ($\Delta = +0.453$ [0.410, 0.492], theorem better in 0/10 units, reverse Holm $p = 0.002$). A constant-1/2 rule (post hoc) matches 0.815 of all instances, against 0.714 for the theorem on G_D and 0.992 on G_{lin}; mean $|o - 1/2| = 0.092$ against 0.166 for the theorem on G_D. The Spearman statistic is 0.341 (theorem on G_D) against 0.408 (sensor indicator). This is why the claim is restricted to the linearly dependent relatives.

Lagged estimator on `citris_1` (secondary). With $L = 5$ and G_{lin}: detection 0.976 [0.953, 0.992], non-relatives declared 0.000, 8/8 passed ($\Delta = -0.505$ to -0.535), theorem exact match 0.998, Spearman 0.996.

On the snapshot sets, relatives detected: 0.933 (`buchholz_1_obs`), 1.000 (`scm_1`), 0.965 (`scm_2`), 1.000 (`scm_4`, `scm_5`); non-relatives declared 0.000 wherever non-relative pairs exist (`scm_4` has none). The chain (`scm_4`) gives 1/2 on all 2,267 instances (observed exactly 1/2 every time); the collider gives 1 on edges into the collider and its sensors and 1/2 elsewhere (all 2,238 instances exact). Pooled, the theorem's strict predictions number 2,768 (observed value 1 on 2,403, 1/2 on 349, 0 on 16: exact 0.868); on them only the sensor indicator disagrees often (390 instances, $\Delta = -0.5$, 35/35 units) and the other competitors coincide with the theorem. The pooled Spearman is 0.913 for the theorem and -0.138 for the sensor indicator. Post hoc, a constant-1/2 rule matches 0.688 of the pooled instances against 0.923.

Wind-tunnel regimes (tertiary, 64 units, no verdict). The same test passes against all eight in each configuration. $D = 10$, $L = 0$: non-relatives declared 0.014, relatives detected 0.592, $\Delta = -0.203$ $[-0.227, -0.179]$ in 63/64 units, exact match 0.539 against 0.432; detected-edge subset $\Delta = -0.169$. $D = 1$, $L = 30$: 0.003, 0.398, $\Delta = -0.225$ $[-0.250, -0.197]$ in 62/64 units, exact 0.443 against 0.353; detected-edge subset $\Delta = -0.134$. $D = 1$, $L = 0$: $\Delta = -0.218$ (64/64 units), exact 0.411 against 0.326; $D = 10$, $L = 3$: $\Delta = -0.218$, exact 0.533 against 0.411. Detection of 0.33–0.59 means that part of the advantage is power-induced, as plan 4 stated in advance. Post hoc, the constant-1/2 rule has mean absolute error 0.311 against 0.363 for the theorem at $D = 1$, $L = 30$ (and 0.313 against 0.294 at $D = 10$, $L = 0$), and matches 0.378 (0.375) of the instances against 0.443 (0.539).

A trivial rule that E1 did not test (post hoc). Each weak component of G_{lin} has a single source (red; pol_1). By the proof of Corollary 6, every node of such a component is reached from that source, so its relatives are the whole component,

and for every node set the predicted counts are the numbers of observed nodes in the endpoints' weak components. The theorem's prediction therefore coincides on all 1,662 instances with the rule "1/2 iff both endpoints lie in the same weak component of G_{lin}, otherwise compare the observed sizes of the two components". That rule was not among the eight competitors, so E1 cannot separate the theory from it; the only data in this appendix where the source structure matters beyond component membership are the snapshot sets with several sources per component, notably the collider `scm_5`.

What E1 does and does not show. It shows that, on one multi-layer light-tunnel series, the estimated relatives counts are tied along the whole chain red $\to$ green $\to$ blue $\to$ sensors and between the polariser block and the angle sensors, and strictly ordered (including against the causal direction) on red, blue $\to \text{pol}_2$ — and that the source-component rule of Theorem 5 and Corollary 6 predicts this while eight naive rules do not. It does not show that the theorem predicts relatives under the published graph; the faithfulness of the graph to linear dependence fails here by design. Most of the separation is in the ties, which naive rules cannot express, and on G_{lin} the prediction equals the weak-component-size rule. The evidence is one device in a favourable regime, and the wind tunnel, where detection is limited by F1, gives weaker support.

Appendix M Experiment E2: Lagged Relatives

Theory (plan 4, addendum A). For $L \in \{0, 1, \dots\}$ let the lag-L relatives of u be u and every v with $\gamma_{uv}(h) \neq 0$ for some $|h| \leq L$. *A1 (proved).* Under the independent-innovations assumption behind Proposition 22, $\text{rel}_0 \subseteq \text{rel}_L \subseteq \text{rel}_\infty \subseteq \text{rel}_S$; for a VAR model all are equal for almost every parameter value; and Theorem 5, Corollary 6 hold verbatim with "relatives" read as rel_L under lag-L faithfulness. Lag-0 faithfulness implies lag-L faithfulness; the converse fails (a white-noise source s with an edge $s \to j$ and no self-loop at s has $\gamma_{sj}(0) = 0 \neq \gamma_{js}(1)$). *Estimator.* This is a Bonferroni max-over-lags variant of the classical cross-correlation tests of independence between two series (Haugh, 1976 cross-correlates prewhitened residuals and uses a portmanteau statistic; we use effective-n Fisher-z statistics without prewhitening). Declare u, v relatives iff $\max_{|h|\leq L} |\operatorname{atanh} \hat{r}_{uv}(h)| \sqrt{\hat{n}_{\text{eff},uv}(h) - 3} > z_c(L)$, with $\hat{n}_{\text{eff},uv}(h) = (n - |h|)/\hat{B}_{uv}$ and $z_c(L) = \Phi^{-1}(1 - \alpha/(2m(2L+1)))$, a Bonferroni correction over the m pairs and the $2L+1$ lags; $L = 0$ is exactly the study-3 test. *A2 (proved for finite n).* For independent stationary Gaussian a, b the null variance of the lag-h cross-covariance has the same bound at every h, so one Bartlett factor serves all lags. *A3 (deterministic part proved, constants sketched).* The recovery bound of F1 holds with $\rho_{\min}$ replaced by $\rho^*_{\min,L} = \min_{\text{relative pairs}} \max_{|h|\leq L} |\rho_{ab}(h)| \geq \rho_{\min}$ (the gain) and a factor $2L + 1$ in the failure probability and the Bonferroni count (the cost): for 105 pairs at $\alpha = 0.01$, z_c rises from 3.90 ($L = 0$) to 4.35 ($L = 3$) and 4.80 ($L = 30$). A lagged correlation must therefore exceed the lag-0 correlation by about 1.12 ($L = 3$) or 1.23 ($L = 30$) times to help.

Design. Held-out data: the wind-tunnel regime-jump runs (`regime_jumps_single`, `regime_jumps_multi`), $2 \times 32 = 64$ regimes of 10,000 steps, in the same archive as the study-3 data but never analysed in study 3. Load_in, load_out, pot_1, pot_2 follow

independent mirrored random walks within a regime; the hatch and sensor parameters are constant within a regime. The first 100 rows of each regime are dropped. Graph: 20 edges without the hatch (19 predicted 1, $\text{pot}_1 \to \text{signal}_1$ predicted 1/2); 54 non-relative and 51 relative pairs; 1,280 subgraphs (10 of size 5 and 10 of size 8 per unit). Primary: $D = 1$ (first differences), $L = 30$ (about 9 s) against $L = 0$; secondary: block means with $D = 10$ and $L = 3$ against $L = 0$. Success required Q7–Q10 all to pass at the primary configuration. Stop conditions: Q7 fails (saturation) or Q10 fails (lags do not help detection, and the lagged estimator is not recommended). Plan 4 states that E2 is not a re-analysis of P9 and that P9's verdict is final.

Detection of the 51 relative pairs: $D = 1$: 0.328 [0.312, 0.345] at lag 0 and 0.398 [0.383, 0.412] with $L = 30$; $D = 10$: 0.592 [0.584, 0.599] and 0.585 [0.576, 0.592] with $L = 3$ (lags reduce detection: the Bonferroni cost exceeds the gain). On `citris_1` with $L = 5$ against $L = 0$: 0.976 against 0.969, difference +0.008 [0, 0.024], so Q10 fails there; detection is already nearly saturated at lag 0. **Verdict.** Q9 fails at the primary configuration, so E2 does not succeed; no stop condition is triggered (Q7 and Q10 pass). Lags raise detection at fine resolution without adding false positives, as A1 and A3 predict. The study-3 estimator (lag 0) rerun on these held-out regimes fails the graded criterion again (0.202 at $D = 10$, 0.120 at $D = 1$): the P9-type failure repeats on held-out data. P9's original verdict (0.267, failed) stands; E2 is not presented as rescuing it. The development check of plan 4 showed that, on the study-3 data with the same estimator, drawing different subgraph seeds moves the P9 statistic from 0.267 to 0.338. We report this as variability of the statistic (about 0.07 from the subgraph draw alone); it is a development run, not a verdict, and it does not change the P9 outcome.

Lags recover the pressure edges at fine resolution (load_in $\to$ upwind pressure 0.562 $\to$ 0.984, load_out $\to$ upwind pressure 0.016 $\to$ 0.812) but lose the pot edges to the larger threshold (0.359 $\to$ 0.156 and 0.266 $\to$ 0.109); every microphone edge, $\text{pot}_2 \to \text{signal}_2$ and the cross-current edges stay undetected. Even at its best ($D = 10$, lag 0) wind-tunnel detection (0.592) is far below that of the light tunnel, so lags are not the main cause of the F1 shortfall.

All six checks pass. With lags the observed detection sits at or slightly below $\bar{p}_{\max}$, consistent with strongly correlated lag statistics; the independence formula overstates both detection and false positives, as the theory said it would. For `citris_1` the predicted accuracy is undefined because no full-set edge of G_{lin} is predicted 1 inside a block.

Appendix N Summary Graphs without Self-Loops

Here $G_S = (V, E_S)$ is an arbitrary directed graph with $|E_S| \geq 1$; no loop assumption is made. B_n is the boolean n-th power of the adjacency matrix ($B_0 = I$), $\Lambda(u, v) = \{n : B_n[u, v] = 1\}$ the set of lengths of walks from u to v, and u is *fresh* if it has no incoming edge in E_S (a self-loop is an incoming edge). Status: Theorem 37 and Corollary 38 are machine-checked; the rest of this section is proved on paper and checked by enumeration, as marked.

Lemma 36 (ancestors and roots without loops) *$u(a) \in \mathrm{anc}(v(t))$ iff $a \le t$ and $t-a \in \Lambda(u,v)$. A node $u(a)$ is a root of G_T iff $a = 1$ or u is fresh.*

Lean: unrolled_reach_iff (already free of the loop hypothesis).

Proof Every edge of G_T raises time by one, so paths from $u(a)$ to $v(t)$ are the lifts of walks of length exactly $t-a$, and every walk lifts. Without loops one cannot pad. A node $u(a)$ with $a \ge 2$ has a parent iff some edge (w,u) exists. □

Theorem 37 (edge value without loop assumption) *Let $(i,j) \in E_S$ and $e_t = (i(t), j(t{+}1))$, $1 \le t < T$. Then $s_\rho(e_t) \in \{\frac{1}{2}, 1\}$ independently of T, and $s_\rho(e_t) = 1$ iff $\mathrm{Strict}_{ij}(t)$ holds, where $\mathrm{Strict}_{ij}(t)$ is*

$$\exists u:\ t \in \Lambda(u,j),\ t{-}1 \notin \Lambda(u,i) \quad \text{or} \quad \exists \text{ fresh } u,\ \exists n \in \{1,\dots,t\}:\ n \in \Lambda(u,j),\ n{-}1 \notin \Lambda(u,i).$$

Lean: isRel_unrolled_iff_NL, reduce_ancestor, relSet_unrolled_eq_iff_NL, edge_score_NL.

Proof As in Theorem 1, $\mathrm{rel}(i(t)) \subseteq \mathrm{rel}(j(t{+}1))$, and the score is 1 iff $\mathrm{roots}(j(t{+}1)) \setminus \mathrm{roots}(i(t)) \neq \emptyset$ (Reisach et al., 2026, Thm. 1, Lemma 1). By Lemma 36 an element of the difference is $u(a)$ with ($a = 1$ or u fresh), $n = t+1-a \in \Lambda(u,j)$ and $u(a) \notin \mathrm{roots}(i(t))$. The case $a = t+1$ would make j fresh, impossible since j has the parent i; so $a \le t$ and, by the lemma again, $u(a)$ is not a root of $i(t)$'s ancestry iff $n-1 = t-a \notin \Lambda(u,i)$. The branch $a = 1$ is $n = t$ with no condition on u; the other is a fresh u with $1 \le n \le t$. (The Lean proof avoids Lemma 1: the extra relative is $u(0)$ in the first case and a node of a fresh u in the second, and reduce_ancestor pushes any common ancestor back along parents to slice 0 or to a fresh node.) □

Corollary 38 *For every summary graph and $T \ge 2$, $s_\rho(G_T) = \frac{1}{2} + \frac{1}{2(T-1)|E_S|}\sum_{(i,j)\in E_S} N_T(i,j)$ with $N_T(i,j) = \#\{t < T : \mathrm{Strict}_{ij}(t)\}$; in particular $s_\rho(G_T) \ge 1/2$. With a loop at every node no node is fresh, $\Lambda(u,j) = [d(u,j),\infty)$, and $\mathrm{Strict}_{ij}(t)$ becomes $\exists u : d(u,j) \le t \le d(u,i)$, i.e. Theorem 1.*

Lean: sRho_unrolled_eq_NL, sRho_unrolled_ge_half_NL.

Exhaustive check of Theorem 37 against brute-force relatives: $p = 2,3$ ($T = 2,\dots,9$), all 65,535 graphs with $p = 4$ ($T = 3,5,8$), and random graphs with $p = 5,6$; 0 mismatches in 120, 3,577, 196,605, 8,928 and 4,494 (graph, T) cases (enumerated, not part of the Lean proof).

Proposition 39 (limit exists and is rational; paper proof, enumerated) *(B_n) is eventually periodic: there are $n_0, \pi \ge 1$ with $B_{n+\pi} = B_n$ for $n \ge n_0$. Hence $L = \lim_T s_\rho(G_T) = \frac{1}{2} + \frac{1}{2|E_S|}\sum_e \delta_e$ with $\delta_e = 1$ if some fresh u has $n \ge 1$, $n \in \Lambda(u,j)$, $n-1 \notin \Lambda(u,i)$, and otherwise δ_e the density, over one period of $t > n_0$, of t with some u, $B_t[u,j] = 1$, $B_{t-1}[u,i] = 0$; and $|s_\rho(G_T) - L| \le (n_0 + \pi + 1)/(2(T-1))$.*

Proof B_{n+1} is a function of B_n and there are finitely many boolean matrices. A fresh witness is monotone in t and has its first n at most $n_0 + 1 + \pi$ (the set $\Lambda(u,j) \setminus (\Lambda(u,i)+1)$ is eventually periodic), so such an edge is strict for all $t \geq n_0 + 1 + \pi$. Without a witness, $\mathrm{Strict}_{ij}(t)$ depends for $t > n_0$ only on $t \bmod \pi$. Counting strict $t < T$ gives the error bound; average over edges. □

Checked on 1,115 graphs (all $15 + 511$ graphs with $p = 2, 3$, 390 random with $p = 4$, 199 random with $p = 5, 6$): the exact value at $T = 600$ is within the bound and $(L - \frac{1}{2})\, 2|E_{\mathrm{S}}|\, \pi$ is an integer in every case.

Proposition 40 (edges inside an SCC; paper proof, enumerated) *Let $(i,j) \in E_{\mathrm{S}}$ with j reaching i and no fresh witness. Then $\mathrm{Strict}_{ij}(t)$ fails for every $t \geq n_0 + 1$.*

Proof Fix u; put $f(n) = B_n[u,j]$, $h(n) = B_{n-1}[u,i]$. The edge gives $h(n) \leq f(n)$. If m is the length of a walk $j \to i$ and $c = m + 1$, then $f(n) \leq h(n+c)$. Take $k \geq 1$ with kc a multiple of π. For $n \geq n_0 + 1$, $h(n) \leq f(n) \leq h(n+c) \leq \cdots \leq h(n+kc) = h(n)$, so $f(n) = h(n)$ for every u. □

Proposition 41 (crossing edge into a cyclic node) *If (i,j) is crossing and j lies on a cycle with period g (gcd of the closed-walk lengths at j), then $\mathrm{Strict}_{ij}(t)$ holds for every $t \in \Lambda(j,j)$, which contains all large multiples of g, so $\delta_e \geq 1/g$, and $\delta_e = 1$ if $g = 1$. A fresh parent $u \neq i$ of j makes every layer strict.*
Lean: *strictNL_of_closed_walk, strictNL_of_fresh_parent, sRho_unrolled_eq_one_NL.*

Proof Take $u = j$: $t \in \Lambda(j,j)$ and $t - 1 \notin \Lambda(j,i) = \emptyset$. Closed-walk lengths at j form a numerical semigroup in $g\mathbb{Z}$, which contains all large multiples of g. □

Corollary 42 (loop-free versions of the loop-dependent statements) *(a) If every weakly connected component of G_{S} is strongly connected (any loop pattern, any periods), then $L = 1/2$ and $s_\rho(G_T) - \frac{1}{2} \leq n_0/(2(T-1))$ (paper proof from Proposition 40; Reisach et al.'s Theorem 4 without loops). (b) The converse and the closed form are false without loops (Table N6). (c) If every crossing edge has a head on a cycle, then $L > 1/2$ iff $E_\times \neq \emptyset$.*

For an acyclic loop-free graph, $L(u,v)$ is finite, $B_n = 0$ for $n \geq p$, and $L = \frac{1}{2} + \#\{e : \text{fresh witness}\}/(2|E_{\mathrm{S}}|)$ (paper proof); for a loop-free in-forest an edge (i,j) has a fresh witness iff j has another parent, so $L = \frac{1}{2} + (m - n_1)/(2m)$ with $m = |E_{\mathrm{S}}|$ and n_1 the number of nodes with one parent, which equals the stationary value below. On 362 random acyclic loop-free graphs $L = S$ in all cases, and the in-forest formula holds on 298 in-forests.

Theorem 43 (stationary value without loops; (a), (c) proved, (b) sketched) *Let G_{S} be arbitrary, $G_{\mathbb{Z}}$ the infinite-past unrolled DAG on $V \times \mathbb{Z}$, and $\Delta(v,w) = \{m - n : \exists u,\ n \in \Lambda(u,v),\ m \in \Lambda(u,w)\}$. Then $\mathrm{rel}_{G_{\mathbb{Z}}}(v(t)) = \{w(t+d) : d \in \Delta(v,w)\}$ (proved), the edge*

*$(i(t), j(t{+}1))$ is strict iff $\Delta(j,w){+}1 \neq \Delta(i,w)$ for some w (proved; always $\Delta(i,w) \subseteq \Delta(j,w){+}1$), and $S^{\mathbb{Z}} = \frac{1}{2} + |E^{\mathbb{Z}}_{\uparrow}|/(2|E_{\mathrm{S}}|)$. (a) Outside $|d| < n_0 + \pi$, $\Delta(v,w)$ is π-periodic, so $E^{\mathbb{Z}}_{\uparrow}$ is detected with witnesses of $|d| \leq n_0 + 2\pi$ (proved). (b) In a finite window after a burn-in $B \geq 2n_0 + 3\pi + 1$ the windowed value equals $S^{\mathbb{Z}} + O(K_0/T)$, $K_0 = n_0 + 2\pi$ (*sketched*). (c) $S^{\mathbb{Z}} = 1/2$ whenever every weak component is strongly connected, whatever its periods (a node colouring* lab *with $\Delta(v,w) = \mathrm{lab}(w) - \mathrm{lab}(v) + g\mathbb{Z}$) (proved). With loops everywhere Δ is $\mathbb{Z}$ or empty and this is Theorem 5.*

Brute-force relatives in G_{B+T} ($B = 40$, $T = 90$, edges at distance ≥ 30 from the window ends) agree with $E^{\mathbb{Z}}_{\uparrow}$ on every edge of all $15 + 511$ graphs with $p = 2, 3$ and 248 random graphs with $p = 4$ (enumerated). Not formalised in Lean: periodicity, Proposition 40, Corollary 42(a), Theorem 43.

Appendix O Lagged Unrollings

Theorems 1, 2 and 5 and their Lean proofs are for VAR(1)-shaped unrollings: every summary edge acts at lag one. Here every non-loop summary edge $e = (i, j)$ carries a lag $\tau_e \in \{1, \ldots, \tau_{\max}\}$ (a pair may carry several lags, each a separate edge), self-loops act at lag one, and G_T has the edges $(i(t), j(t{+}\tau_e))$, $1 \leq t \leq T - \tau_e$. Let $d_\tau(u,v)$ be the smallest total lag of a path from u to v (0 if $u = v$, ∞ if none); $d_\tau(u,v) \leq d(u,v)\,\tau_{\max} \leq \Delta\tau_{\max}$ when finite. Status: paper proof, not machine-checked and not enumerated. Contemporaneous (lag-0) edges are excluded: they break the time-monotonicity used below.

Proposition 44 (lag-weighted version) *(a) $u(s) \in \mathrm{anc}(v(t))$ iff $d_\tau(u,v) \leq t - s$. (b) The copy $(i(t), j(t{+}\tau_e))$ scores 1 iff some u has $d_\tau(u,j) - \tau_e + 1 \leq t \leq d_\tau(u,i)$, and 1/2 otherwise. (c) Crossing edges score 1 in every copy; an edge inside an SCC has at most $\Delta\tau_{\max}$ strict copies, so $\lim_{T\to\infty} s_\rho(G_T) = \frac{1}{2} + |E_\times|/(2|E_{\mathrm{S}}|)$ with $|s_\rho(G_T) - L| = O(\Delta\tau_{\max}/T)$. (d) With a latent burn-in $B \geq \Delta\tau_{\max}$ (or an infinite past), Theorem 5 holds verbatim.*

Proof (a) Every edge raises time by its lag, so a path from $u(s)$ to $v(t)$ projects to a walk of total lag $t - s$; conversely a path of total lag $d_\tau(u,v)$ can be padded by lag-one self-loops at u. (b) As in Lemma 18, the roots of $v(t)$ are $\{u(1) : d_\tau(u,v) \leq t - 1\}$ (every $u(s)$, $s \geq 2$, has the parent $u(s-1)$), and $i(t)$ is a parent of $j(t{+}\tau_e)$, so the score is 1 iff $j(t{+}\tau_e)$ has a root $u(1)$ with $d_\tau(u,j) \leq t + \tau_e - 1$ that $i(t)$ lacks, $d_\tau(u,i) > t - 1$. (c) For a crossing edge take $u = j$: $d_\tau(j,j) = 0 \leq t + \tau_e - 1$ and $d_\tau(j,i) = \infty$. Inside an SCC, $d_\tau(u,i) < \infty$ iff $d_\tau(u,j) < \infty$, so a strict copy needs $t \leq d_\tau(u,i) \leq \Delta\tau_{\max}$. Each edge has $T - \tau_e$ copies; average. (d) By (a), $w(s) \in \mathrm{rel}(v(t))$ iff some u has $d_\tau(u,v) \leq t - 1$ and $d_\tau(u,w) \leq s - 1$; for window nodes $t - 1, s - 1 \geq B \geq \Delta\tau_{\max}$ this is $\mathrm{anc}_S(v) \cap \mathrm{anc}_S(w) \neq \emptyset$, and the proof of Theorem 5 applies. □

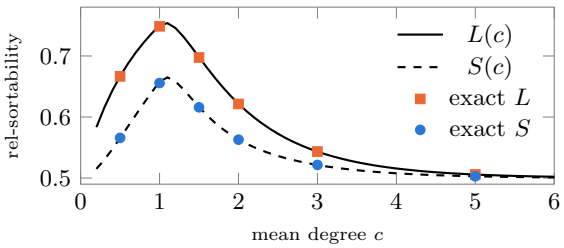


Fig. P1 Proposition 45 (lines) against exact computation on seeded graphs ($p = 2 \cdot 10^4$ to 10^5; points, means over seeds). The standard deviations over seeds are at most 0.0022 (Appendix P).

Appendix P Random Summary Graphs: Derivations and Checks

All graphs here are a non-loop random graph plus a self-loop at every node, so $|E_{\mathrm{S}}| = p + m$ with m the number of non-loop edges, and for acyclic graphs $|E_{\times}| = m$.

Directed ER: statement and figure. Let G_{S} be the directed Erdős–Rényi graph $G(p, c/p)$ plus a self-loop at every node, and x the largest root of $x = 1 - e^{-cx}$ ($x = 0$ for $c \le 1$). By Theorems 2 and 5 only $|E_{\times}|$ and $|E_{\uparrow}|$ matter.

Proposition 45 (ER limits; proved on paper) *Let $L(G_{\mathrm{S}}) = \lim_{T\to\infty} s_\rho(G_T)$ at fixed p and $S(G_{\mathrm{S}}) = s_\rho^W$. In probability and in L^1 as $p \to \infty$, $L(G_{\mathrm{S}}) \to \frac{1}{2} + \frac{c(1-x^2)}{2(1+c)}$ and $S(G_{\mathrm{S}}) \to \frac{1}{2} + \frac{c\,F_\uparrow}{2(1+c)}$ with $F_\uparrow = (1-x)(1-e^{-c}) + x(1-x)\big(1-e^{-c(1-x)d}\big)$, where d is the smallest fixed point in $[0,1]$ of $d = 1 + e^{-c} - e^{-c(1-x)d}$ ($d = 1$ for $c \le 1$).*

Lean: only the deterministic witness lemmas for $E_\uparrow$ (upEdge_witness, upEdges_subset_confluenceNL); the probabilistic proof (Appendix P.1) uses textbook Poisson-coupling, Chernoff and Galton–Watson facts (van der Hofstad, 2016). Not proved: rates, the window $c = 1 + \lambda p^{-1/3}$, a.s. convergence, fixed T as p grows.

The first formula says that the fraction of non-loop edges inside an SCC tends to x^2 (giant strongly connected component, Dorogovtsev, Mendes, & Samukhin, 2001; Dovgal, de Panafieu, Ralaivaosaona, Rasendrahasina, & Wagner, 2020). $L(1) = 3/4$, L peaks at 0.754 near $c = 1.07$ and then decays like e^{-c}; S peaks at 0.665 near $c = 1.09$ (Figure P1). Random summary graphs are thus rel-sortable only near or below criticality, as the giant SCC removes the signal. For in-forests (rivers) the identities are exact (Appendix P).

In-forests (proved identities). Let the non-loop part be an in-forest with m edges and n_1 nodes with exactly one parent. By Corollaries 3 and 6, $L = \frac{1}{2} + \frac{m}{2(m+p)}$, $S = \frac{1}{2} + \frac{m-n_1}{2(m+p)}$, $L - S = \frac{n_1}{2(m+p)}$ (the last checked to machine precision on 18 random forests). For a sequence of in-forests with parent-count distribution π and mean μ: $L \to \frac{1}{2} + \frac{\mu}{2(1+\mu)}$ and $S \to \frac{1}{2} + \frac{\mu-\pi_1}{2(1+\mu)}$. Exact computation at $p = 20{,}000$ (6 seeds each): uniform recursive in-tree ($\pi_1 = 1/4$, $S \to 11/16 = 0.6875$, measured 0.68739 ± 0.00026); preferential attachment (measured $\pi_1 = 0.1656$, $S = 0.70859 \pm 0.00028$); uniform random labelled tree rooted and oriented to the root (Poisson(1) parents,

$\pi_1 \to e^{-1}$, $S \to \frac{1}{2} + (1 - e^{-1})/4 = 0.65803$, measured 0.65783 ± 0.00020). In all three $L = 0.74999$.

Directed ER: summary. $G(p, c/p)$, $m/p \to c$. The finite limit needs the fraction of non-loop edges inside an SCC ($\to x^2$, the two endpoints lie in the giant SCC iff the ancestor tree of i and the descendant tree of j are infinite); the stationary value needs the fraction of up-edges ($\to F_\uparrow$). Both are proved in Appendix P.1 below. Large c: $1 - x \sim e^{-c}$, $L - \frac{1}{2} \sim e^{-c}$, $S - \frac{1}{2} \sim e^{-c}/2$.

Layered DAG (sketch; L proved given edge concentration). M layers of width w, each node of layer $l + 1$ receiving an edge from each node of layer l with probability c/w; as $w \to \infty$ with $\mu = c(M - 1)/M$, $L \to \frac{1}{2} + \frac{\mu}{2(1+\mu)}$ and $S \to \frac{1}{2} + \frac{\mu(1-e^{-c})}{2(1+\mu)}$. For fixed w and deep M the stationary value falls below this (observed 0.7824 against 0.7878 at $w = 50$, $M = 400$, $c = 2$); this correction is not derived.

At the critical point $c = 1$, S converges slowly in p: 0.65574 ± 0.00126 ($p = 2 \cdot 10^4$, 6 seeds), 0.65784 ± 0.00043 ($2 \cdot 10^5$, 6 seeds) and 0.65796 (10^6, 1 seed) against 0.65803, like L (0.74857, 0.74972, 0.74989 against 0.75).

L at $c = 1$ converges as p grows: 0.7466 ($p = 2000$), 0.7486 ($2 \cdot 10^4$), 0.74976 ($2 \cdot 10^5$), 0.74992 (10^6); at $c = 2$, 0.62162 ± 0.00023 at $p = 10^6$ against 0.62170. The giant-SCC fraction matches x^2 (0.6353 against 0.6349 at $c = 2$). The second fixed point matters: with $d = 1$ the formula would give $S = 0.6312$, 0.5766, 0.5246 at $c = 1.5, 2, 3$, outside the error bars, whereas the fixed points $d = 0.485$, 0.221, 0.061 give 0.6163, 0.5632, 0.5214. Further exact rows: $c = 1.5$, $p = 10^5$, 32 graphs, $L = 0.69819 \pm 0.00024$ (formula 0.69810); $c = 2$, $p = 3.2 \cdot 10^5$, $L = 0.62170 \pm 0.00012$ (0.62170). The Barak–Erdős phase-transition paper cited by Reisach et al. (2026) concerns monotone-path reachability at the dense scale $\log n/n$ and is not used here.

P.1 Directed Erdős–Rényi summary graphs: full proof of Proposition 45

Status. Everything in this subsection is proved on paper. The only inputs used without proof are textbook facts (F1)–(F3) below, checked against the author's online version of van der Hofstad (2016) (theorem numbers are those of that version and may differ in print). Lemma 49 is the standard local-coupling argument for sparse random graphs and is written concisely. The deterministic witness lemmas are machine-checked; no probabilistic step is. No human probabilist has yet checked this subsection (see the AI-use statement).

Limits and their order. For a fixed summary graph, Theorem 2 gives $\lim_{T\to\infty} s_\rho(G_T) = L(G_S) = \frac{1}{2} + |E_\times|/(2|E_S|)$, and Theorem 5 gives $s_\rho^W = \frac{1}{2} + |E_\uparrow|/(2|E_S|)$ for every T. For $G_S \sim D(p, c/p)$ we take the limit $T \to \infty$ at fixed p first (this defines $L(G_S)$) and then let $p \to \infty$; for the stationary value no T-limit is needed. The regime of fixed T with $p \to \infty$ is not covered (the bound of Theorem 2 is vacuous for $T \le p$).

Model and notation. $D(p, c/p)$: each ordered pair (i, j), $i \neq j$, is a non-loop edge independently with probability $q = c/p$, and every node has a self-loop (loops change neither reachability nor Src). m is the number of non-loop edges, $|E_S| = p + m$,

N the number of non-loop edges inside an SCC, so $|E_\times| = m - N$. $\mathrm{out}(v)$, $\mathrm{in}(v)$ are the sets of nodes reachable from v, resp. reaching v (v included), T_v^{out}, T_v^{in} their sizes; $v \rightsquigarrow w$ means $w \in \mathrm{out}(v)$. $x = x(c)$ is the largest solution of $x = 1 - e^{-cx}$ (the survival probability of a Poisson(c) Galton–Watson tree), $\theta_K = \theta_K(c)$ the probability that its total progeny is at least K ($\theta_K \downarrow x$), and d the smallest fixed point in $[0,1]$ of $f(d) = 1 + e^{-c} - e^{-c(1-x)d}$. A node is *in-isolated* if it has no non-loop parent. $\mathcal{G}$ denotes the *edge-rooted Poisson tree*: an edge (i,j); i has Poisson(c) parents and Poisson(c) further children, j Poisson(c) further parents and Poisson(c) children; every other node has Poisson(c) parents and Poisson(c) children besides the edge through which it was reached; all independent.

Textbook facts. (F1) Poisson coupling (van der Hofstad, 2016, Thm. 2.10): $d_{\mathrm{TV}}(\mathrm{Bin}(n,q), \mathrm{Poi}(c)) \le nq^2 + |nq - c|$. (F2) Chernoff: for a sum X of independent Bernoulli variables with mean μ and $\delta \in (0,1)$, $P(X \le (1-\delta)\mu) \le e^{-\delta^2\mu/2}$. (From $Ee^{-tX} \le \exp(\mu(e^{-t}-1))$ with $t = -\log(1-\delta)$, Markov's inequality, and $(1-\delta)\log(1-\delta) \ge -\delta + \delta^2/2$; the upper-tail version $P(X \ge \mu + s) \le e^{-s}$ for $s \ge e^2\mu$ follows in the same way.) (F3) Galton–Watson facts (van der Hofstad, 2016, Thms. 3.1, 3.16): the extinction probability is the smallest root of $\eta = e^{c(\eta-1)}$, so the survival probability is $x(c)$, positive iff $c > 1$.

Deterministic lemmas (machine-checked, ErLemmas.lean; no loop hypothesis). A node z is a *source node* if every ancestor of z is reachable from z. By Theorem 5, $(i,j) \in E_\uparrow$ iff some source node reaches j but not i.

(W1) If $(i,j) \in E_\uparrow$, there are $k \notin \{i,j\}$ with $k \to j$ and a source node z with $z \rightsquigarrow k$, $z \not\rightsquigarrow i$ (upEdge_witness).

(W2) If $i \to j$, $k \to j$, and a source node reaches k but not i, then $(i,j) \in E_\uparrow$ (mem_upEdges_of_witness).

(W3) If $i \to j$, $k \to j$ and $\mathrm{in}(k) \cap \mathrm{in}(i) = \emptyset$, then $(i,j) \in E_\uparrow$ (mem_upEdges_of_disjoint_anc).

(W5) A non-loop edge (i,j) lies inside an SCC iff $j \rightsquigarrow i$. The event $\{j \rightsquigarrow i\}$ and the set $\mathrm{out}(j)$ do not change when the edge (i,j) is added or removed (a simple path from j to i does not use it). (Trivial; not formalised.)

Explorations and forced edges. Forward exploration of $\mathrm{out}(v)$ keeps explored, active and unseen sets X_t, A_t, U_t, $a_t = |A_t|$, $\mathrm{disc}_t = t + a_t$, $X_0 = \emptyset$, $A_0 = \{v\}$. Step $t+1$ takes the oldest active node w, reveals the pairs (w,u), $u \in U_t$, and makes the ξ_{t+1} heads of present pairs active; it stops at $T = T_v^{\mathrm{out}}$. Backward exploration of $\mathrm{in}(v)$ reveals the pairs (u,w) instead. Let F be a set of at most three ordered pairs and G^F the graph in which the pairs of F are present and all other pairs are independent $\mathrm{Be}(q)$. *Bookkeeping.* In a forward exploration a pair is revealed only when its tail is explored, and each node is explored once, so each pair is revealed at most once; the pairs revealed at step $t+1$ are disjoint from those revealed before and from all pairs whose tail is not in X_{t+1}. Hence, conditionally on the history $\mathcal{H}_t$, the non-forced pairs revealed at step $t+1$ are independent $\mathrm{Be}(q)$, and $\xi_{t+1} = \#\{u \in U_t : (w,u) \in F\} + \mathrm{Bin}(|U_t| - f_w, q)$ with $f_w \le 3$ the number of forced pairs from w into U_t. Thus, given $\mathcal{H}_t$, ξ_{t+1} stochastically dominates $\mathrm{Bin}(|U_t| - 3, q)$ and is dominated by $3 + \mathrm{Bin}(p,q)$. The same holds for backward explorations (with heads in place of tails), and, by the same argument, for a

backward exploration run after a forward one, conditionally on the forward history, as long as it never discovers a node that the forward exploration discovered (only pairs with tail outside X_f and head explored backward are revealed, and none of them was revealed forward).

Fix $c > 1$ and $\varepsilon \in (0, 1/5)$ with $\beta = c(1-4\varepsilon) > 1$; put $n' = \lceil (1-4\varepsilon)p \rceil$, $\kappa = (\beta-1)^2/(8\beta)$, $\delta = (\beta-1)\varepsilon/(\beta+1)$, $t_0 = C_3 \log p$ with $C_3\kappa \geq 8$. Call an exploration *admissible* if, conditionally on its history, every offspring count ξ_{t+1} taken while $\mathrm{disc}_t < \varepsilon p$ dominates $\mathrm{Bin}(n', q)$ and every offspring count is dominated by $3 + \mathrm{Bin}(p, q)$. By the bookkeeping, forward and backward explorations in G^F are admissible for large p (since $|U_t| - 3 > (1-\varepsilon)p - 3 \geq n'$), and so is the backward exploration of Lemma 47 below.

Lemma 46 (dichotomy) *There is C (depending on c, ε) such that for large p every admissible exploration with final size T satisfies (a) $P(K \leq T < \varepsilon p) \leq Ce^{-\kappa K}$ for all $K \geq 1$; (b) $P\big(T \geq t_0$ and $[\mathrm{disc}_t < \varepsilon p$ for all t, or $a_\tau < \delta p$ at $\tau = \min\{t : \mathrm{disc}_t \geq \varepsilon p\}]\big) \leq Cp^{-8}$. Consequently, with probability at least $1 - Cp^{-7}$, every node v of G has, in both directions, $T_v < t_0$ or ($T_v \geq \varepsilon p$ and $a_\tau \geq \delta p$).*

Proof Let ξ'_t be independent $\mathrm{Bin}(n', q)$ variables coupled so that $\xi_t \geq \xi'_t$ for $t \leq \tau$ (admissibility; after τ or after the exploration stops, ξ'_t uses external randomness), and $W_t = 1 + \sum_{s \leq t}(\xi'_s - 1)$. By induction $a_t \geq W_t$ for $t \leq \min(T, \tau)$. Since $tn'q \geq t\beta$ and $(\beta+1)/2 = (1 - \frac{\beta-1}{2\beta})\beta$, (F2) gives $P(W_t \leq (\beta-1)t/2) \leq P(\mathrm{Bin}(tn', q) \leq t(\beta+1)/2) \leq e^{-\kappa t}$, so $\mathrm{Bad}(t_1) = \{\exists t \geq t_1 : W_t \leq (\beta-1)t/2\}$ has probability at most $e^{-\kappa t_1}/(1-e^{-\kappa})$. (a) If $T < \varepsilon p$ then $\tau = \infty$, so $0 = a_T \geq W_T$, which is $\mathrm{Bad}(K)$ when $T \geq K$. (b) On $\mathrm{Bad}(t_0)^c$ with $T \geq t_0$: for $t_0 \leq t \leq \min(T, \tau)$, $a_t \geq W_t > 0$, so the exploration cannot stop before τ, hence $\tau < \infty$. At τ, $a_\tau + \tau \geq \varepsilon p$, and if $\tau \geq t_0$ also $a_\tau > (\beta-1)\tau/2$; minimising $\max((\beta-1)\tau/2, \varepsilon p - \tau)$ over τ gives $a_\tau \geq \delta p$; if $\tau < t_0$, $a_\tau \geq \varepsilon p - t_0 \geq \delta p$. $e^{-\kappa t_0} \leq p^{-8}$. The consequence is a union bound over p nodes and two directions. □

Lemma 47 (large out-sets reach large in-sets) *With probability at least $1 - Cp^{-5}$: for all v, w, $T^{\mathrm{out}}_v \geq t_0$ and $T^{\mathrm{in}}_w \geq t_0$ imply $v \rightsquigarrow w$.*

Proof Fix (v, w). Explore $\mathrm{out}(v)$ forward up to τ_f (history $\mathcal{H}_f$, discovered set $R_f = X_f \cup A_f$). Since every ξ is dominated by $3 + \mathrm{Bin}(p, q)$, $P(\max_t \xi_t \geq \varepsilon p) \leq e^{-\varepsilon p}$ for large p (F2, upper tail), so $|R_f| \leq 2\varepsilon p$ except on that event. Then explore $\mathrm{in}(w)$ backward, stopping at τ_b or when a node of $R_f \subseteq \mathrm{out}(v)$ is discovered; in the latter case $v \rightsquigarrow w$. Otherwise the backward exploration is the true exploration of $\mathrm{in}(w)$, its unseen set has at least $p - 2\varepsilon p - \varepsilon p - 3 \geq n'$ usable nodes while $\mathrm{disc} < \varepsilon p$, and by the bookkeeping it is admissible conditionally on $\mathcal{H}_f$. By Lemma 46(b) for both explorations, outside probability $2Cp^{-8}$: if $T^{\mathrm{out}}_v, T^{\mathrm{in}}_w \geq t_0$ and no connection occurred, then $|A_f|, |A_b| \geq \delta p$. A_f and A_b are disjoint, and every pair $(a, a') \in A_f \times A_b$ is unrevealed (its tail is not in X_f, its head not in X_b), so given both histories these pairs are independent $\mathrm{Be}(q)$ (or forced present): $P(\text{no edge } A_f \to A_b) \leq (1-q)^{\delta^2 p^2 - 3} \leq e^{-c\delta^2 p/2}$. An edge $a \to a'$ gives $v \rightsquigarrow a \to a' \rightsquigarrow w$. Union bound over p^2 pairs. □

Let Ω be the intersection of the events of Lemma 46 (consequence) and Lemma 47, $P(\Omega^c) \le Cp^{-5}$. Put $\mathcal{S} = \{v : T_v^{\text{out}} \ge t_0\}$, $\mathcal{T} = \{v : T_v^{\text{in}} \ge t_0\}$.

Corollary 48 (structure) *On Ω: $\mathcal{S} \times \mathcal{T}$ is contained in the reachability relation; every SCC other than the one containing $\mathcal{S} \cap \mathcal{T}$ has fewer than t_0 nodes; if $|\mathcal{S} \cap \mathcal{T}| \ge t_0$, this SCC equals $\mathcal{S} \cap \mathcal{T}$.*

Proof The first claim is Lemma 47. A node $u \notin \mathcal{S} \cap \mathcal{T}$ has $T_u^{\text{out}} < t_0$ or $T_u^{\text{in}} < t_0$, and its SCC lies in $\text{out}(u) \cap \text{in}(u)$. If u were in the SCC of $\mathcal{S} \cap \mathcal{T}$, that SCC (of size $\ge t_0$) would lie in $\text{out}(u) \cap \text{in}(u)$, a contradiction. □

(This re-proves the part of the Karp–Łuczak structure theorem we need; the statement is recorded in the introduction of Dovgal et al., 2020.)

Lemma 49 (local concentration; standard coupling) *Fix R and a function φ of the isomorphism class of the radius-R neighbourhood $B_R(e)$ of a non-loop edge e (all edges incident to nodes at undirected distance $< R$ from e, rooted at the oriented edge e), and let $\Phi_p = p^{-1}\sum_e \varphi(B_R(e))$, the sum over non-loop edges. (a) If φ is bounded, $E\Phi_p \to cE_{\mathcal{G}}\varphi$ and $\text{Var}\,\Phi_p \to 0$; the same holds for node-rooted functionals with the node-rooted Poisson tree. (b) If $\varphi \ge 0$ and $\varphi(B_R(e)) \le \text{indeg}(j)$ for $e = (i,j)$, then $\limsup_p E\Phi_p \le cE_{\mathcal{G}}[\min(\varphi, M)] + C/M$ for every M, with C not depending on p or M.*

Proof (a) First moment: $E\sum_e \varphi(B_R(e)) = \sum_{(i,j)} q\, E[\varphi(B_R((i,j); G \cup \{(i,j)\}))]$. Explore the undirected neighbourhood of (i,j) breadth first, revealing for each explored node w all pairs (w,u) and (u,w). On $\{|B_R| \le M\}$, each explored node has $\text{Bin}(|U|, q)$ new children and $\text{Bin}(|U|, q)$ new parents among the unseen nodes, $|U| \ge p - M$, which (F1) couples with $\text{Poisson}(c)$ up to an error $\le Mq^2p + cM/p$ per count, and the pairs between already seen nodes are present with total probability $\le M^2q$. So the truncated ball is within $C(M,R)/p$ in total variation of the truncated $\mathcal{G}$. Tightness: each exploration step is dominated by $\text{Bin}(2p, q)$ (resp. $\text{Poisson}(2c)$), so $E|B_R| \le (1 + 2c + o(1))^R$ and $P(|B_R| > M) \le C^R/M$ uniformly in p; let $p \to \infty$, then $M \to \infty$. Second moment: $E\Phi_p^2 = p^{-2}\big[\sum_e + \sum_{e,f \text{ sharing a node}} + \sum_{e,f \text{ node-disjoint}}\big]$. The first two sums are $O(p)$ since $E\sum_v d(v)^2 = O(p)$. For node-disjoint e, f, explore both balls jointly in $G \cup \{e,f\}$: on {both sizes $\le M$} they are disjoint and non-adjacent except with probability $\le C(M,R)/p$, and then the two explorations are coupled with independent copies of $\mathcal{G}$. Hence $p^{-2}\sum_{\text{disjoint}} \to (cE_{\mathcal{G}}\varphi)^2$ and $\text{Var}\,\Phi_p \to 0$. (b) Apply (a) to $\min(\varphi, M)$ and bound the rest by $E\sum_{(i,j)} \text{indeg}(j)\mathbf{1}[\text{indeg}(j) > M] = E\sum_j \text{indeg}(j)^2\mathbf{1}[\text{indeg}(j) > M] \le p\,E[D^3]/M$, $D \sim \text{Bin}(p-1, q)$, and $E[D^3] \le c^3 + 3c^2 + c$ for all p (uniform integrability). □

Local functionals. All of the following are functions of $B_{2K+1}(e)$ (a breadth-first search of the first K nodes of a set stays within distance $K-1$), are invariant under isomorphism, and count all qualifying nodes, so no labelling enters. $A_e^K = \mathbf{1}[T_j^{\text{out}} \ge K,\ T_i^{\text{in}} \ge K]$, with $E_{\mathcal{G}}A^K = \theta_K^2$ (the descendant tree of j and the ancestor tree of i are independent Poisson Galton–Watson trees). $\Psi_K(e) = \text{cl}_1 \vee \text{cl}_2$, with $\text{cl}_1 = [T_i^{\text{in}} < K$ and j has a parent $k \notin \{i,j\}]$ and $\text{cl}_2 = [$some in-isolated s has $T_s^{\text{out}} < K$, $j \in \text{out}(s)$,

$i \notin \mathrm{out}(s)]$. $\varphi_K^\beta(e) = \#\{k \notin \{i,j\} : k \to j,\ T_k^{\mathrm{in}} < K,\ T_i^{\mathrm{in}} < K,\ \mathrm{in}(k) \cap \mathrm{in}(i) \neq \emptyset\} \leq \mathrm{indeg}(j)$, with $E_{\mathcal{G}} \min(\varphi_K^\beta, M) = 0$ (in $\mathcal{G}$ the ancestor trees of i and of another parent of j are disjoint).

Proposition 50 $(L,\ c > 1)$ *For $c > 1$, $N/p \to cx^2$ in L^1 and in probability.*

Proof By (W5), $N = \sum_e Z_e$ with $Z_e = \mathbf{1}[j \rightsquigarrow i]$; compare with $\tilde{N}_K = \sum_e A_e^K$. (i) $Z = 1, A = 0$: $j \rightsquigarrow i$ and ($T_j^{\mathrm{out}} < K$ or $T_i^{\mathrm{in}} < K$). By (W5) these events do not depend on the presence of e, so the expected number of such edges is $p(p-1)q\,P(i \in \mathrm{out}(j), T_j^{\mathrm{out}} < K)$ plus the same with $\mathrm{in}(i)$, and by exchangeability of the $p-1$ nodes other than j, $P(i \in \mathrm{out}(j), T_j^{\mathrm{out}} < K) = E[(T_j^{\mathrm{out}} - 1)\mathbf{1}(T_j^{\mathrm{out}} < K)]/(p-1) \leq K/(p-1)$. Total $\leq 2cK$. (ii) $Z = 0, A = 1$: on Ω, $T_j^{\mathrm{out}} \geq t_0$ and $T_i^{\mathrm{in}} \geq t_0$ would give $j \rightsquigarrow i$, so T_j^{out} or T_i^{in} lies in $[K, t_0) \subseteq [K, \varepsilon p)$. The expected number of such edges is at most $\sum_{(i,j)} q\, P_{G^{\{e\}}}(\cdot) \leq cp \cdot 2Ce^{-\kappa K}$ by Lemma 46(a) in $G^{\{e\}}$; on Ω^c the count is at most m and $E[m\mathbf{1}_{\Omega^c}] \leq (Em^2)^{1/2} P(\Omega^c)^{1/2} = O(p \cdot p^{-5/2})$. So $E|N - \tilde{N}_K| \leq 2cK + 2cCpe^{-\kappa K} + o(1)$. By Lemma 49(a), $\tilde{N}_K/p \to c\theta_K^2$ in L^2; hence $\limsup_p E|N/p - cx^2| \leq 2cCe^{-\kappa K} + c|\theta_K^2 - x^2| \to 0$ as $K \to \infty$. □

The same argument for the node functional $\mathbf{1}[T_v^{\mathrm{out}} \geq K, T_v^{\mathrm{in}} \geq K]$ gives $|\mathcal{S} \cap \mathcal{T}|/p \to x^2$, so by Corollary 48 the giant SCC has $(x^2 + o_P(1))p$ nodes and all other SCCs fewer than $C_3 \log p$, w.h.p.

Proposition 51 (the tree computation) *In $\mathcal{G}$, with Ψ_∞ defined by $T_i^{\mathrm{in}} < \infty$ and $T_s^{\mathrm{out}} < \infty$, $E_{\mathcal{G}}\Psi_K \uparrow P_{\mathcal{G}}(\Psi_\infty) = F_\uparrow$ as $K \to \infty$.*

Proof Both clauses increase in K to their $K = \infty$ versions (monotone convergence). Let $A(i)$ be the event that the ancestor tree of i is infinite and $D(j)$ that the descendant tree of j is infinite (probability x each, independent). The other parents of j are Poisson(c) many. (1) Not $A(i)$: Ψ_∞ holds iff j has another parent. Clause 1 needs exactly that; clause 2 needs it too, since a path from s to j ends with an edge $k \to j$ and $k = i$ would put i into $\mathrm{out}(s)$. Probability $(1-x)(1-e^{-c})$. (2) $A(i)$: clause 1 fails. Clause 2 needs (a) not $D(j)$, since $\mathrm{out}(s)$ contains the descendant tree of j (probability $1-x$), and (b) another parent k of j and a chain $k = u_0 \leftarrow u_1 \leftarrow \cdots \leftarrow u_r = s$ ending at an in-isolated s, in which every u_t has only finite descendant trees below its children other than u_{t-1} ($u_{-1} = j$); such an s does not reach i, because paths into i stay in the ancestor tree of i. The probability that all other children of a node have finite descendant trees is $E[(1-x)^{\mathrm{Poi}(c)}] = e^{-cx} = 1 - x$. Let $Q(v)$ be the event that v is in-isolated or some parent u of v has all other children finite and $Q(u)$ holds, and $d^* = P(Q)$. Parents are Poisson(c) many and each is *good* (other children finite and Q) independently with probability $(1-x)d^*$, so by Poisson thinning $d^* = e^{-c} + 1 - e^{-c(1-x)d^*} = f(d^*)$. Q is the increasing union of its depth-n truncations Q_n, with $P(Q_0) = e^{-c} = f(0)$ and $P(Q_{n+1}) = f(P(Q_n))$; as f is increasing and continuous, $P(Q_n) = f^{n+1}(0)$ increases to the smallest fixed point d of f (if d' is a fixed point, $f^n(0) \leq f^n(d') = d'$). So $d^* = d$. The number of good other parents of j is Poisson($c(1-x)d$), so (b) has probability $1 - e^{-c(1-x)d}$, independently of $A(i)$ and $D(j)$. Adding (1) and (2) gives $F_\uparrow$. □

(An independent Monte Carlo of the event itself, not of the closed form, gives 0.38756 ± 0.00077 against 0.38773 at $c = 1.5$ and 0.19049 ± 0.00062 against 0.18962 at $c = 2$.)

Proposition 52 (reduction of $E_\uparrow$ to Ψ_K, $c > 1$) *$E \sum_e |\mathbf{1}[e \in E_\uparrow] - \Psi_K(e)| \le (\delta_1(K) + o(1))p$ with $\delta_1(K) = C(c^2 + K^2)e^{-\kappa K} + o_K(1)$, where $o_K(1) \to 0$ as $K \to \infty$ and C does not depend on K or p.*

Proof Let $e = (i, j)$ be a non-loop edge. *Clause 2 never errs*: an in-isolated s is a source node, so $j \in \mathrm{out}(s)$, $i \notin \mathrm{out}(s)$ give $e \in E_\uparrow$. *Case I: $T_i^{\mathrm{in}} < K$.* (I-a) If $e \in E_\uparrow$, (W1) gives another parent of j, so clause 1 holds: no error. (I-b) Clause 1 holds and $e \notin E_\uparrow$. Then $\mathbf{1}[\text{I-b}] \le \sum_{k \notin \{i,j\}} \mathbf{1}[e, (k, j) \text{ present}, \ T_i^{\mathrm{in}} < K, \ e \notin E_\uparrow]$ (a union over *all* other parents; no parent is singled out). Fix the triple (i, j, k) and force $F = \{(i, j), (k, j)\}$. By (W3), $e \notin E_\uparrow$ forces $\mathrm{in}(k) \cap \mathrm{in}(i) \neq \emptyset$, and by (W2) every source node reaching k reaches i. Split on T_k^{in}: (α_1) $K \le T_k^{\mathrm{in}} < \varepsilon p$: probability $\le Ce^{-\kappa K}$ (Lemma 46(a), backward, in G^F). (α_2) $T_k^{\mathrm{in}} \ge \varepsilon p$: the backward exploration of $\mathrm{in}(k)$ explores at least $\varepsilon p/2$ nodes. When a node $y \neq j$ is explored, the pairs (u, y) with u unseen are revealed (all absent with conditional probability $\ge (1 - q)^p$), and the pairs (u, y) with u already seen are never revealed by this exploration and are independent of it and of each other (distinct heads); so, given the history, y is in-isolated with probability $\ge (1-q)^p \ge e^{-c}/2$, and the number of in-isolated nodes among the first $\varepsilon p/2$ explored nodes dominates $\mathrm{Bin}(\varepsilon p/2 - 1, e^{-c}/2)$. By (F2) at least K of them are in-isolated except with probability $e^{-\Omega(p)}$. Each in-isolated node of $\mathrm{in}(k)$ is a source node reaching k, hence reaches i, so lies in $\mathrm{in}(i)$; but $|\mathrm{in}(i)| < K$. So (α_2) has probability $e^{-\Omega(p)}$. (β) $T_k^{\mathrm{in}} < K$ and $\mathrm{in}(k) \cap \mathrm{in}(i) \neq \emptyset$. Summed over k this is $\varphi_K^\beta(e)$. The number of triples is p^3 and each has both pairs present with probability q^2, so cases (α_1), (α_2) contribute at most $c^2 p(Ce^{-\kappa K} + e^{-\Omega(p)})$; case ($\beta$) contributes $E \sum_e \varphi_K^\beta(e) \le (C/M + o(1))p$ for every M by Lemma 49(b), i.e. $o(p)$. *Case II: $T_i^{\mathrm{in}} \ge K$.* Clause 1 fails. (II-0) $K \le T_i^{\mathrm{in}} < \varepsilon p$: at most $cpCe^{-\kappa K}$ such edges in expectation (Lemma 46(a) in $G^{\{e\}}$). (II-1) $T_i^{\mathrm{in}} \ge \varepsilon p$, so $i \in \mathcal{T}$. Work on $\Omega' = \Omega \cap \{\max_v \mathrm{indeg}(v) \le C \log p\}$; $P(\Omega'^c) \le Cp^{-2}$, and the edges counted on Ω'^c contribute at most $E[m\mathbf{1}_{\Omega'^c}] \le (Em^2)^{1/2} P(\Omega'^c)^{1/2} = O(1)$. An error here is an edge $e \in E_\uparrow$ with $\Psi_K(e) = 0$. Take a witness source node z of (W1), in a source SCC σ. On Ω, $z \in \mathcal{S}$ would give $z \rightsquigarrow i$ (Lemma 47), so $T_z^{\mathrm{out}} < t_0$. (A) $\sigma = \{z\}$ (i.e. z in-isolated) and $T_z^{\mathrm{out}} < K$: clause 2 holds with $s = z$; no error. (B1) z in-isolated and $K \le T_z^{\mathrm{out}} < t_0$. The number of such edges is at most $\sum_z \mathbf{1}[K \le T_z^{\mathrm{out}} < t_0] \sum_{j \in \mathrm{out}(z)} \mathrm{indeg}(j)$. Given $\mathrm{out}(z) = O$ with $|O| = r$ (a completed forward exploration), the pairs with tail outside O and head in O are unrevealed, so the expected number of edges into O is at most $r^2 + rpq = r^2 + cr$. With $P(T_z^{\mathrm{out}} = r) \le Ce^{-\kappa r}$ (Lemma 46(a)), the expectation is at most $p \sum_{r \ge K} Ce^{-\kappa r}(r^2 + cr) \le C' p K^2 e^{-\kappa K}$. (B2) $|\sigma| \ge 2$. Then $|\sigma| < t_0$ (as $\sigma \subseteq \mathrm{out}(z)$), every node of σ has out-set $\mathrm{out}(z)$ of size $< t_0$, and the number of edges with head in $\mathrm{out}(z)$ is at most $t_0 \cdot C \log p$ on Ω'. Let N_σ be the number of source SCCs with $2 \le |\sigma| < t_0$. For a set R of size r, "R strongly connected" and "no edge enters R" depend on disjoint sets of pairs; the first requires every node of R to have a parent in R (probability $\le (rq)^r$, distinct heads), the second has probability $(1 - q)^{r(p-r)}$. Hence $EN_\sigma \le \sum_{r=2}^{t_0} \binom{p}{r} (rq)^r (1-q)^{r(p-r)} \le \sum_{r \ge 2} \left(ce^{1-c} e^{cr/p}\right)^r = O(1)$, using $\binom{p}{r} \le p^r/r!$, $r^r/r! \le e^r$, $r \le t_0 = O(\log p)$ and $ce^{1-c} < 1$ for $c \neq 1$. So the (B2) errors have expectation $O(1) \cdot t_0 C \log p = O(\log^2 p) = o(p)$: an expectation on the event Ω' plus the Cauchy–Schwarz bound for Ω'^c above. Adding all cases gives the claim. □

Theorem 53 (S, $c > 1$) *For $c > 1$, $|E_\uparrow|/p \to cF_\uparrow$, and therefore $S = \frac{1}{2} + |E_\uparrow|/(2(p+m)) \to \frac{1}{2} + cF_\uparrow/(2(1+c))$, in L^1 and in probability.*

Proof $E\big||E_\uparrow|/p - cF_\uparrow\big| \le E\sum_e |\mathbf{1}[e \in E_\uparrow] - \Psi_K(e)|/p + E|\Phi_p^{\Psi_K} - cE_{\mathcal{G}}\Psi_K| + c|E_{\mathcal{G}}\Psi_K - F_\uparrow|$; the terms are bounded by Prop. 52, Lemma 49(a) and Prop. 51; let $p \to \infty$, then $K \to \infty$. $m/p \to c$ in L^2 ($m \sim \mathrm{Bin}(p(p-1), q)$), and S and L are bounded continuous functions of $(|E_\uparrow|/p, m/p)$ and $(N/p, m/p)$ on $[0,\infty)^2$. □

Proposition 54 ($c \le 1$) *(a) For $c < 1$, $EN \le c^2/(1-c)$. (b) For $c \le 1$, $N/p \to 0$ in L^1, so $L \to \frac{1}{2} + \frac{c}{2(1+c)}$ ($L(1) = 3/4$). (c) For $c \le 1$, $|E_\uparrow|/p \to c(1-e^{-c})$ in L^1, so $S \to \frac{1}{2} + \frac{c(1-e^{-c})}{2(1+c)}$ ($S(1) = 0.65803$); this is $F_\uparrow$ at $x = 0$.*

Proof (a) A non-loop edge inside an SCC lies on a simple directed cycle, and the expected number of pairs (simple cycle of length k, edge on it) is $k \cdot \frac{(p)_k}{k} q^k \le c^k$; sum over $k \ge 2$. (b) $c < 1$ from (a). $c = 1$: $D(p, 1/p)$ is a subgraph of $D(p, c'/p)$ for every $c' > 1$ under the monotone coupling, and an edge inside an SCC of the subgraph is inside an SCC of the larger graph, so $\limsup_p EN/p \le c'x(c')^2 \to 0$ as $c' \downarrow 1$ (Prop. 50). (c) Let $\Psi_K^1 = \mathrm{cl}_1$. Case I ($T_i^{\mathrm{in}} < K$): (I-a) as before, no error; (I-b) with the same union over other parents k and forced $F = \{(i,j),(k,j)\}$, (W3) forces $\mathrm{in}(k) \cap \mathrm{in}(i) \ne \emptyset$, so either $T_k^{\mathrm{in}} \ge K$ or the event of φ_K^β holds. The expected number of triples with $T_k^{\mathrm{in}} \ge K$ is $E\sum_e \#\{k \notin \{i,j\} : k \to j,\ T_k^{\mathrm{in}} \ge K\}$, a local functional dominated by $\mathrm{indeg}(j)$, so by Lemma 49(b) it is at most $(c \cdot c\,\theta_K + C/M + o(1))p$; the φ_K^β part is $o(p)$ as before. Case II ($T_i^{\mathrm{in}} \ge K$): at most $E\sum_e \mathbf{1}[T_i^{\mathrm{in}} \ge K] = (c\theta_K + o(1))p$ (Lemma 49(a)). Hence $\limsup_p E\sum_e |\mathbf{1}[e \in E_\uparrow] - \Psi_K^1(e)|/p \le c^2\theta_K + c\theta_K + C/M$, which tends to 0 as $K, M \to \infty$ because $\theta_K \to x(c) = 0$ for $c \le 1$ (F3). Finally $E_{\mathcal{G}}\Psi_K^1 = (1 - \theta_K)(1 - e^{-c}) \to 1 - e^{-c}$. □

What is not proved. Rates of convergence; the critical window $c = 1 + \lambda p^{-1/3}$ (Łuczak and Seierstad, recorded in Dovgal et al., 2020; not used); almost sure convergence; the regime of fixed T with $p \to \infty$. The proof does not use the structure theorems of Karp and of Łuczak (not opened by us); it re-proves the part it needs.

Numerical checks of the intermediate claims (fixed seeds, exact computation; supplement `x1proof/verify`). (W4) $E_\uparrow \subseteq$ edges into confluences: 0 violations in 48 graphs. Dichotomy (Lemma 46): at $c = 2$, $p = 10^5$, the mass of $K \le T < 0.02p$ is 0.0097, 0.0008, 0 for $K = 5, 10, 20$; at $c = 1$ it is 0.37, 0.25, 0.17 (critical, no dichotomy, as expected). Structure (Corollary 48): for $c \ge 1.5$ the second-largest SCC never had more than 2 nodes; giant fraction 0.63439 ($p = 10^6$, $c = 2$) against $x^2 = 0.63491$. Prop. 52 edge by edge: at $c = 2$, $p = 2 \cdot 10^5$, $|E_\uparrow| = 76{,}220$ and $\sum_e \Psi_{40}(e) = 76{,}220$ with no edge in either difference; clause 2 never erred in any run. Concentration: the standard deviation of N/p times $\sqrt{p}$ is 2.88–3.07 for $p = 5 \cdot 10^3$ to $3.2 \cdot 10^5$ ($c = 2$), as Lemma 49 requires.

Appendix Q Accuracy Interpretation of s_ρ

An edge $\{x, y\}$ of a DAG is shown in a uniformly random order together with $(\rho_{\text{first}}, \rho_{\text{second}})$. A *pairwise count rule* is $f(a, b) \in [0, 1]$, the probability of declaring "first is the tail"; the *sort rule* is $f_{\text{sort}} = 1, \frac{1}{2}, 0$ for $a <, =, > b$. For the true edge $x \to y$ the accuracy is $\frac{1}{2}f(\rho_x, \rho_y) + \frac{1}{2}(1 - f(\rho_y, \rho_x))$.

Proposition 55 (elementary) *On every DAG, (i) the accuracy of f_{sort} is exactly s_ρ; (ii) every pairwise count rule has accuracy at most s_ρ; (iii) f_{sort} never errs on an edge with $\rho_x \neq \rho_y$.*

Lean: *accRule_sortRule, accRule_sortRule_rho, accRule_le_sOf, accRule_le_sRho, edgeAcc_sortRule_strict, rho_not_gt_of_edge.*

Proof By Reisach et al.'s Theorem 1, $\rho_x \leq \rho_y$ on an edge $x \to y$. If $\rho_x < \rho_y$ the accuracy is at most 1, attained by the sort rule; if $\rho_x = \rho_y = a$ it is $\frac{1}{2}f(a, a) + \frac{1}{2}(1 - f(a, a)) = \frac{1}{2}$ for every f, because the random order forces the tie decision to 1/2. These are the edge scores. (i) restates the definition of s_ρ; (ii) is a three-line consequence of their Theorem 1. □

Combined with Theorem 5, the best accuracy of any pairwise count rule on the stationary unrolled DAG is $\frac{1}{2} + |E_\uparrow|/(2|E_S|)$, at most 0.665 for ER summary graphs and 0.75 for river-like in-forests ($\mu \leq 1$). The restriction to counts is necessary (exhaustive enumeration over labelled DAGs, not proved):

nodes p	2	3	4	5
Bayes rule on the two counts (= edge-pooled s_ρ)	0.5	0.5625	0.592262	0.605547
Bayes rule on counts and the two node degrees	0.5	0.625	0.660714	0.669238

For the skeleton $a - c - b$ the collider has counts $(2, 2, 3)$ (all edges strict) while the chain and the fork have all counts equal to 3; among the tied graphs the centre c is the tail with probability 2/3 under a uniform prior, so a degree-aware rule beats the coin. So s_ρ is the Bayes accuracy of pairwise count rules only. This section is interpretive, not a new theorem.

Appendix R Benchmark Experiment B1: Design, Results and Deviations

Process. The analysis plan (`b1/PLAN.md`) was committed before any simulation that produced an algorithm score; only graph-only exact values were computed before it. Two deviations were recorded in the plan after the validation run (seeds 9000–9004, used only for debugging and timing) and before the test run: D1, the burn-in was raised from 50 to 100 steps because the validation mixing check failed (error $2.8 \cdot 10^{-3}$, threshold 10^{-3}; $1.4 \cdot 10^{-10}$ after the change); D2, SID is undefined when PC returns a cyclic directed part, so PC is dropped from the SID comparison in those cells (930 instead of 1,190 tests). D3: a container restart killed the first test run after about half

of one shard; no results file had been written and no partial score was inspected. The test run was redone from scratch with the same deterministic seeds. No threshold or hypothesis was changed. As with the other studies, the plan is internally time-stamped and agent-executed (Appendix K).

Design. Summary graphs with a self-loop at every node, $p \in \{5, 10\}$, 30 test graphs (seeds 1000–1029) per family: symmetrised ER (expected degree 2) and the directed ring (strongly connected), the chain, random recursive out- and in-trees, and directed ER $G(p, c/p)$ with $c \in \{1, 2\}$. Linear Gaussian VAR(1) with every weight $\sim \mathrm{Unif}(0.5, 1)$, rescaled to spectral radius 0.8 if larger, so relatives in the population equal graph relatives; noise s.d. $\sim \mathrm{Unif}(0.5, 1)$ per node (R-het, primary), equal noise (R-eq), and R-het standardised (R-std). G1: x_1 is noise, then $T - 1$ steps, truth G_T. G2: 100 latent burn-in steps, then the window of T slices, truth the window copy of G_T; the burn-in makes slice-B nodes latent common parents of the first window slice, so we also report SHD excluding pairs inside the first slice. G2s: exact stationary start via the discrete Lyapunov equation (replication of G2). Methods with columns shuffled: rel-SortnRegress (Reisach et al. Algorithms 1–2, $\alpha = 0.05$), R^2-SortnRegress, Var-SortnRegress, a random order followed by the same regression step, FLOP (`flopsearch` 0.3.0, BIC penalty 2, 10 restarts), PC (causal-learn, Fisher z, $\alpha = 0.01$; not run at $p = 10$, $T = 10$ for time), and a labels-visible reference that sorts by time. DAGMA, NOTEARS and GOLEM were run in the separately planned extension of Appendix T. SID by `gadjid`.

Pre-specified hypotheses and verdicts. H1 (start helps rel-SortnRegress on non-strongly-connected families; at least 4/5 Holm-significant and pooled mean ≥ 0.02): supported, 5/5, pooled $+0.125$ $[0.116, 0.134]$; tracking (Spearman of cell means against the exact gap ≥ 0.5): 0.80; the plan's $p = 1.2 \cdot 10^{-10}$ treated the 42 cells as independent, and a graph-cluster permutation gives $p \leq 10^{-4}$ (post hoc, Appendix T); OLS $d = 0.53\,\mathrm{gap} - 0.027$. H2a (gain below 0.03 in every strongly connected cell with $T \geq p$, the review's expectation): fails, as the exact gaps predicted before the run (e.g. ring $p = 5$, $T = 5$: $+0.217$; only symER $p = 5$, $T = 10$ is small, $+0.019$). H2b (non-negligible gain at $T < p$): holds, $+0.13$ to $+0.31$ in all six cells. H2c (symER gain decreases with T): holds, $+0.195$, $p = 1.2 \cdot 10^{-11}$. H3 (some Holm-significant SHD rank swap): supported, 11 swaps: rel- vs R^2-SortnRegress on chains ($p = 10$, $T = 3, 5$), Var-SortnRegress vs FLOP (3 cells) and vs PC (5 cells) on in-trees, and R^2-SortnRegress vs the random order on one ring cell. SID: 9 swaps; orientation score: 9; SHD without first-slice pairs: 9. Swap counting is a crude summary: the post hoc interaction analysis of Appendix T finds Holm-significant G1×G2 interactions in 294 of 595 pair×cell SHD tests. Median Kendall τ between the G1 and G2 rankings: 0.73. FLOP is best on average under both generators (normalised SHD 0.27 vs 0.46).

Two artefacts. Pooled G1−G2 order accuracy over all seven families: rel-SortnRegress $+0.133$ (the H1 value $+0.125$ is over the five non-strongly-connected families; unchanged by standardisation, since it is scale invariant); Var-SortnRegress $+0.202$ raw, $+0.244$ with equal noise, $+0.002$ standardised. Variance accumulates from the noise-only first slice, so the fixed start is var-sortable; standardising removes this but not the relatives artefact. Only the rel- vs R^2-SortnRegress swaps follow the mechanism of Theorems 2 and 5.

Cost of the burn-in. The labels-visible reference has normalised SHD 0.044 (G1) and 0.273 (G2), but 0.041 vs 0.048 without first-slice pairs (G2s identical): stationary-start benchmarks should score the window without first-slice pairs or as an ADMG.

Exact curves and the Δ bound. Figure R2 shows exact $s_\rho(G_T)$ for $T = 2, \ldots, 20$. At $T = 3$ it lies between 0.81 and 1.00; at $T = 20$ symER is still at 0.55–0.57 and the ring with $p = 10$ at 0.74. The Δ form of Theorem 2 was checked against brute-force relatives enumeration on 3,000 random summary graphs × 7 horizons: no violation of either bound, equality in 3,164 cases, strictly sharper than the $p-1$ bound in 10,186. For Theorem 5, $B = \Delta$ gave the stationary value in 726 of 726 graphs and $B = \Delta - 1$ did not in 445, so Δ is sufficient for all of them and necessary for 445 of them (Remark 20).

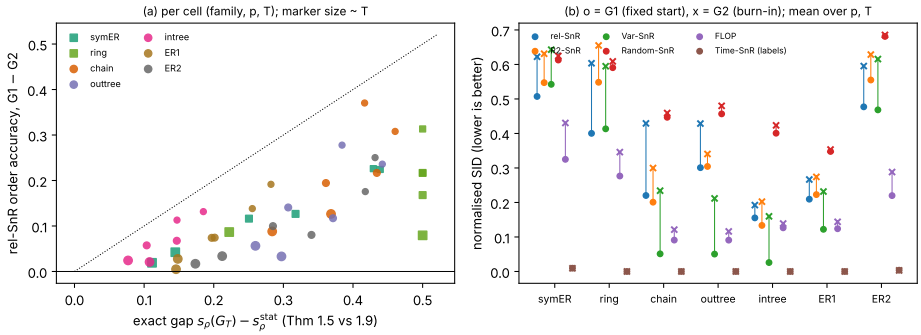


Fig. R2 B1: per-cell G1−G2 gain in rel-SortnRegress order accuracy against the exact population gap (all 42 cells below the diagonal and above 0), and normalised SID per method under G1 and G2. The exact curves of $s_\rho(G_T)$ against T are in Figure 3.

Appendix S B1-NL: Replication under Nonlinear Dynamics, Signed Weights and Uniform Noise

Process. A separate plan (`b1nl/PLAN.md`) was committed before any algorithm score; two notes (a faster but output-identical diagnostic, and the limited power of sanity rule S1) were added after the validation run and before the test run. No threshold, hypothesis or seed changed. The test run (5,040 data sets) was run once.

Design. As B1 (Appendix R) with the B1 code imported unchanged: seven families, $p \in \{5, 10\}$, $T \in \{3, 5, 10\}$, $n = 1000$, G1 against G2 (burn-in 100), R-het and R-std. Three parametrisations. NL: additive VAR with a linear self-loop and cross terms $w\, g(x)$, g drawn from tanh, sin and $(x + x^2/2)/(1 + x^2/4)$, positive w, stability by scaling the Lipschitz-bound matrix to spectral radius at most 0.8. SGN: linear, all weights $\pm\text{Unif}(0.5, 1)$. SGNU: SGN with uniform noise. New seeds (25 graphs per family and p for NL and SGN, 10 for SGNU); PC and R-eq dropped and G2s replaced by a 500-step burn-in, to fit one core (stated in the plan).

Results. H1 (same rule as B1) supported in all three: 5/5 families Holm-significant ($p \leq 1.4 \cdot 10^{-9}$); pooled gain +0.108 [0.100, 0.116] (NL), +0.137 [0.128, 0.146] (SGN), +0.137 [0.123, 0.152] (SGNU); positive in all 126 cells. Tracking: Spearman 0.68, 0.82, 0.78 (slopes 0.34, 0.56, 0.56). Under NL the correlation-count estimator recovers about 65% of the graph relatives (recall 0.67 under G1, 0.62 under G2), which explains the smaller slope, as stated before the run. Under signed weights sample relatives can deviate from graph relatives (out-trees under G2: estimated rel-sortability 0.455), yet the gain is as large as in B1. H3: one Holm-significant SHD swap under NL (R^2-SortnRegress against the random order, chains, $p = 10$, $T = 10$), none under SGN, not testable for SGNU (10 seeds, stated in the plan). By the plan's rule H3 is therefore supported under NL, but on a substantively minor swap, and we do not count it as a replication of B1's swaps; B1's rel- against R^2-SortnRegress chain swap was not detected on SHD at this replication's power, although the interaction replicates (Appendix T). Without PC, 5 of B1's 11 swap types cannot occur. FLOP is the best label-free method run there under both generators in every configuration. The var/rel split holds: Var-SortnRegress gains +0.19, +0.24, +0.24 raw and about 0 standardised; rel-SortnRegress +0.11, +0.15, +0.14. Sanity rules pass (S1 at most 0.015; mixing error at most $5.6 \cdot 10^{-5}$).

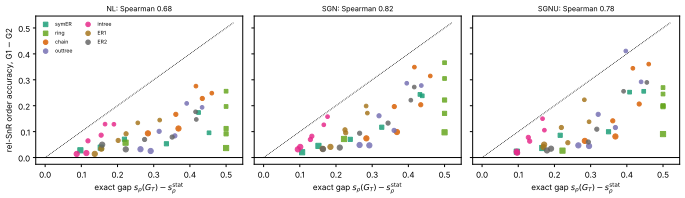


Fig. S3 B1-NL: per-cell G1−G2 gain in rel-SortnRegress order accuracy against the exact population gap, for NL, SGN and SGNU.

Appendix T Competitive Learners, Re-Analysis, Detectability Oracle and Generator Survey

Material in `b1x/` and `lit3/` of the supplement. The learner extension is a planned experiment; the re-analysis and the oracle are post hoc and use no new data.

Learner extension: process and design. The plan (`b1x/PLAN.md`), code and a validation timing were committed and pushed before any score (Appendix K); the test run took 4.21 h on one core and was analysed once, with no deviation (the size rule kept 12 seeds after a projected 4.35 h). B1's generators and SEM, new seeds: families ring, chain, in-tree and ER(c=2), $p \in \{5, 10\}$, $T \in \{3, 5\}$, 12 graphs per (family, p), i.e. 96 summary graphs and 384 data sets (G1 and G2), each analysed raw and standardised with time labels hidden. DAGMA linear (`dagma` 1.1.1, $\lambda = 0.03$), NOTEARS linear (the authors' reference code, $\lambda = 0.1$), GOLEM-EV/NV (a numpy port of the authors' code with the same score, optimiser, defaults and post-processing, gradient-checked;

$p = 5$ only, as planned), all thresholded at 0.3; FLOP, the SortnRegress baselines and the empty graph as comparators; PCMCI+ (tigramite, ParCorr, $\tau_{\max} = 1$) on the time-indexed data as a labels-using reference. Metrics: SHD/$|E|$ (empty graph = 1), normalised SID, ancestral order accuracy; unit of inference the summary graph. T= 10 and the families symER, out-tree and ER(c=1) were excluded for run time.

Verdicts. HX1 (some raw continuous method Holm-significantly favoured by the fixed start): supported; all four raw variants improve on all three metrics (DAGMA SHD/$|E|$ -0.548 $[-0.663, -0.438]$, SID -0.147, order accuracy $+0.197$; NOTEARS -0.419, -0.127, $+0.156$; GOLEM-EV -0.491, -0.182, $+0.225$; GOLEM-NV -0.409, -0.193, $+0.232$; Holm $\leq 4 \cdot 10^{-8}$). HX2 (opposite-sign, both Holm-significant rank change against FLOP on SHD/$|E|$, 64 tests): supported, two changes, both on in-trees: DAGMA$-$FLOP -0.218 under G1 (Holm 0.012) and $+0.306$ under G2 ($1.4 \cdot 10^{-5}$); NOTEARS$-$FLOP -0.176 (0.0011) and $+0.193$ (0.020); the same two appear on SID and order accuracy, and DAGMA's persists without first-slice pairs (0.08 / 0.73 against FLOP's 0.30 / 0.37). HX3 (standardisation removes the gain): DAGMA's SHD gain removed (-0.005 $[-0.027, 0.016]$), all other pairs reduced but not removed by the plan's rule; on order accuracy the standardised effect is significantly negative for DAGMA (-0.064), GOLEM-EV (-0.093) and GOLEM-NV (-0.097), and NOTEARS keeps $+0.036$. rel-SortnRegress replicates ($+0.187$ $[0.169, 0.205]$), and B1's chain swap of rel- against R^2-SortnRegress recurs on the new seeds. No fit failed; NOTEARS on standardised data returned a cyclic graph in 15.6% of data sets (SID missing for 60 of 384). Best label-free method by mean SHD/$|E|$ per (family, p): FLOP in 8/8 cells under G2; under G1 DAGMA on chains with $p = 5$ (0.09 vs 0.10) and on in-trees with $p = 10$ (0.06 vs 0.33), NOTEARS on in-trees with $p = 5$ (0.09 vs 0.27). Raw continuous methods beat the empty graph on average under both generators but are worse than it on chains under G2 (1.14–1.25). Standardised DAGMA and GOLEM orient most tree edges backwards under G1 (order accuracy 0.26–0.27). The gains and the reversals are largest on chains and in-trees, where G1 makes the window almost perfectly var-sortable (Var-SortnRegress 0.95–0.96), and absent or reversed on rings and ER(c=2). This is consistent with a var-sortability effect, but the experiment does not establish it.

Predictions of the plan and what the attribution rests on. The plan stated three predictions. For HX1 it predicted favourable effects on raw data for all continuous methods; this held. For HX2 it predicted no rank change against FLOP on SHD/$|E|$; this failed (two changes). For HX3 it predicted "reduced, not removed" for DAGMA and NOTEARS on order accuracy, because some scale-free gain should survive standardisation. By the plan's rule both were labelled "reduced, not removed", but for DAGMA the standardised effect is negative (-0.064), so no gain survives; for NOTEARS a gain survives ($+0.036$ order accuracy, -0.077 SHD/$|E|$, -0.073 SID, all Holm-significant). The standardised runs are a weak test of the attribution to var-sortability. Standardised DAGMA and GOLEM have order accuracy below 1/2 under both generators (DAGMA 0.35 / 0.41, GOLEM 0.31–0.33 / 0.41–0.42), so they cannot show a scale-free effect of either sign. NOTEARS's remaining gain is a scale-free start effect on a competitive method. The fixed start also raises R^2-sortability, which

is scale free and outside our theory, and we did not separate it. In B1 the gain in Var-SortnRegress order accuracy is the same on rings (+0.224) and chains (+0.231), while the learners gain on chains and not on rings; the explanation by levels (0.73 against 0.93 under G1) was made after the results. A control that rescales each data set to the other generator's marginal variances was planned afterwards and is reported in Section 5.2 and Appendix U. GOLEM did not overtake FLOP on in-trees (G1 differences −0.011 and −0.042, Holm 1 and 0.46). Table 4 omits GOLEM-EV (SHD/$|E|$ 0.44 / 0.93 raw) and standardised GOLEM (0.97 / 1.30 EV, 0.95 / 0.99 NV).

Re-analysis of B1 and B1-NL (post hoc). Graph-cluster bootstrap CIs (10,000 resamples, the three horizons of a graph kept together) are narrower than data-set bootstrap CIs, because T is a fixed design factor; they are used in Table 3. Tracking: cell-level Spearman 0.80 $[0.76, 0.82]$, unit-level 0.62 $[0.59, 0.64]$, slope 0.53 $[0.50, 0.56]$ in B1; 0.68/0.44/0.34 (NL), 0.82/0.64/0.56 (SGN), 0.78/0.62/0.56 (SGNU); graph-cluster permutation $p \leq 10^{-4}$ in all four (the floor). Interactions $I = [m_1 - m_2](\text{G1}) - [m_1 - m_2](\text{G2})$ per graph: Holm-significant in 294/595 (B1), 113/420 (NL), 232/420 (SGN) pair×cell SHD tests and 0 for SGNU (10 graphs cannot pass Holm). In-tree Var-SortnRegress−FLOP on SHD: −68.4 $[-77.0, -60.4]$ (B1), −49.9 (NL), −75.5 (SGN), −75.5 (SGNU); SID −0.12 in all four. Chain rel-−R^2-SortnRegress: SID and orientation interactions replicate at 60–80% of B1's size; the family-pooled SHD interaction is about 0 in B1-NL, but in B1's swap cell ($p = 10$, $T = 3$) it keeps its sign (−52, −16, −32). Empty-graph reference (B1, all families, G1 / G2, SHD/$|E|$): rel- 1.05/1.70, R^2- 1.33/1.96, Var- 0.61/1.39, random 1.72/2.04, FLOP 0.27/0.46, PC 0.40/0.64; the empty graph has normalised SID 0.50. Order accuracy without edges leaving the first window slice (non-SC families): rel-SortnRegress +0.089 $[0.079, 0.100]$ (B1), +0.067 (NL), +0.087 (SGN), +0.095 (SGNU). Multiplicity: each claim of B1, B1-NL, B3 and the extension gets one claim-level p (the k-th smallest Holm p for "k of m families" rules, the best swap for swap rules, the cluster permutation for tracking, the plan's Holm p for B3); Holm over the 21 claims at family-wise 0.05 keeps all but B1-NL H3-SGN ($p = 0.108$); the closest survivors are B1 H3 (Holm 0.0079), B1-NL H3-NL and HX2 (both 0.033). Several claim p are at resolution floors (10^{-4} for permutations, 10^{-5} for B3's sign flips).

Detectability calculation ("oracle"; post hoc). For every B1 test graph, parameter set and horizon we propagated the exact window covariance (G1: $\Sigma_{t+1} = A^\top \Sigma_t A + D$ from $\Sigma_1 = D$; G2: after the same 100 burn-in steps). Population relatives ($|\rho| > 10^{-20}$) reproduce the exact s_ρ of Theorems 2 and 5 in all 2,520 cases. The oracle replaces each sample correlation in rel-SortnRegress's count by the probability that the Fisher-z statistic at $n = 1000$ exceeds the algorithm's threshold ($\varepsilon = 0.062$) and predicts each edge's order accuracy as $P(c_i < c_j) + \frac{1}{2}P(c_i = c_j)$, treating the indicators as independent. It uses no simulated data. Against the observed G1−G2 gain: cell level Spearman 0.99, $R^2 = 0.97$, slope 1.03, intercept −0.010; unit level $R^2 = 0.61$ against 0.34 for the exact gap (difference +0.27 $[0.23, 0.30]$); in a joint regression the gap's coefficient is −0.063 over cells and −0.039 over graph–horizon units. Predicted G1 / G2 levels: chain 0.756/0.528 (observed 0.749/0.532), ring 0.702/0.503 (0.687/0.507), in-tree 0.746/0.677 (0.743/0.673). Without slice-1 edges the non-SC gain is predicted to fall from +0.130 to +0.096 (observed +0.125 to +0.089). What this shows and

does not show: the calculation is the expectation of rel-SortnRegress under the exact covariances of each generator, so the agreement says that the simulated gain has the size those covariances imply and that the independence approximation is adequate. It uses the full start-dependent covariances, so it includes every way in which the start changes pairwise correlations: pairs that are relatives under the fixed start only (the mechanism of Section 3.2) and pairs that are relatives under both starts but differ in correlation strength (a transient that needs no graph argument). We did not decompose the predicted gain into these two parts, so the calculation does not show that the gain is due to the first. The exact gap therefore gives a ranking and a rough scale; all 42 cells have the same sign, so the sign tests nothing.

Generator survey. We read the source code of tigramite 5.2.10.1 (`structural_causal_process`: zeros plus noise, burn-in $\lfloor 0.2T \rfloor$; `var_process`: i.i.d. $N(0,1)$ start, no burn-in), the CDML generator (Lawrence et al., 2021) (zeros, $2\tau_{\max}$ steps dropped), CausalTime (Cheng et al., 2024) (rollouts of 20 steps from a real-data window, no burn-in, in the demo), UUMCdata (Herman et al., 2025) (exact stationary start), causalnex 0.12.1 (zeros, burn-in 100 or more; data with lag labels), the NeuralGC VAR and Lorenz-96 simulators (Tank et al., 2021) (burn-in 100 and 1000) and gCastle 1.0.4 (no continuous time-series simulator; its Hawkes simulator starts with an empty history). cdt, causaldag, causal-learn and lingam ship no time-series benchmark generator. Lohse and Wahl (2025) state a burn-in without its length. The time-unrolled construction of Reisach et al. (2026) starts at $t = 1$ and gives every slice to an i.i.d. method as a separate variable; its code is not public, so only the graph-level start is verified. CauseMe's documentation could not be opened with TLS verification, so we make no claim about it. File and line references are in `lit3/GENERATOR_SURVEY.md`. Reading of exposure: the burn-in lengths above can be compared with $\Delta\tau_{\max}$, which concerns the graph artefact only. The variance artefact needs the start to be far from the stationary distribution relative to the mixing time of the process, which depends on the spectral radius of the generated coefficients; we computed no mixing time for any of these generators, so we do not grade them on it. All of them except the time-unrolled construction produce a single record with time order for time-series methods, so neither artefact applies to them as shipped; they would be exposed only if windows from their first steps were given to a method that ignores time order.

Appendix U Controls: Plan, Design and Verdicts (b1y)

Plan. The controls of Section 5.2 were specified in a time-stamped plan (`b1y/PLAN.md`, commit `5c9f942`) with two dated addenda; it is not a registration and was not posted to a third party. The plan was committed before any simulated data set or learner score of its questions existed; what had been computed (graph-only quantities, unit tests, the exact burn-in computation) is listed in it. Thresholds, verdict rules, the table that maps outcomes to the paper's reading, seeds and the multiplicity family were not changed after the first scored run. Server-side push times could not be read (the activity feed lagged behind the pushes), so the order of commits is the only evidence, and push order cannot rule out that results were computed before a plan was written.

Design. Graphs, weights, noise and the two starts are those of B1 (heterogeneous noise, raw data, $n = 1000$ unless stated, time labels hidden by shuffling columns). Seeds and random streams are disjoint from all earlier studies (a unit test checks this). Primary metrics drop the edges that leave the first slice ($\mathrm{SHD}_{x1}/|E|$, OA_{x1}); full-window metrics are secondary. Units are summary graphs; intervals are 10,000-resample graph-cluster bootstrap intervals. Q1 (counterfactual starts and dose): 7 families, $p \in \{5, 10\}$, 30 graphs each, $T \in \{3, 5, 10\}$, 11 arms; dose study on rings, chains and ER(c=2) graphs with $n \in \{200, 1000, 5000, 25000\}$ and spectral-radius caps $\{0.5, 0.8, 0.95\}$. Q2 (variance-matched control): chains, in-trees, rings, ER(c=2), six (p, T) cells, 12 test graphs per (family, p), four arms (the two starts and two variance-rescaled arms). Q5 (baselines): 12 new graphs per (family, p) for rings, chains, in-trees and ER(c=2), 384 data sets. The real-data analysis (Q4) and the first-slice re-analysis (Q3) use committed data only.

Environment and learners. Addendum 1 moved the runs to Python 3.12 (Python 3.10 for DYNOTEARS, which uses `causalnex` 0.12.1; all versions pinned in `requirements.txt` and the environment lists) and added GOLEM-EV and GOLEM-NV, DYNOTEARS, GES, GRaSP and BOSS (GES from Tetrad 7.6.10 with its SEM-BIC penalty 2, because causal-learn's GES was too slow at that penalty), DirectLiNGAM and a pooled VAR-LiNGAM for replicated short trajectories: pooled least squares for the lag-one matrix on all transitions, DirectLiNGAM on the pooled residuals for the contemporaneous matrix, and the combination $(I - B_0)M$ as in Hyvärinen et al. (2010). It is not the packaged estimator, which expects one long series, and it is a time-labelled reference only. A unit test recovers a known SVAR with non-Gaussian noise.

Deviations (all recorded). (1) Addendum 2: NOTEARS runs only for $d \leq 30$ after measured fit times on burn-in data (about 150 s per fit at $d = 30$, against 3.5 s on fixed-start data); the pairing of NOTEARS and GOLEM with FLOP uses only the cells where they ran. (2) The first tuning run was stopped; two shards lost their CSV header and were repaired, one truncated shard was recomputed. (3) SID is set to missing for estimates with self-loops or cycles, after one such output made `gadjid` panic; SHD and the other metrics are unaffected. (4) The circular block bootstrap for E1 was not run: its 15 wind-tunnel units are independent runs, and the light-tunnel replication is ten segments of one file and is reported without inference. (5) The `secs` column double-counts GOLEM time, so summed fit time overstates the cost; the whole study took about 21 CPU-hours of a 25-hour cap. (6) Server push times were not obtainable. (7) For the river data the earlier text quoted the station counts of an earlier subgraph set; the corrected counts are in Table 6.

Exact burn-in computation (Q6). The results quoted in Remark 20 come from `results/q6_burnin_exact.txt`, an exact enumeration of window relatives that does not use the theorems of Section 3.2; the reference burn-in was $\Delta + 2p$.

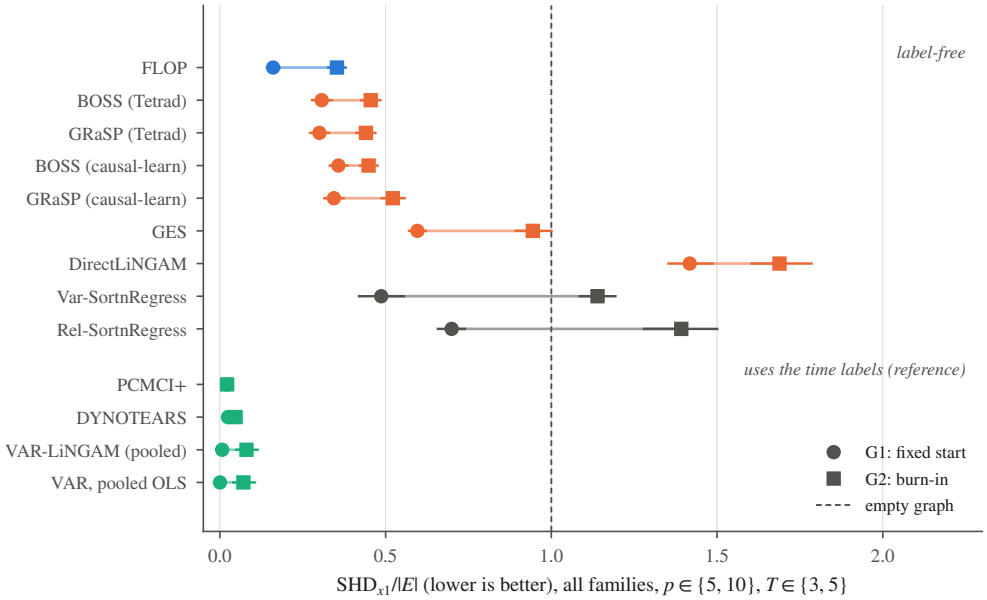


Fig. U4 Q5: label-free methods (orange, blue, grey) and references that use the time labels (green) under the fixed start (circles) and a burn-in (squares); SHD per true edge without first-slice edges, all families, $p \in \{5, 10\}$, $T \in \{3, 5\}$, 95% graph-cluster intervals (`figs/fig_learners.py`, from `b1y/results/q5.csv.gz`). FLOP is the best label-free method in 7 of the 8 (family, start) groups; every label-free method is worse than the labelled references, and only the SortnRegress variants and DirectLiNGAM get worse than the empty graph (dashed) under the burn-in.

Appendix V B3: Discriminating Test of the Source-SCC Rule on Held-Out Data

Process. The plan (`b3/PLAN.md`, with an OSF-formatted copy) was pushed before any byte of the held-out data was read; zip CRC-32 values were taken from the remote archive directory before the plan and SHA-256 hashes at download. The plan is internally time-stamped only: it was not posted to an external registry before data access, and results were committed 3–4 minutes after the plan. A disclosed development check on study-3 data from the same device had shown every colour edge to be detectable, so the primary outcome was expected.

Data. Causal Chambers experiments not previously downloaded by the project. H1 (primary): `lt_camera_walks_v1`, experiments `ar_1_uniform_{ref, pol_1_90, pol_12_30}`: five independent stationary AR(1) actuators drive the light sensors, so the colour block is one weak component with three source SCCs; $3 \times 30{,}000$ steps, 30 units of 3,000. H2 (co-primary): `lt_camera_v1/scm_5_red` and `scm_5_blue`, i.i.d. data from the collider R→G←B, 20 units. Secondary: `scm_5_green`, `scm_4_green`, `scm_4_blue`, and `scm_4_red` as a single-source control.

Test. E1's estimator unchanged. Competitors: constant 1/2, the weak-component-size rule, the theorem on the design graph G_D, and E1's eight naive rules. Primary set D^*: instances where the theorem and the weak-component rule disagree (56% of

instances for H1, 47% for H2, from graph-only computation before the plan); on D^* the prediction is the same on G_{lin} and G_D. Failure of the estimator pushes observed values towards the tie rules, so a pass on D^* cannot come from ties. Sign-flip tests over units, Holm, family $\alpha = 0.025$; pass requires $\Delta \leq -0.10$ against every competitor that disagrees on D^*.

Results. H1: on 2,361 instances, counts sorted as predicted in 2,356; $\Delta = -0.498$ against both tie rules, 30/30 units, $p \approx 10^{-5}$; non-relatives declared dependent 0.054, relatives detected 0.990. H2: $\Delta = -0.50$ against the tie rules and the sensor indicator, 20/20 units; R→G sorted 264/264 when B is observed and tied 123/123 when not. S1 (predicted tie, detected edge; 1,251 instances): theorem beats the eight naive rules ($\Delta = -0.43$, 30/30 units) and the theorem on G_D narrowly ($\Delta = -0.21$, 21/30 units, Holm $p = 0.021$). No stop condition met.

Failures. `scm_4_blue`: the tie rule is significantly better ($\Delta = +0.115$, 0/10 units, reverse $p = 0.001$); post hoc, the intervened blue channel ranges over 1–13 of 255 levels, so none of its edges is detected, and the pooled detection prerequisite (0.837) did not catch this. The pooled secondary analysis still passes, in 20/30 units. In `ar_1_uniform_pol_1_90` the polariser linear-zero refinement of Appendix L fails (polarisers correlate with ir_3, vis_3 at $|r| \approx 0.2$–0.28; G_D matches 0.959 against 0.041 for G_{lin}); H1 does not depend on it. Deviations: the inputs of two experiments do not reproduce from the published generator seeds (ranges match; not gating).

Appendix W Proofs for Sections 3.3 and 3.4

Proposition 8 is proved in the main text. The proofs below are paper proofs; none is machine-checked. Notation as in Section 3.4.

Proof of Proposition 9 (a) Every edge of G_T goes from a slice t to slice $t+1$, so a directed path ending in $v(t)$ has at most $t-1$ edges, and the self-chain $v(1) \to \cdots \to v(t)$ has exactly $t-1$. Graph isomorphisms preserve height, so two labellings of the same column DAG as unrollings with self-loops put the same columns in slice t, and T is the largest height and p the number of columns of height one. (b) The column DAG of the ring $i \to i+1$ with self-loops has the edges $(i,t) \to (i,t{+}1)$ and $(i,t) \to (i{+}1,t{+}1)$. With the new unit label $i' = i - t \pmod p$ the first family reads $i' \to i' - 1$ and the second $i' \to i'$, which is the unrolling of the reversed ring with self-loops. (c) Without loops the unrolling of $0 \to 1$ has the edges $0(t) \to 1(t{+}1)$, $t < T$, which are pairwise disjoint; every one of them has a tail of height one and a head of height two. □

Proof of Proposition 10 Two parents $i(t)$ and $k(t)$ of a node $j(t{+}1)$ lie in the same slice and are not adjacent, so every node with at least two parents is an unshielded collider, and all its incoming edges have the same direction in every member of the class. The remaining edges are those into nodes with exactly one parent; since every node $j(t{+}1)$ has the parent $j(t)$, these are the self-chain edges of loop-only units. Fix such a unit j and a member of the class. On the chain $j(1), \ldots, j(T)$ the member has no collider (a collider at a chain node would be unshielded, as chain neighbours at distance two are not adjacent, and G_T has none there), so the chain is oriented away from one of its nodes $j(s)$. Conversely, re-rooting the chain at any $j(s)$ keeps the skeleton, creates no collider (every chain node still has at most one parent, because a loop-only unit has no parent outside its chain), removes none, and creates no cycle

(a directed cycle would have to enter the chain from outside, and no edge does). The chains of different loop-only units are disjoint and can be re-rooted independently, which gives T^{m_0} members. □

Proof of Proposition 11 (a) Let G' be the column DAG of a sparsest representation. Σ is Markov to G', because the representation is a linear structural equation model on G' with independent noise terms (the start covariance is diagonal under the fixed-start law). G' has $(T-1)|\operatorname{supp} A'|$ edges, at most as many as G_T, because the true parameter is itself a representation. As the distribution is faithful to G_T, a DAG to which it is Markov has at least as many edges as G_T, with equality only if it is Markov equivalent to G_T (Chickering, 2002, Theorem 4). By Proposition 10, G' is G_T with the chains of some loop-only units re-rooted. G' is also an unrolling with self-loops, so by Proposition 9(a) every edge of G' joins consecutive heights, and every node of height at least two has a self-parent, with different nodes having different self-parents. Re-root the chain of a loop-only unit j at $j(s)$. If $1 < s < T$, then $j(s)$ has no parent in G', and $j(s-1)$ and $j(s+1)$ both have $j(s)$ as their only parent, so both would have the self-parent $j(s)$: impossible. If $s = T$, then $j(1)$ has height T in G'; if j has a child $k \neq j$ in G_S, the edge $j(1) \to k(2)$ is an edge of G' (it enters a node with two parents or more, namely $j(1)$ and $k(1)$), which gives $k(2)$ the height $T+1$: impossible. So only isolated units can be reversed, and only as a whole. (b) Let v_t be the variance of $x_j(t)$, $v_{t+1} = a^2 v_t + d$. The reversed chain is a representation iff the regression of $x_j(t)$ on $x_j(t+1)$ has the same coefficient $a v_t / v_{t+1}$ and the same residual variance $d\, v_t / v_{t+1}$ for all $t \in \{1, \dots, T-1\}$, i.e. iff $v_t / v_{t+1} = r$ does not depend on t. For $T = 2$ there is one transition and nothing to compare. For $T \geq 3$, $v_3 - v_2 = a^2(v_2 - v_1)$ gives $\frac{1}{r}(\frac{1}{r} - 1) = a^2(\frac{1}{r} - 1)$, so $r = 1$, which is stationarity, or $1/r = a^2$, which forces $d = 0$ and is excluded. If $r = 1$ for the first two transitions then v_t is constant for all t. (c) follows from (a) and (b). □

Proof of Proposition 12 (a) In the truth, $a_2 = \alpha a_1 + \varepsilon_a$ and $b_2 = \beta a_1 + \gamma b_1 + \varepsilon_b$, with an arbitrary covariance of (a_1, b_1). Take the units $a' = (a_1, a_2)$ and $b' = (b_2, b_1)$, i.e. the first slice $\{a_1, b_2\}$ and the second slice $\{a_2, b_1\}$. Then $a_2 = \alpha a_1 + 0 \cdot b_2 + \varepsilon_a$ and $b_1 = \beta' a_1 + \gamma' b_2 + \eta$, where η is the residual of the linear regression of b_1 on (a_1, b_2). By construction η is uncorrelated with the new first slice; it is a linear function of $(a_1, b_1, \varepsilon_b)$ and therefore uncorrelated with ε_a, and ε_a is uncorrelated with (a_1, b_2). So the new innovation covariance is diagonal and uncorrelated with the new first slice, whose covariance is free. The new self-loop coefficients are $\alpha \neq 0$ and γ', which is non-zero whenever $\gamma \neq 0$ (the partial correlation of b_1 and b_2 given a_1 is non-zero); the support of the new transition matrix has three entries when $\beta' \neq 0$ and two otherwise, so the new representation has no more edges than the truth. (b) A stationary Gaussian AR(1) process has a symmetric autocovariance, so the reversed window has the same covariance, and an isolated unit is uncorrelated with all other columns. (c) $\Gamma = \sigma^2 (I - A^2)^{-1}$ solves $\Gamma = A \Gamma A^\top + \sigma^2 I$ and commutes with A. The time-reversed stationary process has transition matrix $\Gamma A^\top \Gamma^{-1} = A$ and innovation covariance $\Gamma - A \Gamma A^\top = \sigma^2 I$, so the reversed window is a representation with the same support. □

On Conjecture 13. Let H be G_T plus a complete DAG on the first slice. Under the stated faithfulness, the argument of Proposition 11(a) shows that the DAG of a sparsest representation (its unrolling plus a complete DAG on its first slice) is Markov equivalent to H, hence has the same skeleton U. In H the edges between slices 1 and 2 are not fixed by colliders, because the parents of a slice-2 node lie in the first-slice clique. If the only partition of the nodes of U into T layers of size p, with edges inside

layer 1 or between consecutive layers and a perfect matching between consecutive layers, is the true one, the conjecture follows; the reversed layering is excluded when slice T is an independent set of U and slice 1 is not. We checked uniqueness by hand for the summary edge $a \to b$ with self-loops and $T = 3$; Proposition 12(a) shows that it fails for $T = 2$. We have no proof for general graphs, and the case in which the first slice is not a clique in U is not covered.

References

Afyouni, S., Smith, S.M., Nichols, T.E. (2019). Effective degrees of freedom of the Pearson's correlation coefficient under autocorrelation. *NeuroImage*, *199*, 609–625,

Andrews, B., & Kummerfeld, E. (2026). Better simulations for validating causal discovery with the DAG-adaptation of the onion method. *Journal of Machine Learning Research*, *27*(147), 1–42, (arXiv:2405.13100)

Arcones, M.A. (1994). Limit theorems for nonlinear functionals of a stationary Gaussian sequence of vectors. *The Annals of Probability*, *22*(4), 2242–2274, https://doi.org/10.1214/aop/1176988503

Assaad, C.K., Devijver, E., Gaussier, E. (2022). Survey and evaluation of causal discovery methods for time series. *Journal of Artificial Intelligence Research*, *73*, 767–819, https://doi.org/10.1613/jair.1.13428

Bartlett, M.S. (1935). Some aspects of the time-correlation problem in regard to tests of significance. *Journal of the Royal Statistical Society*, *98*(3), 536–543,

Bello, K., Aragam, B., Ravikumar, P. (2022). DAGMA: Learning DAGs via M-matrices and a log-determinant acyclicity characterization. *Advances in neural information processing systems* (Vol. 35).

Boege, T., Kubjas, K., Misra, P., Solus, L. (2025). Colored Gaussian directed acyclic graphical models. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, , https://doi.org/10.1093/jrsssb/qkaf068 (arXiv:2404.04024)

Bollen, K.A. (1989). *Structural equations with latent variables*. New York: John Wiley & Sons.

Breitung, J., & Swanson, N.R. (2002). Temporal aggregation and spurious instantaneous causality in multiple time series models. *Journal of Time Series Analysis*,

23(6), 651–665, https://doi.org/10.1111/1467-9892.00284

Brockwell, P.J., & Davis, R.A. (1991). *Time series: Theory and methods* (Second ed.). New York: Springer.

Brouillard, P., Squires, C., Wahl, J., Kording, K.P., Sachs, K., Drouin, A., Sridhar, D. (2025). The landscape of causal discovery data: Grounding causal discovery in real-world applications. *Proceedings of the fourth conference on causal learning and reasoning* (Vol. 275). (arXiv:2412.01953)

Cheng, Y., Wang, Z., Xiao, T., Zhong, Q., Suo, J., He, K. (2024). CausalTime: Realistically generated time-series for benchmarking of causal discovery. *International conference on learning representations.*

Chickering, D.M. (2002). Optimal structure identification with greedy search. *Journal of Machine Learning Research*, *3*, 507–554,

Clifford, P., Richardson, S., Hémon, D. (1989). Assessing the significance of the correlation between two spatial processes. *Biometrics*, *45*(1), 123–134,

Dang, Q.-V., Le, D., Dinh, M.N. (2026). SPADE: Spline additive-noise DAG estimation for interpretable nonlinear causal discovery in managerial decision support. *Array*, *32*, 101335, https://doi.org/10.1016/j.array.2026.101335

de Moura, L., & Ullrich, S. (2021). The Lean 4 theorem prover and programming language. *Automated deduction – cade 28* (pp. 625–635).

Dorogovtsev, S.N., Mendes, J.F.F., Samukhin, A.N. (2001). Giant strongly connected component of directed networks. *Physical Review E*, *64*(2), 025101,

Dovgal, S., de Panafieu, É., Ralaivaosaona, D., Rasendrahasina, V., Wagner, S. (2020). *The birth of the strong components.*

Friedman, N., Murphy, K., Russell, S. (1998). Learning the structure of dynamic probabilistic networks. *Proceedings of the 14th conference on uncertainty in artificial intelligence (uai).* (arXiv:1301.7374)

Gamella, J.L., Peters, J., Bühlmann, P. (2025). Causal chambers as a real-world physical testbed for AI methodology. *Nature Machine Intelligence*, *7*(1), 107–118,

Gentzel, A., Garant, D., Jensen, D. (2019). The case for evaluating causal models using interventional measures and empirical data. *Advances in neural information processing systems* (Vol. 32, pp. 11722–11732). (arXiv:1910.05387)

Ginovyan, M.S., & Taqqu, M.S. (2021). *A survey on limit theorems for Toeplitz type quadratic functionals of stationary processes and applications.*

Gong, M., Zhang, K., Schölkopf, B., Glymour, C., Tao, D. (2017). Causal discovery from temporally aggregated time series. *Proceedings of the 33rd conference on uncertainty in artificial intelligence.*

Gong, M., Zhang, K., Schölkopf, B., Tao, D., Geiger, P. (2015). Discovering temporal causal relations from subsampled data. *Proceedings of the 32nd international conference on machine learning* (Vol. 37, pp. 1898–1906).

Granger, C.W.J. (1969). Investigating causal relations by econometric models and cross-spectral methods. *Econometrica*, *37*(3), 424–438, https://doi.org/10.2307/1912791

Guan, V., Janssen, J., Rahmani, H., Warren, A., Zhang, S., Robeva, E., Schiebinger, G. (2026). Identifying drift, diffusion, and causal structure from temporal snapshots. *Journal of Machine Learning Research*, *27*, 1–56, (arXiv:2410.22729)

Hamilton, J.D. (1994). *Time series analysis.* Princeton University Press.

Haugh, L.D. (1976). Checking the independence of two covariance-stationary time series: A univariate residual cross-correlation approach. *Journal of the American Statistical Association*, *71*(354), 378–385, https://doi.org/10.1080/01621459.1976.10480353

Henckel, L., Würtzen, T., Weichwald, S. (2024). Adjustment identification distance: A gadjid for causal structure learning. *Proceedings of the fortieth conference on uncertainty in artificial intelligence* (Vol. 244, pp. 1569–1598).

Herman, R.J., Wahl, J., Ninad, U., Runge, J. (2025). Unitless unrestricted Markov-consistent SCM generation: Better benchmark datasets for causal discovery. *Proceedings of the fourth conference on causal learning and reasoning* (Vol. 275, pp. 1506–1531).

Hsu, D., Kakade, S.M., Zhang, T. (2012). A tail inequality for quadratic forms of subgaussian random vectors. *Electronic Communications in Probability*, *17*, , (arXiv:1110.2842)

Hyttinen, A., Plis, S., Järvisalo, M., Eberhardt, F., Danks, D. (2016). Causal discovery from subsampled time series data by constraint optimization. *Proceedings of the eighth international conference on probabilistic graphical models* (Vol. 52, pp. 216–227). (arXiv:1602.07970)

Hyvärinen, A., Zhang, K., Shimizu, S., Hoyer, P.O. (2010). Estimation of a structural vector autoregression model using non-Gaussianity. *Journal of Machine Learning Research*, *11*(56), 1709–1731,

Kaiser, M., & Sipos, M. (2022). Unsuitability of NOTEARS for causal graph discovery when dealing with dimensional quantities. *Neural Processing Letters*, *54*(3), 1587–1595, https://doi.org/10.1007/s11063-021-10694-5

Kilian, L., & Lütkepohl, H. (2017). *Structural vector autoregressive analysis.* Cambridge University Press.

Lanne, M., & Lütkepohl, H. (2008). Identifying monetary policy shocks via changes in volatility. *Journal of Money, Credit and Banking*, *40*, 1131–1149, https://doi.org/10.1111/j.1538-4616.2008.00151.x

Laurent, B., & Massart, P. (2000). Adaptive estimation of a quadratic functional by model selection. *The Annals of Statistics*, *28*(5), 1302–1338,

Lawrence, A.R., Kaiser, M., Sampaio, R., Sipos, M. (2021). *Data generating process to evaluate causal discovery techniques for time series data.* arXiv:2104.08043; Causal Discovery & Causality-Inspired Machine Learning Workshop at NeurIPS 2020.

Lewis, D.J. (2021). Identifying shocks via time-varying volatility. *The Review of Economic Studies*, *88*, 3086–3124, https://doi.org/10.1093/restud/rdab009

Lohse, C., & Wahl, J. (2025). Sortability of time series data. *Transactions on Machine Learning Research*, , (arXiv:2407.13313)

Lütkepohl, H. (2005). *New introduction to multiple time series analysis.* Berlin: Springer.

Montagna, F., Mastakouri, A.A., Eulig, E., Noceti, N., Rosasco, L., Janzing, D., ... Locatello, F. (2023). Assumption violations in causal discovery and the robustness of score matching. *Advances in neural information processing systems* (Vol. 36).

Mooij, J.M., Janzing, D., Schölkopf, B. (2013). From ordinary differential equations to structural causal models: the deterministic case. *Proceedings of the 29th conference on uncertainty in artificial intelligence.* (arXiv:1304.7920)

Müller, U.K., & Elliott, G. (2003). Tests for unit roots and the initial condition. *Econometrica*, *71*(4), 1269–1286, https://doi.org/10.1111/1468-0262.00447

Nagarajan, R., & Scutari, M. (2025). *Benchmarking constraint-based Bayesian structure learning algorithms: Role of network topology.* (arXiv:2501.02019)

Ng, I., Ghassami, A., Zhang, K. (2020). On the role of sparsity and DAG constraints for learning linear DAGs. *Advances in neural information processing systems* (Vol. 33).

Nourdin, I., Peccati, G., Podolskij, M. (2011). Quantitative Breuer–Major theorems. *Stochastic Processes and their Applications*, *121*(4), 793–812, https://doi.org/10.1016/j.spa.2010.12.006

Ormaniec, W., Sussex, S., Lorch, L., Schölkopf, B., Krause, A. (2025). Standardizing structural causal models. *International conference on learning representations.*

Pamfil, R., Sriwattanaworachai, N., Desai, S., Pilgerstorfer, P., Georgatzis, K., Beaumont, P., Aragam, B. (2020). DYNOTEARS: Structure learning from time-series data. *Proceedings of the twenty third international conference on artificial intelligence and statistics* (Vol. 108, pp. 1595–1605).

Peters, J., Janzing, D., Gretton, A., Schölkopf, B. (2009). Detecting the direction of causal time series. *Proceedings of the 26th annual international conference on machine learning (icml)* (pp. 801–808).

Reisach, A.G., Chambaz, A., Blanchard, G., Weichwald, S. (2026). *A topological sorting criterion for random causal directed acyclic graphs.* arXiv:2605.06288.

Reisach, A.G., Seiler, C., Weichwald, S. (2021). Beware of the simulated DAG! Causal discovery benchmarks may be easy to game. *Advances in neural information processing systems* (Vol. 34).

Reisach, A.G., Tami, M., Seiler, C., Chambaz, A., Weichwald, S. (2023). A scale-invariant sorting criterion to find a causal order in additive noise models. *Advances in neural information processing systems* (Vol. 36).

Rigobon, R. (2003). Identification through heteroskedasticity. *Review of Economics and Statistics*, *85*, 777–792, https://doi.org/10.1162/003465303772815727

Roy, R. (1989). Asymptotic covariance structure of serial correlations in multivariate time series. *Biometrika*, *76*(4), 824–827,

Rudelson, M., & Vershynin, R. (2013). *Hanson–Wright inequality and sub-Gaussian concentration.* (arXiv:1306.2872)

Runge, J. (2020). Discovering contemporaneous and lagged causal relations in autocorrelated nonlinear time series datasets. *Proceedings of the 36th conference on uncertainty in artificial intelligence* (Vol. 124, pp. 1388–1397).

Runge, J. (2026). *Tigramite: causal inference for time series.* Python package version 5.2.10.1, https://github.com/jakobrunge/tigramite. (Released 2026-01-14 on PyPI)

Runge, J., Bathiany, S., Bollt, E., Camps-Valls, G., Coumou, D., Deyle, E., ... Zscheischler, J. (2019). Inferring causation from time series in Earth system sciences. *Nature Communications*, *10*(1), 2553, https://doi.org/10.1038/s41467-019-10105-3

Runge, J., Nowack, P., Kretschmer, M., Flaxman, S., Sejdinovic, D. (2019). Detecting and quantifying causal associations in large nonlinear time series datasets. *Science Advances*, *5*(11), eaau4996,

Spirtes, P., Glymour, C., Scheines, R. (2000). *Causation, prediction, and search* (2nd ed.). Cambridge, MA: MIT Press.

Stein, G., Shadaydeh, M., Blunk, J., Penzel, N., Denzler, J. (2025). CausalRivers – scaling up benchmarking of causal discovery for real-world time-series. *International conference on learning representations.*

Tan, J., & Syrgkanis, V. (2026). *CausalSmith: A formally grounded, self-improving agentic framework for automated research in causal inference.*

Tank, A., Covert, I., Foti, N., Shojaie, A., Fox, E.B. (2021). Neural Granger causality. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, , https://doi.org/10.1109/TPAMI.2021.3065601 (Early access; arXiv:1802.05842)

The mathlib Community (2020). The Lean mathematical library. *Proceedings of the 9th acm sigplan international conference on certified programs and proofs* (pp. 367–381).

van der Hofstad, R. (2016). *Random graphs and complex networks, volume 1.* Cambridge University Press.

Vowels, M.J., Camgoz, N.C., Bowden, R. (2022). D'ya like DAGs? A survey on structure learning and causal discovery. *ACM Computing Surveys*, *55*(4), 1–36, https://doi.org/10.1145/3527154

Wang, Y., Geng, X., Huang, W., Huang, B., Gong, M. (2023). Generator identification for linear SDEs with additive and multiplicative noise. *arXiv preprint arXiv:2310.19491*, ,

Weichwald, S., Jakobsen, M.E., Mogensen, P.B., Petersen, L., Thams, N., Varando, G. (2020). Causal structure learning from time series: Large regression coefficients may predict causal links better in practice than small p-values. *Proceedings of the neurips 2019 competition and demonstration track* (Vol. 123, pp. 27–36).

Weiss, A.A. (1984). Systematic sampling and temporal aggregation in time series models. *Journal of Econometrics*, *26*(3), 271–281, https://doi.org/10.1016/0304-4076(84)90022-8

Weiss, G. (1975). Time-reversibility of linear stochastic processes. *Journal of Applied Probability*, *12*, 831–836, https://doi.org/10.2307/3212735

Westfall, J., & Yarkoni, T. (2016). Statistically controlling for confounding constructs is harder than you think. *PLOS ONE*, *11*(3), e0152719,

Wienöbst, M., Henckel, L., Weichwald, S. (2026). Embracing discrete search: A reasonable approach to causal structure learning. *International conference on learning representations.* (arXiv:2510.04970)

Zheng, X., Aragam, B., Ravikumar, P., Xing, E.P. (2018). DAGs with NO TEARS: Continuous optimization for structure learning. *Advances in neural information processing systems* (Vol. 31).

Zheng, Y., Huang, B., Chen, W., Ramsey, J., Gong, M., Cai, R., ... Zhang, K. (2024). Causal-learn: Causal discovery in Python. *Journal of Machine Learning Research*, *25*(60), 1–8,

Table 1 Status of the claims of this paper. "Lean" means a theorem proved in Lean 4 (the map from Lean statements to paper statements has not been checked by a human); "paper proof" means a proof in the appendix that no human has read line by line; "enumerated" means exact computation on many instances, not a proof.

claim	evidence	status
exact per-edge value, limit, necessity of the condition (Q1)	proof; Lean; 46,724 enumerated cases	proved; Lean-checked
stationary value after burn-in $B \geq \Delta$ (Q2)	proof; Lean	proved; Lean-checked
sharpness of $B \geq \Delta$ (Q2)	Lean for one edge; paper proof for the relatives of chains; exact enumeration for chains ($p \leq 8$), rings and 396 random graphs; 445 of 726 random graphs fail at $\Delta - 1$	Lean for $\Delta = 1$; proved on paper for chain relatives; enumerated for chain s_ρ; not sharp in general
limits for random summary graphs (Q2)	paper proof; numerics	paper proof, unread by a probabilist
slices, equivalence class, non-identifiability (Q1)	paper proofs; small-case checks	paper proofs, conditional on faithfulness
effective sample size, saturation, conditioning (Q3)	paper proofs; simulations; algebra in Lean	paper proofs, asymptotic; bounds loose
fixed start raises rel-sortability (Q2, Q4)	simulation, 420 graphs, replicated	supported, positive in all 42 conditions
exact gap and the size of the gain (Q4)	post hoc detectability calculation; designed counterfactual and dose tests (Section 5.2)	the realised effect follows detectable correlation; the exact gap is the large-n ceiling (0.76 of it at $n = 25{,}000$)
DAGMA and NOTEARS overtake FLOP under a fixed start (Q4)	learner extension, 96 graphs; tuned, variance-matched control, 288 graph-seeds	observed and survives tuning; variance channel by the plan's rule but not sufficient; mechanism not explained; plan predicted otherwise
relatives count orients edges in real data (Q3, Q4)	CausalRivers; two Causal Chambers devices (P7–P9, E1, E2, B3)	mostly failed or non-discriminating

Table 2 Rel-sortability of 3-node summary graphs (with self-loops, $|E_S| = 5$). Finite: $\lim_T s_\rho(G_T)$ (Theorem 2); stationary: s_ρ^W (Theorem 5). Lean: chain3_limit, chain3_stat, fork_limit, fork_stat, collider_limit, collider_stat; 3-cycle by reisach_theorem4.

summary graph	finite G_T (limit)	stationary
chain $a \to b \to c$	0.7	0.5
fork $a \leftarrow b \to c$	0.7	0.5
collider $a \to c \leftarrow b$	0.7	0.7
3-cycle (6 edges)	0.5	0.5

Table 3 B1, heterogeneous noise, raw data; means over 180 data sets per family. OA: order accuracy of rel- and Var-SortnRegress, G1 / G2. ΔOA: G1−G2 for rel, with 95% graph-cluster bootstrap CI. SC: strongly connected.

family	exact s_ρ	rel OA	ΔOA rel	var OA
symER (SC)	.78/.50	.63/.50	+.126 [.115,.137]	.64/.50
ring (SC)	.95/.50	.69/.51	+.180 [.167,.193]	.73/.50
chain	.89/.50	.75/.53	+.217 [.209,.226]	.93/.70
out-tree	.84/.50	.69/.55	+.144 [.127,.161]	.94/.71
in-tree	.79/.66	.74/.67	+.069 [.057,.081]	.95/.72
ER(c=1)	.77/.56	.67/.59	+.085 [.068,.103]	.86/.65
ER(c=2)	.83/.52	.70/.59	+.109 [.092,.127]	.74/.59

Table 4 Competitive learners (Appendix T), raw data unless noted; 96 graphs, G1 / G2. OA: ancestral order accuracy; Δ: G1−G2 with graph-cluster 95% CI. GOLEM-NV at $p = 5$ only.

method	SHD/$\lvert E\rvert$	OA	ΔOA
FLOP	.19 / .52	.89 / .77	–
DAGMA	.35 / .90	.80 / .60	+.197 [.177,.216]
NOTEARS	.47 / .89	.74 / .59	+.156 [.136,.175]
GOLEM-NV	.39 / .80	.79 / .56	+.232 [.205,.258]
DAGMA, std.	.87 / .87	.35 / .41	−.064 [−.079,−.048]
NOTEARS, std.	.76 / .84	.49 / .46	+.036 [.022,.049]
in-trees: DAGMA	.08 / .87		
NOTEARS / FLOP	.13 / .76; .30 / .56		
PCMCI+ (labels)	.02 / .02	1.0 / 1.0	–

Table 5 Scoring without the edges that leave the first slice (post hoc re-analysis of committed B1 and learner-extension data; no new data). Gains are G1−G2 in order accuracy; shifts are G2−G1 in SHD/$|E|$; the last three rows (four methods) are the in-tree interactions of the rank change against FLOP (G1 value minus G2 value of [method − FLOP]). Graph-cluster 95% CIs for the first rows.

quantity	all edges	without slice-1 edges
rel-SortnRegress, OA gain (300 non-SC graphs)	$+0.125$ $[0.116, 0.134]$	$+0.089$ $[0.079, 0.100]$
R^2-SortnRegress, OA gain	$+0.075$	$+0.096$
Var-SortnRegress, OA gain	$+0.210$	$+0.195$
FLOP, SHD/$\|E\|$ shift (420 graphs)	$+0.197$ $[0.185, 0.209]$	$+0.093$ $[0.084, 0.102]$
rel-SortnRegress, SHD/$\|E\|$ shift	$+0.649$	$+0.531$
in-tree interaction vs FLOP, DAGMA	-0.524	-0.587
NOTEARS	-0.369	-0.462
GOLEM-EV / GOLEM-NV	-0.384 / -0.318	-0.484 / -0.421

Table 6 What the inference of each study rests on. "Units" are the units over which the reported intervals and p-values are computed.

study	units used for inference	what is independent
B1, B1-NL, learner extension (simulation)	summary graphs (cluster bootstrap)	graphs; each data set has $n = 1000$ replicate trajectories of a short window
CausalRivers (P1–P6)	120 subgraphs	900 station slots, 445 distinct stations, one station used up to 7 times (112 of the 120 subgraphs share a station with another); two regions share a confounding driver; the 12 h and 24 h columns are the same records
wind tunnel (P7–P9)	300 subgraphs, 15 runs	runs
E1	10 segments	one series (`citris_1`)
B3	30 units	three runs of 30,000 steps
E2	64 held-out regimes, 1,280 subgraphs	regimes within runs; dependence between regimes not examined

Table 7 Tests fixed in the time-stamped analysis plans. Pass/fail is by the point rule fixed in the plan. Brackets: 95% cluster-bootstrap CI. "non-rel. dep.": fraction of graph non-relative pairs declared dependent. "Acc.": orientation accuracy on edges predicted to be strictly sorted.

study (units)	test (threshold)	result (12 h / 24 h)	verdict
CausalRivers, marginal (120 subgraphs)	P1: confluence subgraphs have $\tilde{s} > 1/2$	112 of 114	fails as worded[a]
	P2: Acc. > 0.5, CI excl. 0.5	$0.501\,[0.492, 0.509]$ / $0.498\,[0.489, 0.507]$	**fail**
	P3: Spearman$(\tilde{s}, \text{acc.}) \geq 0.3$	-0.018 / 0.018	**fail**
	(non-rel. dep., nominal test)	0.984 / 0.978	stop condition met
CausalRivers, conditioned on regional proxy (120 subgraphs)	P4: non-rel. dep. < 0.5 given proxy	0.741 / 0.735	**fail**
	P5: Acc. > 0.5, CI excl. 0.5	$0.539\,[0.513, 0.565]$ / $0.540\,[0.514, 0.567]$	pass, not interp.[b]
	P6: Spearman ≥ 0.3	-0.113 / -0.121	**fail**
Causal Chambers, wind tunnel (15 runs, primary)	P7: non-rel. dep. < 0.5	$0.039\,[0.031, 0.047]$	pass
	P8: Acc. > 0.5, CI excl. 0.5	$0.656\,[0.637, 0.675]$	pass
	P9: Spearman ≥ 0.3 (300 subgraphs)	$0.267\,[0.163, 0.368]$	**fail**
light tunnel (10 segments, replication)	R7 / R8 / R9 (same thresholds)	0.000 / 1.000 / $0.779\,[0.696, 0.864]$	pass (replication)[c]
E1, light tunnel (10 units, G_{lin})	E1a: non-rel. dep. < 0.5; det. ≥ 0.8	0.000; $0.969\,[0.945, 0.992]$	pass
	E1b: $\Delta \leq -0.10$, Holm $p \leq 0.05$, all 8	$\Delta = -0.502$ to -0.533 (10/10)	pass (8/8)[d]
	E1c: as E1b, detected edges	$\Delta = -0.500$	pass (8/8)
E2, wind regimes (64 units; D=1, lags L=30 vs. 0)	Q7: non-rel. dep. < 0.5	$0.003\,[0.001, 0.005]$	pass
	Q8: Acc. > 0.5, CI excl. 0.5	$0.553\,[0.527, 0.579]$	pass
	Q9: Spearman ≥ 0.3 (1,280 subgr.)	$0.243\,[0.193, 0.294]$	**fail**
	Q10: detection gain, CI lower bound > 0	$+0.069\,[0.060, 0.079]$	pass
	E2 success (Q7–Q10 all)		**not met**

[a] The two exceptions are split-and-rejoin channels with a single source, for which the theorem itself predicts 1/2. [b] Stop condition (P4 failed at both resolutions): P5 is reported, not interpreted. [c] Fixed in the plan as not counting towards a positive demonstration (i.i.d. inputs, weak time-series aspect). [d] G_{lin} is the primary graph fixed in plan 4. It omits six real causal edges that, under a symmetry lemma and an assumption of ideal polariser sensors, carry no linear dependence. On the published graph the theorem still beats all eight overall, but not on its strictly ordered edges, and on G_{lin} the prediction equals a weak-component-size rule (Section 5.3, Appendix L).

Table 8 Out-of-sample F-theory check on the Causal Chambers: predictions from the first half of each unit, observed rates on the second half. Accuracy is predicted by Monte Carlo of independent pair decisions.

	wind tunnel		light tunnel	
	pred.	obs.	pred.	obs.
non-relatives declared (F2)	0.031	0.028	0.006	0.000
relatives detected (F1)	0.360	0.343	0.679	0.681
orientation accuracy	0.630	0.620	0.999	1.000

Table L1 E1 primary, `citris_1`, graph G_{lin}, lag 0, 1,662 instances. Δ [95% CI]; all rows: 10/10 units, Holm $p = 0.0078$. "Exact": share of instances whose predicted value equals the observed one. "Spearman": agreement of the mean predicted edge value with observed accuracy over 200 subgraphs.

predictor	Δ [CI]	Δ, detected edges	exact	Spearman
theorem	—	—	0.992	0.962
sensor indicator	-0.502 $[-0.507, -0.496]$	-0.500	0.203	0.408
ancestors, full graph	-0.515 $[-0.519, -0.510]$	-0.500	0.147	0.334
ancestors, subgraph	-0.509 $[-0.513, -0.505]$	-0.500	0.120	0.121
in-degree, full graph	-0.515 $[-0.520, -0.511]$	-0.500	0.180	0.349
in-degree, subgraph	-0.510 $[-0.514, -0.505]$	-0.500	0.143	0.105
depth, full graph	-0.515 $[-0.519, -0.510]$	-0.500	0.147	0.334
depth, subgraph	-0.510 $[-0.514, -0.505]$	-0.500	0.114	0.113
every edge 1	-0.533 $[-0.536, -0.530]$	-0.500	0.115	undefined

Table L2 E1 on i.i.d. structural-model sets (no verdict; secondary). Competitors' exact-match range over the eight predictors. Units: `scm_1` 5, the others 10 each (45 pooled).

set	instances	theorem exact	competitors exact	Δ range	result
`buchholz_1_obs`	1,712	0.791	0.328–0.375	-0.499 to -0.496	pass 8/8
`scm_1_reference`	867	1.000	0.314–0.392	-0.500	not met[†]
`scm_2_reference`	1,650	0.812	0.324–0.408	-0.443 to -0.426	pass 8/8
`scm_4_reference` (chain)	2,267	1.000	0.000–0.088	-0.500	pass 8/8
`scm_5_reference` (collider)	2,238	1.000	0.464–0.495	-0.500	pass 8/8
pooled (45 units)	8,734	0.923	0.283–0.322	-0.489 to -0.487	pass 8/8, Holm $p = 8 \times 10^{-5}$

[†]Five units: the smallest attainable Holm-adjusted p is $8/32 = 0.25$.

Table M3 E2, 64 held-out regimes, 95% unit-bootstrap CIs. Q8: accuracy on the 1,216 instances predicted 1. Q10: relatives detected, lagged minus lag 0 (non-relatives declared difference in the last column).

configuration	Q7	Q8	Q9	Q10	non-rel. diff.
primary D=1, L=30	0.003 [0.001, 0.005] pass	0.553 [0.527, 0.579] pass	0.243 [0.193, 0.294] **fail**	+0.069 [0.060, 0.079] pass	+0.001 [−0.001, 0.002]
secondary D=10, L=3	0.011 [0.008, 0.014]	0.711 [0.692, 0.729]	0.202 [0.152, 0.253] fail	−0.007 [−0.012, −0.003] fail	−0.002 [−0.004, −0.0003]
lag 0, D=10 (replication of study 3)	0.014 [0.010, 0.018]	0.742 [0.722, 0.760]	0.202 [0.145, 0.259] fail	—	—
lag 0, D=1	0.002 [0.001, 0.004]	0.552 [0.532, 0.570]	0.120 [0.066, 0.176] fail	—	—

Table M4 Detection rate per edge over 64 regimes: share of regimes in which the pair was declared dependent.

edge	D=1, L=0	D=1, L=30	D=10, L=0
load_in → rpm_in	0.859	0.875	1.000
load_in → rpm_out	0.000	0.172	0.094
load_in → current_in	1.000	1.000	1.000
load_in → current_out	0.000	0.000	0.047
load_out → rpm_in	0.000	0.016	0.000
load_out → rpm_out	0.844	0.969	1.000
load_out → current_out	1.000	1.000	1.000
load_out → current_in	0.000	0.000	0.000
load_in → pressure_intake	0.000	0.094	0.953
load_out → pressure_intake	0.000	0.031	0.906
load_in → pressure_upwind	0.562	0.984	1.000
load_out → pressure_upwind	0.016	0.812	1.000
load_in → pressure_downwind	0.734	0.953	1.000
load_out → pressure_downwind	0.750	1.000	1.000
pot_1 → signal_1	0.359	0.156	0.969
pot_1 → signal_2	0.266	0.109	0.969
pot_2 → signal_2	0.000	0.000	0.266
pot_1, load_in, load_out → mic	0.000	0.000	0.016 (pot_1), 0.000

Table M5 Out-of-sample F-theory checks: the first half of each unit predicts the second half. Predicted detection is a bracket $[\bar{p}_{\max}, \bar{p}_{\mathrm{ind}}]$ when lags are used. Pass if the observed value lies within 0.03 of the bracket (criterion fixed in plan 4).

configuration	predicted detection	observed	non-rel. declared, pred. / obs.	accuracy, pred. / obs.
wind D=1, L=0	0.250	0.242	0.007 / 0.003	0.573 / 0.590
wind D=1, L=30	[0.344, 0.426]	0.332	[0.021, 0.074] / 0.003	0.504 / 0.506
wind D=10, L=0	0.548	0.533	0.014 / 0.008	0.649 / 0.632
wind D=10, L=3	[0.536, 0.554]	0.524	[0.016, 0.027] / 0.008	0.628 / 0.618
`citris_1`, L=0	0.935	0.945	0.007 / 0.000	undefined
`citris_1`, L=5	[0.940, 0.988]	0.945	[0.010, 0.045] / 0.000	undefined

Table N6 Loop-free examples (L: finite limit, S: stationary). Lean: edgeNoLoops_sRho, colliderNL_sRho, edgeLoopHead_sRho; the others are computed values (example scripts in the supplement).

summary graph	without loops	with loops
single edge $0 \to 1$	$s_\rho(G_T) = 1/2$ for all T	$L = 2/3$, $S = 1/2$
chain $0 \to 1 \to 2$	$L = 1/2$	$L = 0.7$
2-cycle $\{0, 1\}$ plus $0 \to 2$	$L = 1/2$	loop at 2 only: $L = 5/8$
collider	$L = S = 1$	$L = S = 0.7$
edge, loop at head only	$L = S = 1$	–

Table P7 Exact computation (means $\pm$ s.d. over seeds; SCCs by scipy and source sets by bitsets) against the limits of Proposition 45, proved in Appendix P.1.

c	p	L exact	L formula	S exact	S formula
0.5	20000	0.6666 ± 0.0011	0.6667	0.5658 ± 0.0014	0.5656
1.0	20000	0.7486 ± 0.0008	0.7500	0.6557 ± 0.0013	0.6580
1.5	50000	0.6975 ± 0.0020	0.6981	0.6159 ± 0.0018	0.6163
2.0	50000	0.6214 ± 0.0022	0.6217	0.5629 ± 0.0014	0.5632
3.0	100000	0.5434 ± 0.0010	0.5433	0.5215 ± 0.0006	0.5214
5.0	20000	0.5059 ± 0.0005	0.5058	0.5028 ± 0.0003	0.5029

Table U8 Verdicts of the controls plan, by the rules fixed in advance (metric OA_{x1} for Q1, $SHD_{x1}/|E|$ for Q2 and Q5; graph-cluster 95% intervals). A: $SHD_{x1}/|E|$ of FLOP minus that of the method, so positive means the method is better; R: ratios defined in the plan.

item	rule fixed in advance	result	verdict
V1	negative control: |difference| < 0.01	-0.004 $[-0.011, 0.002]$	valid
S1–S4	burn-in mixed (< 0.02); CF2 definable ($\leq 10\%$); detectability of $CF2_{m1}$ equals the fixed start's ($\leq 2\%$ of pairs flip in $\geq 90\%$ of cases); gain reproduces	0.001; 1.2%; 95.5%; $+0.090$ $[0.080, 0.101]$	all hold
V2	retention of $CF2_{m1}$: ≥ 0.75 with CI lower bound ≥ 0.50 means "detectable strength"; ≤ 0.35 with upper bound ≤ 0.60 means "population pattern"	0.955 $[0.884, 1.030]$; $CF2_{m3}$ 0.919, $CF2_{m10}$ 0.788	detectable strength
V3	strength dose at the fixed start's pattern: CI excludes 0 and $\geq 0.25\times$ gain	$+0.008$ $[0.001, 0.016]$ vs 0.023	flat
V4	$Q(25000) - Q(1000) \geq 0.10$, CI above 0, $Q(25000) \geq 0.60$	$0.18, 0.40, 0.57, 0.76$ at $n = 200, \ldots, 25000$; $+0.360$ $[0.294, 0.425]$	supported
V5	$Q(0.95) - Q(0.5) \geq 0.10$, CI above 0	$0.09 \to 0.51$; $+0.418$ $[0.349, 0.489]$	supported
slope	within-family slope of gain on exact gap (no threshold)	0.569 $[0.402, 0.731]$ (SC families 0.046 $[-0.347, 0.444]$)	reported
story table	V2 "detectable strength" and V4 supported	exact gap is the $n \to \infty$ ceiling; realised effect is its detectable image	U-A
Q2 pre-condition	in-trees: $A(\text{G1}) - A(\text{G2}) \geq 0.10$, CI above 0	DAGMA 0.564, NOTEARS 0.628, GOLEM-NV 0.448, GOLEM-EV 0.513	met for all four
HY2	$A(\text{G1}) > 0 > A(\text{G2})$, both CIs exclude 0	DAGMA $+0.200/-0.365$; NOTEARS $+0.234/-0.393$; GOLEM-NV $+0.123/-0.325$; GOLEM-EV $0.000/-0.513$	holds except GOLEM-EV
HY1	$R_{VM} \leq 0.40$ with CI upper bound ≤ 0.60 means "variance channel"	-1.25, -0.44, -0.41, -1.50 (DAGMA, NOTEARS, GOLEM-NV, GOLEM-EV)	variance channel (over-correction)
injection	$R_{inj} \geq 0.60$ means variance alone suffices	0.43, 0.49, 0.08, -0.19	not sufficient
HY3	FLOP's own shift, in-trees (no threshold)	-0.099 $[-0.119, -0.080]$ (G1−G2)	reported
check	Var-SortnRegress order accuracy within 0.03 of the target start	0.008, 0.001	passed
HB1	FLOP has the lowest mean SHD_{x1} among label-free methods in ≥ 6 of 8 groups	7 of 8; exception in-tree, fixed start (Var-SortnRegress 0.136 vs FLOP 0.251)	supported
C1–C7	Holm over seven claims at 0.05	adjusted p: 0.016 (V3), 0.0007 (V4, V5, C5, C6, C7), 0.0010 (slope)	all reject